
To Each (Textual Sequence) Its Own: Improving Memorized-Data Unlearning in Large Language Models

George-Octavian Bărbulescu¹ Peter Triantafillou¹

Abstract

LLMs have been found to memorize training textual sequences and regurgitate verbatim said sequences during text generation time. This fact is known to be the cause of privacy and related (e.g., copyright) problems. Unlearning in LLMs then takes the form of devising new algorithms that will properly deal with these side-effects of memorized data, while not hurting the model’s utility. We offer a fresh perspective towards this goal, namely, that each textual sequence to be forgotten should be treated differently when being unlearned based on its degree of memorization within the LLM. We contribute a new metric for measuring unlearning quality, an adversarial attack showing that SOTA algorithms lacking this perspective fail for privacy, and two new unlearning methods based on Gradient Ascent and Task Arithmetic, respectively. A comprehensive performance evaluation across an extensive suite of NLP tasks then mapped the solution space, identifying the best solutions under different scales in model capacities and forget set sizes and quantified the gains of the new approaches.

1. Introduction

LLMs are models typically built over transformer architectures and trained on a vast corpus of data (with up to trillions of tokens) (Radford et al., 2019; Chowdhery et al., 2023; Gao et al., 2020; Fedus et al., 2022; Kassem et al., 2023). Massive training data and a (lengthy) training process allow LLMs to establish factual associations and memorize language semantics and grammar (Zhang et al., 2021b).

Unfortunately, previous research has shown that LLMs memorized training examples which causes them to emit

¹Department of Computer Science, University of Warwick, Coventry, United Kingdom. Correspondence to: George-Octavian Bărbulescu <george-octavian.barbulescu@warwick.ac.uk>.

information verbatim as it exists in the training corpus when prompts are carefully designed. This, in turn, violates privacy regardless of any sophisticated adversarial attacks. These problems are juxtaposed against laws and regulations protecting the rights of individuals to be forgotten (RTBF) (Mantelero, 2013; Graves et al., 2021), that is, to have their information removed from the LLM. A straightforward approach to achieve this would be to retrain the LLM from ‘scratch’. Alas, this is infeasible in the context of LLMs, even with small/medium-sized models (millions of parameters). The problem is further exacerbated by the fact that the tendency of LLMs to regurgitate training data increases proportionally with the number of model parameters (Carlini et al., 2022b). Hence, the need to research efficient removal mechanisms of memorized data from pretrained LLMs, without requiring training the LLMs from scratch.

As mentioned, LLMs have been shown to (i) memorize and (ii) emit memorized training data at generation time, which causes privacy and copyright problems. As an example, GPT-3 (Brown et al., 2020) has been shown to generate PII information verbatim, raising significant concerns given its wide commercial usage (Heikkilä, 2022). Interestingly, memorized training data points can also decrease the quality of the model (Zhang et al., 2021b), by linking tokens which do not have a general meaning or a justifiable factual association, but only belong to an entity or an individual. But, given that memorization is the root cause of these problems, how can one intuitively define ‘memorization’? For LLMs, typically, if there exists a sufficient overlap between a training data point and a generated text response, then that data point is considered to have been memorized (Tirumala et al., 2022; Carlini et al., 2021; 2022b).

The focus of this research is to enrich LLMs with the ability to unlearn memorized textual sequences (a.k.a. training data points, examples, or samples), while satisfying three constraints: First, do so while maintaining the model’s utility for downstream applications unrelated to the forgotten information. Second, do so without having access to the rest of the pre-training dataset, as this may be either intractable or too large/costly to fine-tune over. Third, do so while avoiding the aforementioned problems with respect to privacy, copyrights, etc. As a real-world example, consider the need

to remove the effect of one or more textual sequences from the LLM for copyright reasons (Graves et al., 2021; Vilaronga et al., 2018). One approach to address this problem would be to place constraints on the original data (e.g. to obfuscate it) at inference time (Majmudar et al., 2022; Lu et al., 2020). However, such a solution would not protect from adversarial attacks to expose PII; to deal with such attacks, one must also erase the contribution of the targeted textual sequences from the model’s parameters. To protect against adversaries with white-box access to the model’s parameters, the community has developed algorithms that ensure differential privacy (DP) (Abadi et al., 2016; Yu et al., 2021; Li et al., 2021). Unfortunately, it has been discovered that training LLMs with DP can be very expensive (Anil et al., 2021). Also, given that the memorized information that needs to be removed is not known during training, it is impossible to draw a clear privacy boundary (Brown et al., 2022), and training with DP on the entire training dataset introduces noise which hurts model utility (Anil et al., 2021; Nguyen et al., 2022).

We focus on the problem of deleting memorized textual sequences from the models’ weights, as we deem it is a fundamental first step in unlearning in generative language models. We claim that there exists no satisfactory solution to this fundamental first step in unlearning in LLMs, as it depends, on the one hand, on the number of samples (a.k.a. data points, text sequences) to be forgotten, and, on the other hand, on the capacity (number of parameters) of the model. Existing state-of-the-art solutions include self-supervised learning techniques (Jang et al., 2022; Liu & Kalinli, 2023; Chen & Yang, 2023) and reinforcement learning methods (Kassem et al., 2023; Lu et al., 2022). To quantify the success of unlearning, previous methods holistically characterize the memorization of the whole *forget set*, that is, the set of textual sequences that are to be ‘unlearned’. Hence, the forget set is given a “memorization score” which is the average of the memorization scores of each textual sequence (Jang et al., 2022). If the average score is within an expected threshold, then the unlearning algorithm is deemed successful. Such expected thresholds are defined as the level of memorization one would expect to see for arbitrary text sequences unseen during training. We hypothesized and show with this work that looking at average memorization scores is not sufficient to conclude that forget-set examples are no longer extractable at generation time. This is in principle akin to the finding by Carlini et al. that average values are misleading when defending against Membership Inference Attacks (MIAs) (Carlini et al., 2022a). In our problem setting, intuitively, the distribution of memorization scores may be non-uniform (and indeed long-tailed). As such, it is possible that a subpopulation of textual sequences in the forget set is still memorized, while another subpopulation may have been completely forgotten. This leads to important

privacy loopholes: (i) highly memorized data points can be easily extracted; and (ii) even data points that are unlearned “too well” can still be extractable with a minimum-effort white-box access to the LLM. Alas, current state-of-the-art solutions are susceptible to these phenomena.

In light of the above, we propose to quantify the success of unlearning by looking at sequence-level memorization. We quantify unlearning success by counting the number of outliers with respect to a (desired) memorization threshold. For example, if the average memorization score across a large population of unseen (during training) textual sequences is, say, 30%, then the subpopulation of the forget set that scores above 60% after unlearning is still extractable. Furthermore, we propose two unlearning algorithms designed to meet this new criterion. The key aim behind them lies in providing fine-grained (per textual sequence) memorization control.

Contributions:

- We introduce a new metric to quantify the success of forgetting textual sequences in LLMs, which focuses on example-level memorization as opposed to aggregate memorization across forget-set examples.
- We devise a new Membership Inference Attack (MIA) for unlearning memorized data in LLMs, which shows evidence that existing SOTA unlearning algorithms, that focus on the aggregate memorization in the forget-set examples, are prone to privacy violations as a result of the presence of under- and over-memorized subpopulations of data points after unlearning.
- We introduce a novel metric to quantify model indistinguishability for unlearning algorithms in LLMs. Intuitively, indistinguishability is the closeness in the behaviour of two models: (i) the unlearned model and (ii) models retrained from scratch without the textual samples to be removed. To map behaviour, we inspect the model’s loss, which is a principled approach in SOTA methods from privacy literature.
- We introduce two (fairly simple) algorithms that provide the desired fine-grained, per-example unlearning.
- The two new algorithms piggyback upon known SOTA algorithms for unlearning in LLMs, and for detoxifying LLMs. Interestingly, we show that a solution for detoxifying LLMs can become a SOTA solution for unlearning in LLMs in specific circumstances.
- We study the solution space across model and forget-set sizes using nine NLP classification and four text generation tasks. We identify the best approach in different circumstances. We compare the new algorithms against existing SOTA unlearning solutions in LLMs, showing their superiority for both model utility and privacy metrics. So, in this case, “simple does it”!

Put together, we hope the above constitute actionable and significant steps towards understanding and solving

memorized-data unlearning in LLMs.

2. Background

2.1. Memorization in Large Language Models

Memorization in LLMs has been studied extensively as it comes with significant privacy and legal concerns (Carlini et al., 2021; 2022b; Jagielski et al., 2022; Tirumala et al., 2022). Pretrained LLMs have a tendency to regurgitate verbatim some of the training data, and their memorization is proportional to their number of parameters (Carlini et al., 2021; Lee et al., 2021). For instance, GPT-J which features 6 billion parameters has been shown to have memorized at least 1% (Carlini et al., 2022b; Wang, 2021; Black et al., 2021) of Pile (Gao et al., 2020), the dataset used during its training. This memorization, albeit unintended, is crucial to model natural language, and for generalisation (Zhang et al., 2021b; Feldman, 2020). A range of metrics have been proposed to compute the level of memorization across data points in causal language modelling (Tirumala et al., 2022; Carlini et al., 2022b; Jagielski et al., 2022), based on the likelihood of emitting textual sequences verbatim.

2.2. Machine Unlearning for Memorized Data

Machine unlearning has been receiving increasing attention due primarily to ensuring RTBF, but also for other applications (Kurmanji et al., 2023) - albeit unlearning research for image classification has matured more than unlearning in LLMs (see for instance the NeurIPS 2023 unlearning competition (Triantafillou et al., 2023)). Unlearning memorized data from LLMs is the field of study concerned with efficient methods for removing the contributions of training examples from the model (Maini et al., 2024; Neel & Chang, 2023; Zhang et al., 2024). Exact unlearning retrains LLMs from scratch by first processing the dataset (i.e., removing the training data points requested to be removed). However, this is a costly process, especially when it comes to re-training very large transformers, the current state-of-the-art architecture in NLP (Brown et al., 2020; Abadi et al., 2016). To mitigate the costs associated with retraining LLMs for each forget request, approximate unlearning methods have been proposed (Pawelczyk et al., 2023; Chen & Yang, 2023; Meng et al., 2022; Jang et al., 2022; Kassem et al., 2023) which remove the contribution of the target data points from the pretrained model.

3. Related Work

Related works can be taxonomized based on whether they: (i) require access to training datasets; (ii) target language generation; (iii) alter model weights or simply alter LLM output; and (iv) target unlearning memorized data or ‘higher-level’ entities such as entity or factual associations.

We focus on: **(i) algorithms which avoid accessing the pretraining data.** These are solutions which assume that retain data (training data minus forget set) is too large/costly to be acted upon, in contrast to unlearning research leveraging training datasets (Chen & Yang, 2023; Liu & Kalinli, 2023); **(ii) on language generation**, unlike unlearning methods designed for text classification (e.g., via *label flipping* (Pawelczyk et al., 2023)); **(iii) on methods that alter model weights.** Differential Privacy Decoding (DPD) methods (e.g., (Majmudar et al., 2022)) are inference-time techniques. They remain relevant to this work since they could be used to solve, for instance, copyright issues, oblivious to the fact that they do not alter model weights. Nonetheless, we will compare against DPD to showcase its privacy problems; **(iv) on unlearning memorized training text**, instead of removing factual associations or higher-level semantics such as bias and toxicity (Meng et al., 2022; Wang et al., 2023; Lu et al., 2022; Kassem et al., 2023). For instance, (Kassem et al., 2023) derive a paraphrasing policy by repeatedly sampling answers from the original LLM with respect to parts of the original data points (used as prompts), picking the sample least similar to the original. Akin to detoxifying or removing biases from LLMs, the method mitigates *approximate memorization*, i.e., higher-level similarity to the original training text (Ippolito et al., 2022), instead of *exact memorization* (Tirumala et al., 2022), namely, the likelihood of extracting or emitting a training textual sequence verbatim, which is a closer concern to data copyright and adversarial attacks.

The current SOTA approach for our task at hand (also regarded as SOTA in (Liu & Kalinli, 2023; Kassem et al., 2023; Pawelczyk et al., 2023) is the work of (Jang et al., 2022), who uses gradient ascent (GA) oblivious to the pre-training dataset to minimize exact memorization over a set of textual sequences. It shows that simply maximizing the loss with respect to the textual sequences to be forgotten produces a competitive trade-off between privacy and model utility. We will compare against this SOTA algorithm quantifying further improvements. Finally, Task Arithmetic (TA) (Ilharco et al., 2022) is a recent learning paradigm that builds on the intuition that learning a new task can be seen as a vector of weights in neural networks. Task arithmetic has shown impressive results for removing biases and detoxifying LLMs. Although these are very similar tasks to unlearning, using TA in LLMs for unlearning has never been investigated before. We will show this is unfortunate.

4. The New Approach

The codebase is available for reproducibility: ¹

¹https://github.com/GeorgeOctavian/selective_unlearning




Figure 1. The distribution of *exact memorization* scores across the forget set D_f when the average memorization threshold is $\leq 33\%$ (as exhibited by GPT-Neo with 1.3B parameters over unseen/arbitrary textual sequences (Jang et al., 2022)). While Equation (2) holds, 5 (8) samples from D_f have an exact memorization score of over 60% (under 10%).

4.1. Quantifying Memorization in LLMs

To quantify the dynamics of memorization during unlearning we adopt a variation of *label memorization* (Zhang et al., 2021a; Pondenkandath et al., 2018) produced by (Tirumala et al., 2022) to study the properties of neural LMs during training, referred to as *exact memorization*. The motivation for using *exact memorization* to characterise LLMs’ knowledge over a textual sequence is twofold: 1) it is prominent in previous work that explores copyright issues in LLMs (Jang et al., 2022; Tirumala et al., 2022), and 2) it works well in practical settings as it is inexpensive to compute.

Definition 4.1. Let $x = (x_1, x_2, \dots, x_T)$ be a textual sequence of T tokens with $x \in X$, where X is the set of all possible textual sequences over the LMs’ vocabulary. Let f_θ be a neural LM with parameters θ . Then *exact memorization* is a function $g : X \rightarrow \mathbb{R}^+$ where:

$$g(x) = \frac{\sum_{i=1}^{T-1} \mathbf{1}\{f_\theta(\cdot|x_{<i}) = x_i\}}{T-1} \quad (1)$$

Exact memorization is used in prior research to define/compute the success of forgetting in LLMs.

Definition 4.2. Let $D_f \subset X$ be the set of textual sequences to be forgotten and $D_t \subset X$, $D_t \cap D_f = \emptyset$, be a test set of textual sequences (i.e. text not used during the training of the model), then a *forgetting* algorithm is successful from a memorization standpoint if the following inequality holds:

$$\frac{1}{|D_f|} \sum_{x_f \in D_f} g(x_f) \leq \frac{1}{|D_t|} \sum_{x_t \in D_t} g(x_t) \quad (2)$$

The intuition behind the unlearning success threshold is that the population of textual sequences in the forget set D_f should be characterized by a memorization score lower or similar to that of the population of samples unseen during the LLMs training (see Table 1 for forgetting thresholds across the GPT-Neo family (Black et al., 2021)).

However, we intuit that an average memorization score over the whole forget set will in general not provide enough information needed to deem an unlearning algorithm’s execution as successful and is thus inappropriate. This is in principle akin to Carlini et al’s argument that average values are inappropriate to judge successful defence against Membership Inference Attacks (MIAs) (Carlini et al., 2022a). We showcase this with an experiment we conducted where the forget set D_f is composed of 32 textual sequences of $T = 200$ tokens, each picked arbitrarily from the Pile corpora, known to be easily extractable from the GPT-Neo family of models. Then, we apply gradient ascent (i.e., we maximise $\mathcal{L}_\theta(x) = -\sum_{t=1} \log f_\theta(x_t|x_{<t})$, where $x \in D_f$) to unlearn the textual sequences. The distribution of the memorization scores in D_f after unlearning is shown in Figure 1.

The figure depicts a clear non-uniform and highly variable distribution of the memorization scores across examples in D_f . Note that the algorithm has met the empirical threshold for unlearning success. Nonetheless, as discussed earlier (and as we will show later) such highly variable memorization scores of examples in D_f constitute loopholes that can lead to copyright liabilities and/or can be exploited to reveal private information. While some variation in memorization is to be expected across D_f , as some data points contain more common or generalizable n-grams, large variations in memorization scores are inherently problematic. This led our research to the question of whether there exists a class of unlearning algorithms that can further reduce the likelihood of *memorization-score outliers* (that is, training data points with very low or very high memorization scores during unlearning) which will ensure lower susceptibility to privacy attacks and copyright violations and without impacting model utility. For this we propose to use the following metric which simply counts such outliers after unlearning (i.e., when Equation (2) becomes true).

$$c(D_f) = \sum_{x \in D_f} \mathbf{1}\{g(x) \in ([0, lb] \cup [ub, 1])\} \quad (3)$$

Equation (3) counts the number of points with memorization scores after unlearning, m , where $m < lb$ or $ub < m$. The lower and upper bound thresholds lb and ub are found empirically and depend on the training data set. In our case, for example, we have chosen empirically $lb = 15\%$ and $ub = 60\%$. The motivation for counting data points with very low memorization scores is that these become extractable via adversarial attacks (as first shown using the Streisand effect in (Golatkar et al., 2020)), and we hypothesize these also impact model utility (e.g. by unlearning useful, recurrent n-grams). We have tested this hypothesis by performing a membership inference attack on the unlearned textual sequences (see Appendix G) and by measuring model indistinguishability on these (see Appendix H).

4.2. Selective Unlearning

Building on our insight that reducing memorization for privacy is a data-point-level problem and that unlearning success cannot be appropriately measured with an aggregate over the forget set examples (which is the approach commonly taken in computer vision unlearning (Kurmanji et al., 2023) and in previous unlearning research in LLMs (Jang et al., 2022; Kassem et al., 2023)), we propose two alternative unlearning algorithms to provide fine-grained exact memorization control.

4.2.1. SELECTIVE GRADIENT ASCENT

En route to devising our first algorithm, we visit gradient ascent (GA). GA is an ever-present method for unlearning in various settings. It is found, for instance, in unlearning in image classification literature as "NegGrad" in (Kurmanji et al., 2023) and in forgetting text in LLMs (Jang et al., 2022). Gradient ascent works by simply maximizing the loss function for examples in the forget set, instead of minimizing it, where the loss function is dependent on the task. In causal language modelling, unlearning minimizes:

$$\mathcal{L}_\theta(x) = \sum_{t=1}^{T-1} \log f_\theta(x_t | x_{<t}), \text{ where } x \in D_f \quad (4)$$

Note Equation (4) is the same as maximising the original log-likelihood function used in causal language modelling ($\mathcal{L}_\theta(x) = -\sum_{t=1}^{T-1} \log f_\theta(x_t | x_{<t})$, where $x = (x_1, x_2, \dots, x_T)$ is a sequence of tokens). If the number of examples in the forget set is larger than one, gradient ascent becomes oblivious to the individual memorization scores and stops when the average memorization across the forget set meets the empirical threshold for unlearning success.

To control the level of memorization we introduce Selective Gradient Ascent (SGA), which performs a two-step optimisation routine. Let γ be an upper limit for the acceptable exact memorization score of any given sample from the forget set after the process of unlearning, and let D_γ be the set of data points above the limit: $D_\gamma = \{x | g(x) > \gamma, x \in D_f\}$, where $g(x)$ computes the memorization score at the example level as per Equation (1). The unlearning algorithm begins by performing gradient ascent with respect to D_γ . After each epoch, D_γ is recomputed based on the entire forget set D_f . If there are no more elements above the unacceptable memorization threshold, $D_\gamma = \emptyset$, and inequality Equation (2) is false, we continue unlearning with gradient ascent with respect to the top-k memorized samples from the forget set, $D_{tk} = \arg \text{top}_k \{(x, g(x)) | x \in D_f\}$, recomputed after each epoch until the average memorization across the population of textual sequences reaches the desired level.

Discussion: SGA aims to provide fine-grained unlearning. The same cannot be achieved via sequential gradient ascent, i.e., unlearning each sequence one by one, as the data points

Table 1. The forgetting thresholds for GPT-Neo models on the Pile validation corpora, from (Jang et al., 2022). Average memorization across 10,000 unseen textual sequences.

NUMBER OF PARAMETERS	MEMORIZATION THRESHOLD
GPT-NEO (125M)	29.94%
GPT-NEO (1.3B)	33.27%
GPT-NEO (2.7B)	34.02%

in the forget set may share n-grams and their memorization scores may be linked. In this case, the order in which to unlearn data points is crucial and requires considering all possible orderings and suffering the associated combinatorial explosion in time. In this research, we merely touch the surface of our proposed paradigm, *selective unlearning*, which can be adapted to any gradient-based method.

4.2.2. TASK ARITHMETIC FOR UNLEARNING

A contrasting approach to unlearning information is to identify in the model’s parameter space the weights related to the forget set. This is a technique that has been used for unlearning in image classification settings (Golatkar et al., 2020; Thudi et al., 2022; Zhang et al., 2022; Chundawat et al., 2023; Kurmanji et al., 2023), as well as in LLMs. For LLMs, one technique to achieve this is to observe which parameters change when fine-tuning the model on the forget set. A technique that adopts this paradigm is task arithmetic (Ilharco et al., 2022).

Definition 4.3. Let $\theta \in \mathbb{R}^d$ be the parameters of a neural language model, where d is the number of parameters. Let \mathcal{A} be a deep learning algorithm such as fine-tuning with gradient descent. Then $\theta_{ft} = \mathcal{A}(\theta, D_f)$ are the model parameters after fine-tuning on the forget set. To forget this information, *task arithmetic* is expressed as:

$$\theta_u = \theta - (\theta_{ft} - \theta) \quad (5)$$

The intuition behind task arithmetic is to identify the direction in which the model’s parameters move when learning a new task. This is analogous to identifying the important weights for the task, which form a vector $\delta_{ft} = (\theta_{ft} - \theta)$. While removing this task vector from the original model as per Equation (5) has shown promise in detoxifying neural LMs (Ilharco et al., 2022), applying it for our unlearning task is not trivial. The forget set D_f is generally comprised of data points that are highly memorized, otherwise, there would be no need for unlearning in the first place. Therefore, it is crucial to reflect on whether the drift δ_{ft} is a good approximation of the learning movement that occurred when D_f was initially memorized.

We hypothesize the ‘drift’ when the level of memorization across the forget set is already very high (with respect to Equation (1)) cannot be used for unlearning purposes. The intuition is that the data is already learned well, and further

Table 2. Performance overview across the suite of unlearning algorithms. Unlearning 16 samples. RUNTIME is shown in minutes. At least one of the new algorithms we present (TA, TAU, and/or SGA) outperform GA for F1, CE, and ACC. For EXTRACTABLE, SGA is significantly superior to all, followed by TA/TAU, while GA has poor performance.

MODEL	ALGO	F1 ↑	CE ↓	ACC ↑	EXTRACTABLE ↓	RUNTIME ↓	ΔMEM ↓
GPT-NEO(125M)	DEFAULT	10.27	3.82	44.44	13.8	N/A	50.14
	DPD	6.23	N/A	N/A	8.8	N/A	15.44
	GA	2.28	6.18	40.34	5.0	7.22	3.34
	SGA	1.86	5.63	41.67	1.8	8.18	2.78
	TA	7.18	4.82	39.51	2.2	7.91	0.55
	TAU	7.37	4.44	43.68	3.4	50.74	6.18
GPT-NEO(1.3B)	DEFAULT	12.86	3.29	51.39	16.0	N/A	60.34
	DPD	6.52	N/A	N/A	2.4	N/A	10.95
	GA	11.34	3.52	49.37	4.6	16.23	4.06
	SGA	12.63	3.37	49.30	1.0	20.17	1.86
	TA	9.60	3.69	48.75	2.0	14.80	1.80
	TAU	12.52	3.39	49.24	4.0	29.96	0.81
GPT-NEO(2.7B)	DEFAULT	11.23	3.09	56.94	16.0	N/A	60.55
	DPD	6.27	N/A	N/A	1.0	N/A	3.73
	GA	10.99	3.27	56.25	7.0	22.14	6.22
	SGA	10.94	3.32	57.22	0.8	29.96	4.28
	TA	10.03	3.81	49.58	1.6	19.95	2.49
	TAU	11.34	3.27	55.14	4.0	36.54	1.52

fine-tuning on the forget set may lead to small changes in the parameter space. Furthermore, even if the average memorization score is not high, there could be outliers (see Figure 1), and these highly-memorized members become impossible to unlearn for the same reason.

Therefore, we first run a couple of epochs using SGA with a twofold objective in mind: 1) to standardize the distribution of memorization scores across the forget set, and 2) to decrease the average memorization.

Definition 4.4. Let $\theta_{sga} \in \mathbb{R}^d$ be the weights of the target model after applying SGA for a number of epochs k and let λ be a hyperparameter to adjust the size of the vector, then Task Arithmetic for Unlearning (TAU) is defined as:

$$\theta_u = \theta_{sga} - \lambda(\mathcal{A}(\theta_{sga}, D_f) - \theta_{sga}) \quad (6)$$

5. Experimental Evaluation

5.1. General Setup

5.1.1. MODEL FAMILY & FORGET SET

We use the GPT-Neo model family with 125 million (M), 1.3 billion (B), and 2.7 billion (B) parameters since: 1) they are known to emit verbatim responses and 2) they are the only large-enough language models used in previous research on unlearning for copyright purposes (Kassem et al., 2023; Jang et al., 2022). Thus, with the GPT-Neo family a fair comparison with previous work can be secured.

The textual sentences we unlearn from the models (i.e., set D_f) are picked arbitrarily from the Training Data Extraction Challenge, with samples publicly available at (Jang, 2022).

These samples contain 200 tokens each, and they are easy-to-extract sequences from the Pile Corpora (Gao et al., 2020), the dataset used to pretrain the GPT-Neo family.

5.1.2. BASELINES

We use SOTA unlearning methods in LLMs for unlearning memorized examples in the context of language generation. Given our taxonomy from Section 3, GA (Jang et al., 2022) is the SOTA. We also use TA (Ilharco et al., 2022) - by adapting it for unlearning training text from LLMs. We also consider differential privacy decoding (DPD) (Majmudar et al., 2022) to shed light in utility-privacy trade-offs in text generation tasks.

Our hardware setup and the hyperparameters are reported in Appendix C.

5.1.3. EVALUATION METRICS

To evaluate privacy, we report the number of ‘extractable’ outliers after applying each algorithm, computed as per Equation (3) with $lb = 15\%$ and $ub = 60\%$. We stop unlearning when Equation (2) becomes true or if the algorithm stops making progress. Given generally, our privacy objective is to mimic the average memorization score of an unseen population of textual sequences $\frac{1}{|D_f|} \sum_{x_f \in D_f} g(x_f) \approx \frac{1}{|D_t|} \sum_{x_t \in D_t} g(x_t)$; we report the distance between the two averages at the end.

To study the privacy-utility trade-offs, we use the target LLMs after unlearning to generate text in four popular




Figure 2. Average F1 scores for dialogue tasks. Unlearning 8 (left), 32 (middle), and 128 (right) samples. On the 125M model. The key discovery here is that TA and/or TAU perform best in small LLMs.

benchmarks: Wizard of Wikipedia (WoW)(Dinan et al., 2018), Wizard of Internet (WoI) (Komeili et al., 2021), Blended Skill Talk (BST) (Smith et al., 2020), and Empathetic Dialogues (ED) (Rashkin et al., 2018). We compute the F1 score (harmonic mean between precision and recall) and report the cross entropy (CE) with respect to the ground truth continuation for each dialogue.

Furthermore, we evaluate the target model for common sense and scientific reasoning post-unlearning by using a suite of popular NLP classification tasks: ARC-Easy and ARC-Challenge (Clark et al., 2018), COPA (Gordon et al., 2012), Lambada (Paperno et al., 2016), Winogrande (Sakaguchi et al., 2021), PubmedQA (Jin et al., 2019), Piqa (Bisk et al., 2020), Hellaswag (Zellers et al., 2019), and MathQA (Amini et al., 2019). Across all the classification tasks, we report the average accuracy.

5.2. Performance Overview

Experiment setup for Table 2. Five forget sets D_f^i where $i \in [0, 4]$, are constructed with 16 textual sequences, $\|D_f^i\| = 16$, sampled from the Pile corpora. Each unlearning algorithm operates on a forget set. For each algorithm the results are averaged over the 5 forget sets. At the end, we report the mean across the average results for the dialogue tasks (F1, cross-entropy (CE)), and the classification tasks (ACC). We also report Δ_{MEM} as the difference between the memorization threshold and the memorization average at the end, to uncover outliers (but do not use it to report privacy quality). Instead, we rely on the number of outliers (EXTRACTABLE) to quantify protection.

Key insights. We unveil six core insights. 1) GA and SGA destructively ensure privacy in small LLMs, sacrificing 80% of the language generation capabilities in the 125M parameter model. In contrast, the task arithmetic-based methods (TA, TAU) reduce the F1 score across the four dialogue tasks by 30%. The destructive nature of GA in small models has been independently recognised in previous work (Jang et al., 2022) (the work that proposed GA). Uncovering reli-

able unlearning solutions for small-to-medium-sized LLMs is paramount to enabling language modelling on hardware with limited computational resources. 2) Larger models are more robust to the process of unlearning. Even simple methods such as GA perform reasonably well when the number of tunable parameters is high. We highlight that our algorithms (TAU and SGA) ensure higher utility across all tasks than the baselines. For our algorithms, we remark that the average F1 score and the cross-entropy are aligned with the default model (i.e., the model before unlearning). For the 2.7B model, we observe TAU surpasses the language modelling capabilities of the original model. 3) The strongest privacy guarantees are brought by SGA across all experiments, with a maximum of 1.8 samples extractable under our memorization thresholds. While the thresholds are set empirically, this gives us a good intuition on whether the individual memorization scores across the individuals in the forget set are aligned with the forgetting threshold. 4) In contrast to the large drops in the ability to generate text, the model’s reasoning abilities (i.e., the average accuracy across the classification tasks) is more robust to unlearning (results in Appendix I). This is true for all model sizes. 5) Differential private decoding (DPD) and Task Arithmetic (TA) have a good privacy-to-utility trade-off in small models (125M) but fail when the number of parameters exceeds 1B. This can also be observed by comparing the average loss when moving from small to large models, in the case of TA. 6) On average, SGA and TAU incur a higher runtime overhead than GA and TA. Nonetheless, the runtime penalty is insignificant when putting it in the context of retraining from scratch or training with DP learning, and given the high utility and privacy benefits.

Appendix A has an in-depth evaluation of the language generation capabilities varying forget-set and model sizes.

5.3. Text Generation vs Forget Set Size

Experiment setup for Figure 2. We construct 15 forget sets, 5 for each cardinality from $\{8, 32, 128\}$. Text generation




Figure 3. Average F1 scores for dialogue tasks. Unlearning 32 samples. On the 1.3B (left) and 2.7B (right) models. The key discovery here is that SGA catches up and surpasses the rest for larger models.

after unlearning is tested on the four dialogue tasks.

Key insights. There are three overarching findings concerning the trade-off between text generation and the size of the forget set. 1) Generally, task arithmetic (TA), and task arithmetic for unlearning (TAU) perform best across different forget-set cardinalities. Task arithmetic performs best when the number of samples is relatively small (we refer to Table 2 for the results on 16 samples), and we observe from the left-most spider plot in Figure 2 it greatly surpasses the other baselines when the number of samples forgotten at once is 8. TAU outperforms TA when the number of samples increases. We observe TAU outperforms the baselines if 32 and 128 samples, respectively, are forgotten at once. 2) Notably, gradient-based unlearning algorithms improve their trade-off between privacy and utility when the forget set size increases. On average, GA surpasses TAU when the cardinality of the forget set is 128, in the BST task. 3) There is a steep contrast in the ability to generate text between unlearning methods. When the forget set is small, gradient-based methods lose ca. 80% in F1 for the 125M vs the 1.3B and 2.7B models (see Table 2).

5.4. Text Generation vs Number of Model Parameters

Experiment setup for Figure 3. We investigate unlearning methods w.r.t. their ability to generate text with different model sizes. We reuse the forget sets of 32 samples from above and unlearn these from the 1.3B and 2.7B GPT-Neo models. Table 2 has results for forget sets of 16 samples.

Key insights. There are 2 key findings: 1) SGA outperforms competitors in text generation when the number of parameters exceeds 1B. This is in contrast to our findings with small models (see Figure 2), where gradient-based methods underperform regardless of the forget set size. Looking at F1 scores at task level (e.g., empathetic dialogues, etc.) we conclude that SGA sacrifices less utility while ensuring the

strongest privacy guarantee, based on Equation (3)). For instance, when the forget sets have a cardinality of 16, only 1 textual sample is found outside the interval (15%, 60%). The same results are seen in big models as we vary the forget set size (see Appendix B). 2) In contrast to our findings from Table 2, increasing the size of the forget set in large models (from 16 to 32 here), TAU has a worse text generation utility-privacy trade-off than SGA.

6. Lessons Learned

To our knowledge, this is the first work to contribute the following findings: 1) The forget set may indeed contain extractable data points after applying various types of promising unlearning algorithms (GA, TA, DPD). We designed a new MIA against GA to show this (for space reasons, results are in Appendix G). 2) There exist unlearning algorithms that innately minimize the number of forget-set outliers (e.g., SGA), which performs excellently when the model and the forget-set size scale up. 3) While TA, as proposed by (Ilharco et al., 2022), does not scale well to large forget-set and model sizes, it is a strong performer for smaller models. This fills an important gap as SOTA unlearning solutions so far sacrifice the model’s utility in smaller models, hindering pushing LMs to devices with limited resources. 4) TAU becomes on par with SGA and even exceeds its performance when the model is large and forget sets are smaller. 5) There exists a unlearning algorithm (TAU) that has a 3X performance improvement in text generation than previous SOTA. 6) After unlearning, the LLMs’ reasoning abilities appear to degrade much less than the ability to generate text. This is an important result, which warrants more research.

7. Limitations

First, while we discover new unlearning algorithms with SOTA performance for unlearning memorized data from LLMs, we also unveil there is no single solution that performs best across the entirety of our problem space (i.e. different forget set sizes and model sizes). This finding is a consistent limitation in the Machine Unlearning research community (Kurmanji et al., 2023; 2024; Maini et al., 2024). Second, our research aims to pave the way for efficient unlearning of textual sequences. To this end, we quantify privacy as the propensity of the model to regurgitate the training data verbatim (a.k.a. memorization). An adjacent problem setting to ours is unlearning factual associations from the model. In this line of work, an attacker may try to extract the link between two entities (e.g. PII data associated with a particular individual from the pre-training corpora). While we do our best to show empirically an unlearned textual sequence (containing the PII information) will not be revealed verbatim by the LLM, a robust auditing process must be in place to ensure there are no further links between entities and private information. We deem our work as one of the first steps towards removing private information, and higher-level problems, such as eroding associations between entities, benefit from the initial steps we make in this direction.

8. Conclusion

We explored the problem of unlearning memorized data in LLMs from a fresh perspective: Each textual sequence to be forgotten is treated differently, based on its memorization score. We adopted the TA method (for detoxifying LLMs) and showed that it can be SOTA in certain important cases. We proposed two new methods: SGA and TAU, which introduce the new unlearning perspective into GA and TA. We devised a new MIA that shows that SOTA methods (like GA) that do not adopt this perspective are vulnerable to privacy attacks. We demonstrated clear performance gains of these methods for model utility and for our new metric (for quantifying success for privacy and copyright reasons) across an extensive suite of NLP tasks. This testifies to the importance of this fresh perspective.

Impact Statement

Recent advancements in deep learning and LLMs in particular have arrived at a fast pace, but come with increased responsibility: As researchers, we need to ensure the unintended behaviours of deep learning models are addressed thoroughly, be they harmful biases, privacy or copyright violations etc. Our research aims to take a fundamental step towards alleviating a class of privacy concerns from LLMs. For similar reasons, we do not store, process, or report any

qualitative data related to the Pile corpora (Gao et al., 2020), and only use the publicly available samples at (Jang, 2022).

References

- Abadi, M., Chu, A., Goodfellow, I., McMahan, H. B., Mironov, I., Talwar, K., and Zhang, L. Deep learning with differential privacy. In *Proceedings of the 2016 ACM SIGSAC conference on computer and communications security*, pp. 308–318, 2016.
- Amini, A., Gabriel, S., Lin, P., Koncel-Kedziorski, R., Choi, Y., and Hajishirzi, H. Mathqa: Towards interpretable math word problem solving with operation-based formalisms. *arXiv preprint arXiv:1905.13319*, 2019.
- Anil, R., Ghazi, B., Gupta, V., Kumar, R., and Manurangsi, P. Large-scale differentially private bert. *arXiv preprint arXiv:2108.01624*, 2021.
- Bisk, Y., Zellers, R., Gao, J., Choi, Y., et al. Piqa: Reasoning about physical commonsense in natural language. In *Proceedings of the AAAI conference on artificial intelligence*, volume 34, pp. 7432–7439, 2020.
- Black, S., Gao, L., Wang, P., Leahy, C., and Biderman, S. Gpt-neo: Large scale autoregressive language modeling with mesh-tensorflow. *If you use this software, please cite it using these metadata*, 58, 2021.
- Brown, H., Lee, K., Mireshghallah, F., Shokri, R., and Tramèr, F. What does it mean for a language model to preserve privacy? In *Proceedings of the 2022 ACM Conference on Fairness, Accountability, and Transparency*, pp. 2280–2292, 2022.
- Brown, T., Mann, B., Ryder, N., Subbiah, M., Kaplan, J. D., Dhariwal, P., Neelakantan, A., Shyam, P., Sastry, G., Askell, A., et al. Language models are few-shot learners. *Advances in neural information processing systems*, 33: 1877–1901, 2020.
- Carlini, N., Tramer, F., Wallace, E., Jagielski, M., Herbert-Voss, A., Lee, K., Roberts, A., Brown, T., Song, D., Erlingsson, U., et al. Extracting training data from large language models. In *30th USENIX Security Symposium (USENIX Security 21)*, pp. 2633–2650, 2021.
- Carlini, N., Chien, S., Nasr, M., Song, S., Terzis, A., and Tramèr, F. Membership inference attacks from first principles. In *2022 IEEE Symposium on Security and Privacy (SP)*, pp. 1897–1914. IEEE, 2022a.
- Carlini, N., Ippolito, D., Jagielski, M., Lee, K., Tramer, F., and Zhang, C. Quantifying memorization across neural language models. *arXiv preprint arXiv:2202.07646*, 2022b.

- Chen, J. and Yang, D. Unlearn what you want to forget: Efficient unlearning for llms. *arXiv preprint arXiv:2310.20150*, 2023.
- Chowdhery, A., Narang, S., Devlin, J., Bosma, M., Mishra, G., Roberts, A., Barham, P., Chung, H. W., Sutton, C., Gehrmann, S., et al. Palm: Scaling language modeling with pathways. *Journal of Machine Learning Research*, 24(240):1–113, 2023.
- Chundawat, V. S., Tarun, A. K., Mandal, M., and Kankanhalli, M. Can bad teaching induce forgetting? unlearning in deep networks using an incompetent teacher. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 37, pp. 7210–7217, 2023.
- Clark, P., Cowhey, I., Etzioni, O., Khot, T., Sabharwal, A., Schoenick, C., and Tafjord, O. Think you have solved question answering? try arc, the ai2 reasoning challenge. *arXiv preprint arXiv:1803.05457*, 2018.
- Dinan, E., Roller, S., Shuster, K., Fan, A., Auli, M., and Weston, J. Wizard of wikipedia: Knowledge-powered conversational agents. *arXiv preprint arXiv:1811.01241*, 2018.
- Fedus, W., Zoph, B., and Shazeer, N. Switch transformers: Scaling to trillion parameter models with simple and efficient sparsity. *The Journal of Machine Learning Research*, 23(1):5232–5270, 2022.
- Feldman, V. Does learning require memorization? a short tale about a long tail. In *Proceedings of the 52nd Annual ACM SIGACT Symposium on Theory of Computing*, pp. 954–959, 2020.
- Gao, L., Biderman, S., Black, S., Golding, L., Hoppe, T., Foster, C., Phang, J., He, H., Thite, A., Nabeshima, N., et al. The pile: An 800gb dataset of diverse text for language modeling. *arXiv preprint arXiv:2101.00027*, 2020.
- Gliwa, B., Mochol, I., Biesek, M., and Wawer, A. SAM-Sum corpus: A human-annotated dialogue dataset for abstractive summarization. In *Proceedings of the 2nd Workshop on New Frontiers in Summarization*, pp. 70–79, Hong Kong, China, November 2019. Association for Computational Linguistics. doi: 10.18653/v1/D19-5409. URL <https://www.aclweb.org/anthology/D19-5409>.
- Golatkar, A., Achille, A., and Soatto, S. Eternal sunshine of the spotless net: Selective forgetting in deep networks. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 9304–9312, 2020.
- Gordon, A., Kozareva, Z., and Roemmele, M. Semeval-2012 task 7: Choice of plausible alternatives: An evaluation of commonsense causal reasoning. In **SEM 2012: The First Joint Conference on Lexical and Computational Semantics—Volume 1: Proceedings of the main conference and the shared task, and Volume 2: Proceedings of the Sixth International Workshop on Semantic Evaluation (SemEval 2012)*, pp. 394–398, 2012.
- Graves, L., Nagisetty, V., and Ganesh, V. Amnesiac machine learning. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 35, pp. 11516–11524, 2021.
- Heikkilä, M. What does gpt-3 “know” about me, 2022.
- Ihharco, G., Ribeiro, M. T., Wortsman, M., Gururangan, S., Schmidt, L., Hajishirzi, H., and Farhadi, A. Editing models with task arithmetic. *arXiv preprint arXiv:2212.04089*, 2022.
- Ippolito, D., Tramèr, F., Nasr, M., Zhang, C., Jagielski, M., Lee, K., Choquette-Choo, C. A., and Carlini, N. Preventing verbatim memorization in language models gives a false sense of privacy. *arXiv preprint arXiv:2210.17546*, 2022.
- Jagielski, M., Thakkar, O., Tramer, F., Ippolito, D., Lee, K., Carlini, N., Wallace, E., Song, S., Thakurta, A., Papernot, N., et al. Measuring forgetting of memorized training examples. *arXiv preprint arXiv:2207.00099*, 2022.
- Jang, J. Knowledge Unlearning for Mitigating Privacy Risks in Language Models. <https://github.com/joeljjang/knowledge-unlearning>, 2022.
- Jang, J., Yoon, D., Yang, S., Cha, S., Lee, M., Logeswaran, L., and Seo, M. Knowledge unlearning for mitigating privacy risks in language models. *arXiv preprint arXiv:2210.01504*, 2022.
- Jin, Q., Dhingra, B., Liu, Z., Cohen, W. W., and Lu, X. Pubmedqa: A dataset for biomedical research question answering. *arXiv preprint arXiv:1909.06146*, 2019.
- Kassem, A., Mahmoud, O., and Saad, S. Preserving privacy through dememorization: An unlearning technique for mitigating memorization risks in language models. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pp. 4360–4379, 2023.
- Komeili, M., Shuster, K., and Weston, J. Internet-augmented dialogue generation. *arXiv preprint arXiv:2107.07566*, 2021.

- Kurmanji, M., Triantafillou, P., Hayes, J., and Triantafillou, E. Towards unbounded machine unlearning. In *Thirty-seventh Conference on Neural Information Processing Systems*, 2023. URL <https://openreview.net/forum?id=OveBaTtUAT>.
- Kurmanji, M., Triantafillou, E., and Triantafillou, P. Machine unlearning in learned databases: An experimental analysis. *Proceedings of the ACM on Management of Data*, 2(1):1–26, 2024.
- Lee, K., Ippolito, D., Nystrom, A., Zhang, C., Eck, D., Callison-Burch, C., and Carlini, N. Deduplicating training data makes language models better. *arXiv preprint arXiv:2107.06499*, 2021.
- Li, X., Tramer, F., Liang, P., and Hashimoto, T. Large language models can be strong differentially private learners. *arXiv preprint arXiv:2110.05679*, 2021.
- Liu, Z. and Kalinli, O. Forgetting private textual sequences in language models via leave-one-out ensemble. *arXiv preprint arXiv:2309.16082*, 2023.
- Lu, X., West, P., Zellers, R., Bras, R. L., Bhagavatula, C., and Choi, Y. Neurologic decoding:(un) supervised neural text generation with predicate logic constraints. *arXiv preprint arXiv:2010.12884*, 2020.
- Lu, X., Welleck, S., Hessel, J., Jiang, L., Qin, L., West, P., Ammanabrolu, P., and Choi, Y. Quark: Controllable text generation with reinforced unlearning. *Advances in neural information processing systems*, 35:27591–27609, 2022.
- Maini, P., Feng, Z., Schwarzschild, A., Lipton, Z. C., and Kolter, J. Z. Tofu: A task of fictitious unlearning for llms. *arXiv preprint arXiv:2401.06121*, 2024.
- Majmudar, J., Dupuy, C., Peris, C., Smaili, S., Gupta, R., and Zemel, R. Differentially private decoding in large language models. *arXiv preprint arXiv:2205.13621*, 2022.
- Mantelero, A. The eu proposal for a general data protection regulation and the roots of the ‘right to be forgotten’. *Computer Law & Security Review*, 29(3):229–235, 2013.
- Mattern, J., Mireshghallah, F., Jin, Z., Schölkopf, B., Sachan, M., and Berg-Kirkpatrick, T. Membership inference attacks against language models via neighbourhood comparison. *arXiv preprint arXiv:2305.18462*, 2023.
- Meng, K., Bau, D., Andonian, A., and Belinkov, Y. Locating and editing factual associations in gpt. *Advances in Neural Information Processing Systems*, 35:17359–17372, 2022.
- Neel, S. and Chang, P. Privacy issues in large language models: A survey. *arXiv preprint arXiv:2312.06717*, 2023.
- Nguyen, T. T., Huynh, T. T., Nguyen, P. L., Liew, A. W.-C., Yin, H., and Nguyen, Q. V. H. A survey of machine unlearning. *arXiv preprint arXiv:2209.02299*, 2022.
- Paperno, D., Kruszewski, G., Lazaridou, A., Pham, Q. N., Bernardi, R., Pezzelle, S., Baroni, M., Boleda, G., and Fernández, R. The lambada dataset: Word prediction requiring a broad discourse context. *arXiv preprint arXiv:1606.06031*, 2016.
- Pawelczyk, M., Neel, S., and Lakkaraju, H. In-context unlearning: Language models as few shot unlearners. *arXiv preprint arXiv:2310.07579*, 2023.
- Pondenkandath, V., Alberti, M., Puran, S., Ingold, R., and Liwicki, M. Leveraging random label memorization for unsupervised pre-training. *arXiv preprint arXiv:1811.01640*, 2018.
- Radford, A., Wu, J., Child, R., Luan, D., Amodei, D., Sutskever, I., et al. Language models are unsupervised multitask learners. *OpenAI blog*, 1(8):9, 2019.
- Rashkin, H., Smith, E. M., Li, M., and Boureau, Y.-L. Towards empathetic open-domain conversation models: A new benchmark and dataset. *arXiv preprint arXiv:1811.00207*, 2018.
- Sakaguchi, K., Bras, R. L., Bhagavatula, C., and Choi, Y. Winogrande: An adversarial winograd schema challenge at scale. *Communications of the ACM*, 64(9):99–106, 2021.
- Shetty, J. and Adibi, J. The enron email dataset database schema and brief statistical report. *Information sciences institute technical report, University of Southern California*, 4(1):120–128, 2004.
- Smith, E. M., Williamson, M., Shuster, K., Weston, J., and Boureau, Y.-L. Can you put it all together: Evaluating conversational agents’ ability to blend skills. *arXiv preprint arXiv:2004.08449*, 2020.
- Thudi, A., Deza, G., Chandrasekaran, V., and Papernot, N. Unrolling sgd: Understanding factors influencing machine unlearning. In *2022 IEEE 7th European Symposium on Security and Privacy (EuroS&P)*, pp. 303–319. IEEE, 2022.
- Tirumala, K., Markosyan, A., Zettlemoyer, L., and Aghajanyan, A. Memorization without overfitting: Analyzing the training dynamics of large language models. *Advances in Neural Information Processing Systems*, 35:38274–38290, 2022.

- Triantafillou, E., Pedregosa, F., Hayes, J., Kairouz, P., Guyon, I., Kurmanji, M., Dziugaite, G. K., Triantafillou, P., Zhao, K., Hosoya, L. S., Junior, J. C. S. J., Dumoulin, V., Mitliagkas, I., Escalera, S., Wan, J., Dane, S., Demkin, M., and Reade, W. Neurips 2023 - machine unlearning, 2023. URL <https://kaggle.com/competitions/neurips-2023-machine-unlearning>.
- Villaronga, E. F., Kieseberg, P., and Li, T. Humans forget, machines remember: Artificial intelligence and the right to be forgotten. *Computer Law & Security Review*, 34(2): 304–313, 2018.
- Wang, B. Mesh-Transformer-JAX: Model-Parallel Implementation of Transformer Language Model with JAX. <https://github.com/kingoflolz/mesh-transformer-jax>, May 2021.
- Wang, S., Zhu, Y., Liu, H., Zheng, Z., Chen, C., et al. Knowledge editing for large language models: A survey. *arXiv preprint arXiv:2310.16218*, 2023.
- Yao, Y., Xu, X., and Liu, Y. Large language model unlearning, 2024.
- Yu, D., Naik, S., Backurs, A., Gopi, S., Inan, H. A., Kamath, G., Kulkarni, J., Lee, Y. T., Manoel, A., Wutschitz, L., et al. Differentially private fine-tuning of language models. *arXiv preprint arXiv:2110.06500*, 2021.
- Zellers, R., Holtzman, A., Bisk, Y., Farhadi, A., and Choi, Y. Hellaswag: Can a machine really finish your sentence? *arXiv preprint arXiv:1905.07830*, 2019.
- Zhang, C., Bengio, S., Hardt, M., Recht, B., and Vinyals, O. Understanding deep learning (still) requires rethinking generalization. *Communications of the ACM*, 64(3):107–115, 2021a.
- Zhang, C., Ippolito, D., Lee, K., Jagielski, M., Tramèr, F., and Carlini, N. Counterfactual memorization in neural language models. *arXiv preprint arXiv:2112.12938*, 2021b.
- Zhang, J., Chen, S., Liu, J., and He, J. Composing parameter-efficient modules with arithmetic operations, 2023.
- Zhang, N., Yao, Y., Tian, B., Wang, P., Deng, S., Wang, M., Xi, Z., Mao, S., Zhang, J., Ni, Y., et al. A comprehensive study of knowledge editing for large language models. *arXiv preprint arXiv:2401.01286*, 2024.
- Zhang, Z., Zhou, Y., Zhao, X., Che, T., and Lyu, L. Prompt certified machine unlearning with randomized gradient smoothing and quantization. In Koyejo, S., Mohamed, S., Agarwal, A., Belgrave, D., Cho, K., and Oh, A. (eds.), *Advances in Neural Information Processing Systems*, volume 35, pp. 13433–13455. Curran Associates, Inc., 2022.

A. Accuracy for Text Generation Results

Table 3. Average F1 scores for the dialogue tasks: WoI, BST, ED, and WoW. Averages computed across five forget sets. TA is better for small models (125M). As forget-set sizes increase, TAU catches up to TA and surpasses it for larger models. SGA is weaker for smaller models and smaller forget sets. SGA catches up and surpasses GA as models and/or forgets sets get larger, where it becomes the winner overall. Overall, one or more of the new methods examined here, namely TA, TAU, and SGA, is the best performer.

MODEL	ALGO	WoI/F1 ↑	BST/F1 ↑	ED/F1 ↑	WoW/F1 ↑	AVG F1 ↑	FORGET SIZE
GPT-NEO(125M)	GA	5.25	6.11	4.87	4.40	5.16	
	SGA	3.68	3.28	2.88	3.11	3.24	
	TA	6.85	7.27	7.74	7.05	7.23	
	TAU	4.44	5.78	5.11	4.75	5.02	
GPT-NEO(1.3B)	GA	11.83	10.06	10.81	12.70	11.35	8
	SGA	12.52	11.51	12.21	12.85	12.27	
	TA	9.28	6.69	6.10	9.52	7.90	
	TAU	11.81	10.48	10.52	12.24	11.26	
GPT-NEO(2.7B)	GA	9.42	18.25	9.92	9.79	11.84	
	SGA	9.54	18.71	8.39	7.84	11.12	
	TA	5.77	14.56	6.43	7.43	8.55	
	TAU	7.74	19.18	9.74	8.22	11.22	
GPT-NEO(125M)	GA	2.72	2.23	2.37	2.23	2.28	
	SGA	2.77	0.90	1.71	2.06	1.86	
	TA	7.95	6.24	8.10	6.43	7.18	
	TAU	9.27	6.26	8.68	5.29	7.37	
GPT-NEO(1.3B)	GA	13.17	10.01	10.33	11.83	11.34	16
	SGA	13.45	12.08	10.94	14.03	12.63	
	TA	10.90	7.48	8.56	11.46	9.60	
	TAU	13.50	11.67	10.69	14.21	12.52	
GPT-NEO(2.7B)	GA	8.67	14.56	10.15	10.59	10.99	
	SGA	10.04	14.29	10.05	9.38	10.94	
	TA	7.26	13.72	7.81	11.34	10.03	
	TAU	10.39	14.57	9.35	11.03	11.34	
GPT-NEO(125M)	GA	2.20	3.37	2.94	2.88	2.85	
	SGA	4.07	5.64	4.24	4.62	4.64	
	TA	5.28	6.44	6.06	5.24	5.76	
	TAU	5.88	7.40	6.85	6.54	6.67	
GPT-NEO(1.3B)	GA	11.77	7.80	10.02	11.83	10.36	32
	SGA	11.38	10.20	10.68	12.39	11.17	
	TA	9.20	6.69	7.56	9.42	8.22	
	TAU	10.67	9.21	8.71	11.90	10.12	
GPT-NEO(2.7B)	GA	7.58	10.06	8.32	8.98	8.73	
	SGA	9.62	10.06	10.37	8.82	9.72	
	TA	6.16	8.35	8.38	5.96	7.21	
	TAU	7.91	9.17	8.73	7.32	8.28	

B. Extended Performance Evaluation

Table 4. Performance overview across the suite of unlearning algorithms when forgetting 8 samples at once (averaged over 5 forget sets). All algorithms perform comparably for CE and ACC. For F1, as before, in the vast majority of times, one of TA/TAU/SGA is the clear winner. For the other metrics there exist large differences. Note the EXTRACTABLE column: SGA outperforms all. And that TA/TAU outperform GA.

MODEL	ALGO	F1 \uparrow	CE \downarrow	ACC \uparrow	EXTRACTABLE \downarrow	RUNTIME \downarrow	Δ MEM \downarrow
GPT-NEO(125M)	GA	5.16	5.31	38.85	2.4	9.14	1.22
	SGA	3.24	6.02	38.36	0.2	10.66	1.96
	TA	7.23	4.66	40.10	1.0	8.12	0.97
	TAU	5.02	6.36	39.09	3.2	31.03	1.35
GPT-NEO(1.3B)	GA	11.35	3.50	50.83	3.0	15.02	4.04
	SGA	12.27	3.37	49.65	0.4	22.25	3.19
	TA	7.90	3.85	46.25	1.0	18.12	1.31
	TAU	11.26	3.46	50.66	1.6	39.83	2.21
GPT-NEO(2.7B)	GA	11.84	3.55	51.11	3.6	23.62	4.86
	SGA	11.12	3.42	53.33	0.6	25.88	3.62
	TA	8.55	4.19	43.34	1.0	19.93	4.03
	TAU	11.22	3.71	50.00	1.4	52.46	1.29

Table 5. Performance overview across the suite of unlearning algorithms when forgetting 32 samples at once (averaged over 5 forget sets). The conclusions are the same as in the results of Table 4.

MODEL	ALGO	F1 \uparrow	CE \downarrow	ACC \uparrow	EXTRACTABLE \downarrow	RUNTIME \downarrow	Δ MEM \downarrow
GPT-NEO(125M)	GA	2.85	6.22	41.56	9.0	9.06	2.20
	SGA	4.64	5.60	39.48	3.2	12.22	0.79
	TA	5.76	4.85	38.12	5.0	12.14	0.56
	TAU	6.67	4.56	42.22	6.2	71.28	7.63
GPT-NEO(1.3B)	GA	10.36	3.60	51.46	9.4	19.82	3.93
	SGA	11.17	3.48	49.51	2.2	29.09	1.29
	TA	8.22	4.00	44.37	2.6	19.69	1.10
	TAU	10.12	3.59	49.23	8.8	44.16	0.62
GPT-NEO(2.7B)	GA	8.73	3.70	53.95	11.6	23.10	5.36
	SGA	9.72	3.54	55.20	2.0	29.45	3.11
	TA	7.21	4.32	44.09	3.2	27.33	3.15
	TAU	8.28	3.63	53.89	10.0	40.60	1.49

C. Hyperparameters & Hardware Settings

The hyperparameters used in our experiments are reported in Table 6. Differential privacy decoding (DPD) uses a linear interpolation $\lambda = 0.2$ that allows the decoder to achieve the desired *exact memorization* thresholds (see Table 1 for the thresholds). We use the same learning rate for gradient ascent (GA) as in the original paper (Jang et al., 2022), given unlearning is performed over the same data distribution. For SGA we have a memorization upper limit of 70%. Subsequently, the loss is maximised with respect to the most memorised sample $k = 1$. For task arithmetic, we perform a grid search to find the size of the drift that can reduce the average memorization to the threshold $\lambda = 3$, which is more aggressive than the unlearning proposal from the original paper (Ilharco et al., 2022). For TAU, we perform SGA with the same hyperparameters as above until the memorization score becomes $Y_s = 50\%$. Subsequently, we scale the drift by a factor of $\lambda = 1$ during task arithmetic for the last part of the unlearning process.

In terms of hardware, we use two NVIDIA RTX A5000 with 24GB across all experiments.

Table 6. Hyperparameters

ALGO	HYPERPARAMETER	SYMBOL	VALUE
DPD	LINEAR INTERPOLATION FACTOR	λ	0.2
GA	LEARNING RATE	α	5E-5
SGA	LEARNING RATE	α	5E-5
	MEMORIZATION UPPER LIMIT	γ	70%
	TOP-K	K	1
TA	LEARNING RATE	α	5E-5
	DRIFT SCALING FACTOR	λ	3
TAU	LEARNING RATE	α	5E-5
	MEMORIZATION UPPER LIMIT	γ_u	70%
	MEMORIZATION TO SWITCH	γ_s	50%
	DRIFT SCALING FACTOR	λ	1

D. Comparing Against GA+Mismatch for Unlearning Memorized Data in LLMs

In (Yao et al., 2024) the GA algorithm (as inspired by (Thudi et al., 2022)) and a variation of GA, coined GA+Mismatch, were presented and studied as methods for unlearning in LLMs. GA and GA-Mismatch were shown to be top performers against finetuning. As the proposed and studied methods in our work (namely, TA, TAU, and SGA) outperform GA (Jang et al., 2022), we expected to see similar behaviour against GA+Mismatch. To substantiate this, we conducted experiments using GPT-NEO(1.3B) with the results presented and summarized in Table 7, where it is shown that GA+Mismatch (or MIS as we call it for brevity) lacks behind the new methods proposed here across all metrics.

Table 7. Comparing GA+Mismatch (Yao et al., 2024) (MIS) against GA, SGA, TA, and TAU in GPT-NEO(1.3B). We see that MIS is never the top performer and most often a bad performer. Specifically: For F1, it is either the 4th or 5th best (worst). For CE and ACC it is always the 5th best (worst). For EXTRACTABLE it is either 2nd- or 3rd-best. In addition, MIS is always worse than GA for F1, CE, and ACC. Notably, however, MIS outperforms GA for EXTRACTABLE, which, in hindsight, is to be expected as the added noise by the mismatch term in the loss function of MIS tends to protect privacy.

GPT-NEO(1.3B)	ALGO	F1 \uparrow	CE \downarrow	ACC \uparrow	EXTRACTABLE \downarrow	Δ MEM \downarrow
8 SAMPLES	GA	11.35	3.50	50.83	3.0	4.04
	SGA	12.27	3.37	49.65	0.4	3.19
	TA	7.90	3.85	46.25	1.0	1.31
	TAU	11.26	3.46	50.66	1.6	2.21
	MIS	8.51	4.03	43.23	0.8	3.15
16 SAMPLES	GA	11.34	3.52	49.37	4.6	4.06
	SGA	12.63	3.37	49.30	1.0	1.86
	TA	9.60	3.69	48.75	2.0	1.80
	TAU	12.52	3.39	49.24	4.0	0.81
	MIS	8.99	4.02	42.15	3.6	0.88
32 SAMPLES	GA	10.36	3.60	51.46	9.4	3.93
	SGA	11.17	3.48	49.51	2.2	1.29
	TA	8.22	4.00	44.37	2.6	1.10
	TAU	10.12	3.59	49.23	8.8	0.62
	MIS	8.83	4.03	41.74	2.4	4.44

E. The Effect of Parameter-Efficient Methods on Unlearning Memorized Data in LLMs

New work on parameter-efficient (finetuning) methods (a.k.a. PEFT) (Zhang et al., 2023) has shown that PEFT methods can help to detoxify LLMs. The studied methods there were LoRA and a new method proposed, called $(IA)^3$. As LoRA is shown to outperform $(IA)^3$, we conducted experiments to show how LoRA would affect the unlearning performance of GA, SGA, and TA in our problem setting. The results are shown and summarized in Table 8.

Table 8. Investigating the effect of a Parameter Efficient FineTuning (PEFT) method (Zhang et al., 2023) (LoRA) on unlearning memorized data in LLMs. The results show the effect of adding LoRA to GA, SGA, and TA versus GPT-NEO(125M). PEFT methods like LoRA cannot be viewed as unlearning algorithms per se. Nonetheless, they do affect the efficacy of unlearning. The results show that adding LoRA to (i) TA never improves TA for F1, ACC, CE, and EXTRACTABLE; (ii) SGA only benefits F1, but not ACC, CE, EXTRACTABLE or Δ_{MEM} ; (iii) GA benefits for F1, CE, and EXTRACTABLE, but not for ACC or Δ_{MEM} . As expected, LoRA can indeed improve significantly the runtime of algorithms: by ca. 2X for GA and SGA. For TA the improvement is much smaller. In general, the big picture is expected to be nuanced and a separate comprehensive investigation should be conducted to assess the impact of PEFT methods on unlearning algorithms.

PARAMETERS	ALGO	F1 \uparrow	CE \downarrow	ACC \uparrow	EXTRACTABLE \downarrow	RUNTIME \downarrow	Δ_{MEM} \downarrow
GPT-NEO(125M)	GA	2.85	6.22	41.56	9.0	9.06	2.20
	SGA	4.64	5.60	39.48	3.2	12.22	0.79
	TA	5.76	4.85	38.12	5.0	12.14	0.56
LoRA (R=32, $\alpha=64$)	GA+LoRA	7.39	5.88	36.11	7.0	5.26	3.70
	SGA+LoRA	8.56	5.68	36.08	5.0	6.23	3.90
	TA+LoRA	4.68	5.78	33.99	5.0	13.07	0.69

F. The Effect of Large Forget Sets

Table 9. The effect of larger forget sets. We notice that for larger forget sets, the conclusions with respect to the utility metrics (i.e., F1, CE, and ACC) among GA, SGA, TA, and TAU remain as before. Namely, the SOTA GA is never the winner - at least one of the new solutions we propose outperforms GA. Similarly for the forget quality metrics i.e., EXTRACTABLE, where SGA and TA outperform GA.

FORGET SET SIZE	ALGO	F1 \uparrow	CE \downarrow	ACC \uparrow	EXTRACTABLE \downarrow	RUNTIME \downarrow	Δ_{MEM} \downarrow
256	GA	8.42	4.40	42.40	11.6	10.82	4.29
	SGA	8.59	4.33	42.36	8.4	11.06	2.91
	TA	8.01	4.20	42.53	3.8	35.67	0.73
	TAU	9.15	4.10	44.90	6.4	65.45	0.38
512	GA	9.02	4.24	44.10	9.4	11.35	2.61
	SGA	9.17	4.20	44.06	6.6	11.30	1.93
	TA	7.08	4.46	43.61	4.6	16.21	1.67
	TAU	8.48	4.18	44.58	10.8	20.98	1.59

G. Adversarial Attacks for Textual Sequences

Why do textual sequences completely forgotten during unlearning raise privacy concerns? To answer this question, we revert our attention to membership inference attacks (MIAs), adversarial attacks aimed at extracting training data points. We set up a loss-based MIA for Unlearning in LMs, which we coin MIAU, inspired from (Mattern et al., 2023).

Definition G.1. Let $x \in X$ be a textual sequence, where X is the set of all possible textual sequences over the model’s vocabulary \mathcal{V} . Then we define \tilde{x} as a **neighbour** to x if \tilde{x} is constructed with a bidirectional encoder g from the corrupted version of \hat{x} (x with a percentage of tokens w_i masked). Formally, for each masked token in \hat{x} replace it with $\tilde{w}_i = g(\mathcal{V}^i, x)$

Definition G.2. Let $x = (w_1, w_2, \dots, w_T)$ be a textual sequence of T tokens with $x \in X$, where X is the set of all possible textual sequences over the model’s vocabulary \mathcal{V} . Let f_θ be a neural LM with parameters θ optimised under the causal language modeling loss function $\mathcal{L}_\theta(x) = -\sum_{t=1}^{T-1} \log f_\theta(w_t | w_{<t})$. Let $\tilde{x}_1, \tilde{x}_2, \dots, \tilde{x}_n$ be a set of n neighbours to x . Then we define **MIA for Unlearning (MIAU)** as:

$$\left| \mathcal{L}_\theta(x) - \frac{1}{n} \sum_{i=1}^n \mathcal{L}_\theta(\tilde{x}_i) \right| > \gamma \tag{7}$$

The intuition behind the decision law in Equation (7) is simple. The distance between the loss of the original textual sequence to the average loss of its neighbours is higher than a privacy threshold γ when the model has memorised the

original sequence or when the sequence has been completely removed, as we expect better generalisation over its neighbours in the latter case. This is because we treat the neighbours as validation points or, in Machine Unlearning parlance, as a proxy for ‘retrain from scratch’. Note that the memorization score defined in 1 is an approximation of the causal language modelling loss function, hence the reason we use loss and memorization interchangeably here.

It is evident from Figure 4 and Figure 5 that the members (i.e. the textual sequences) of the forget set become more extractable after unlearning if their memorization score is very low. This artifact of unlearning has been observed in previous Machine Unlearning work in computer vision and coined the “Streisand effect” (Golatkar et al., 2020). We are the first, to the best of our knowledge, to prove it occurs in LLMs and to take a first step towards devising unlearning solutions that prevent it.




Figure 4. From left to right: the loss distribution across 256 textual sequences by querying the original model (1.3B GPT-Neo checkpoint) and the model after unlearning. On the x-axis we see the distance of training samples (members) to their artificial neighbours (Equation (7)). We use 128 samples from the enron email dataset (Shetty & Adibi, 2004) as members and 128 samples from the SAMSsum training dataset (Gliwa et al., 2019) as nonmembers. After unlearning the members with gradient ascent, they become easily extractable (less overlap with nonmembers) as the distance between their loss and that of their neighbours’ becomes increasingly large.

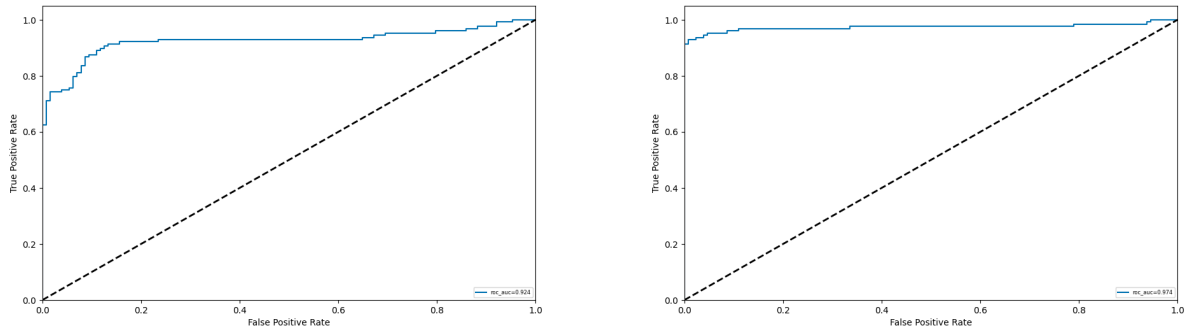


Figure 5. From left to right: the ROC curves for the experiment setup in Figure 4 for the original model and the unlearned model. Remarkably, there is a 50% increase in the success of the adversarial attack after unlearning, from a true positive rate (TPR) of 60% to a TPR of 90%, at a false positive rate (FPR) of 0%. And there is a clear difference on FPR in the FNR region [0, 0.2], which includes the regions we should attend to, according to Carlini et al (Carlini et al., 2022a).

H. Unlearning Effectiveness via Loss-Distance Magnitude

The membership inference strategy presented in Appendix G is used to ‘count’ the number of extractable points from the forget set after unlearning is applied. For a textual sequence to be qualified as a training or unlearned ‘member’, the absolute distance to its neighbours (Equation (7)) needs to be higher than a certain appetite for false positive rate, which is generally very low. We note the same distance can be used as a measurement of model ‘indistinguishability’ with respect to the forget set. We refer to model indistinguishability as the closeness in the behaviour of two models: (i) the unlearned model and (ii) models retrained from scratch without the forget set. In essence, the higher (lower) indistinguishability to retrain-from-scratch, implies that MIAs will be less (more) successful, leading to better (worse) privacy.

To quantify model behaviour we inspect the model’s loss, which is a principled approach and in the spirit of SOTA methods from privacy literature (Carlini et al., 2022a). To compare indistinguishability between two unlearning algorithms, A and B, we introduce the loss-distance magnitude (LOMA). We remark LOMA is independent of memorization (Equation (1)). LOMA is a fine-grained metric that tracks the distance to retrain-from-scratch by fixing the forget set and using the parameters at the end of the unlearning process in Equation (7) (i.e. θ_A and θ_B). Subsequently, we compare the two loss-distance distributions to the loss distances we would get for unseen data by the language model, which serves as our reference point for high indistinguishability (better privacy). A key finding is that feeding unseen data through Equation (7) results in consistent distances oblivious to θ . This is easily explained by the fact that unseen data had no impact on the parameters neither during pretraining nor unlearning.




Figure 6. We investigate the difference between the loss of our unlearned model (with GA and SGA) and a retrain-from-scratch proxy (proxy ‘data neighbours’ computed as shown in Appendix C). We do this for every sample in the forget set! Naturally, higher indistinguishability means much smaller distances (i.e., between 0 and 0.5 in our experiments on unseen data). The two figures show the loss distances of the forget set for the 125M and 1.3B checkpoints of GPT-Neo, aggregated across several forget sets/ unlearning runs.

It becomes obvious from Figure 6 that SGA provides crystal clear privacy benefits. In the GPT-Neo (125M) we observe 23 samples in the forget set reside within $[0, 0.25]$ with 15 samples in the first bin for SGA, compared to 6 for GA. Furthermore, we see GA has a long-tail distribution, with 13 sequences above a loss distance of 1, compared to only 5 for SGA. Similarly, the plot for 1.3B speaks for itself. There are 19 data points in the first bin $[0, 0.5]$ for SGA and only 7 for GA. Furthermore, there are 16 examples with a loss distance higher or equal to 4 for GA, whereas there are zero such examples for SGA. Generally, the population of the unlearned samples resides within small distances from the retrain-from-scratch proxy for SGA, while the distribution of distances for GA is much worse/flat with distances up to 9.5. While inspecting memorization can be used to extract data points (e.g. the points from the long-tail above), through this experiment we quantify the ‘magnitude’ of extractability and the distance in model behaviour from the optimal unlearned model (i.e. retrain-from-scratch).

I. Language Reasoning (Classification) Results




Figure 7. The average accuracy scores across nine language reasoning tasks. From left to right: unlearning 8 samples, 16 samples, and 32 samples at once, respectively. Results based on the 125 million checkpoint.




Figure 8. The average accuracy scores across nine language reasoning tasks. From left to right: unlearning 8 samples, 16 samples, and 32 samples at once, respectively. Results based on the 1.3B checkpoint.




Figure 9. The average accuracy scores across nine language reasoning tasks. From left to right: unlearning 8 samples, 16 samples, and 32 samples at once, respectively. Results based on the 1.3B checkpoint.

J. Loss on Dialogue Tasks




Figure 10. The average loss across the four dialogue tasks. From left to right: unlearning 8 samples, 16 samples, and 32 samples at once, respectively. Results based on the 125 million checkpoint.

To Each (Textual Sequence) Its Own: Improving Memorized-Data Unlearning in Large Language Models




Figure 11. The average loss across the four dialogue tasks. From left to right: unlearning 8 samples, 16 samples, and 32 samples at once, respectively. Results based on the 1.3B checkpoint.




Figure 12. The average loss across the four dialogue tasks. From left to right: unlearning 8 samples, 16 samples, and 32 samples at once, respectively. Results based on the 2.7B million checkpoint.