# Improving Factuality and Reasoning in Language Models through Multiagent Debate 

Yilun Du ${ }^{1}$ Shuang Li ${ }^{1}$ Antonio Torralba ${ }^{1}$ Joshua B. Tenenbaum ${ }^{1}$ Igor Mordatch ${ }^{2}$


#### Abstract

Large language models (LLMs) have demonstrated remarkable capabilities in language generation, understanding, and few-shot learning in recent years. An extensive body of work has explored how their performance may be further improved through the tools of prompting, ranging from verification, self-consistency, or intermediate scratchpads. In this paper, we present a complementary approach to improve language responses where multiple language model instances propose and debate their individual responses and reasoning processes over multiple rounds to arrive at a common final answer. Our findings indicate that this approach significantly enhances mathematical and strategic reasoning across a number of tasks. We also demonstrate that our approach improves the factual validity of generated content, reducing fallacious answers and hallucinations that contemporary models are prone to. Our approach may be directly applied to existing black-box models and uses identical procedure and prompts for all tasks we investigate. Overall, our findings suggest that such "society of minds" approach has the potential to significantly advance the capabilities of LLMs and pave the way for further breakthroughs in language generation and understanding.


## 1. Introduction

Large language models (LLMs) have demonstrated remarkable language generation, understanding, and few-shot learning capabilities in recent years. These methods are trained on a massive corpus of text on the internet, where the quality and accuracy of extracted natural language may not be ensured. Thus, current models may suffer from confidently

[^0]hallucinating facts or making implausible jumps in chains of reasoning. An extensive body of recent work has focused on improving factual accuracy and reasoning in language models. These range from prompting models with few or zero-shot chain-of-thought demonstrations, use of verification, self-consistency, or intermediate scratchpads.

We note that these techniques are applied over a single model instance. Instead, we propose a complementary approach inspired by The Society of Mind (Minsky, 1988) and multiagent settings, where multiple language model instances (or agents) individually propose and jointly debate their responses and reasoning processes to arrive at a common answer. More specifically, given a query, multiple instances of a language model first generate individual candidate answers to a query. Then each individual model instance reads and critiques the responses of all other models and uses this content to update its own answer. This step is then repeated over several rounds. This process induces models to construct answers that are consistent with both their internal critic as well as sensible in light of the responses of other agents. The resulting quorum of models can hold and maintain multiple chains of reasoning and possible answers simultaneously before proposing the final answer.

We find that our debate approach outperforms single model baselines such as zero-shot chain of thought (Kojima et al., 2022) and reflection (Shinn et al., 2023; Madaan et al., 2023) on a variety of six reasoning, factuality, and questionanswering tasks. Using both multiple model agents and multiple rounds of debate are important to achieve the best performance. Given an initial query, we find that individual model instances propose a diverse range of answers despite being the same model class (although we also investigate the case of mixing different model types, such as chatGPT (OpenAI, 2022) and Bard (Pichai, 2023)). After debating and examining the responses of other model instances, we find that the population almost always converges on a single and more accurate common answer. Debate results are also less likely to include false facts that models are internally uncertain of. This is because as the debate progresses, individual model instances tend to disagree on uncertain facts and omit them from the answer (Figure 12). Lastly, we find


Figure 1. Multiagent Debate Improves Reasoning and Factual Accuracy. Accuracy of traditional inference and our multiagent debate over six benchmarks (chess move optimality reported as a normalized score)
that debate does not just act to amplify one correct answer in a model quorum - we find many cases where all the models initially make incorrect predictions, but then arrive at the correct answer as the debate progresses (Figure 4,13).

We use the same methodology and prompt templates for all our tasks and require only black-box access to language model generations - no model-internal information such as likelihoods or gradients is needed. This allows our method to be used with common public models serving interfaces. The method is also orthogonal to other model generation improvements such as retrieval or prompt engineering (in fact, we combine our debate method with zero-shot chain of thought). While the debate process is more costly, requiring multiple model instances and rounds, it arrives at significantly improved answers and may be used to generate additional model training data, effectively creating a model self-improvement loop.

To help evaluate the effect of our approach on factual accuracy, we introduce a new benchmark and dataset evaluating the factual accuracy of famous computer scientist biographies. We find that contemporary language models have an especially high tendency to hallucinate factually incorrect biographies, often misrepresenting the relevant institutions and dates. Moreover, these facts are often inconsistent across different language model instances. By asking models to come to a consensus across their answers, such inconsistent facts may be either removed or corrected.

In summary, our work contributes the following. First, we present multiagent debate, an illustrate how it improves the reasoning accuracy in language models. Second, we illustrate how this similarity improves the factuality of language models and present a new benchmark of factual correctness which language models struggle with. Finally, we evaluate the performance of our debate procedure in language generation, both in terms of the number of agents, the underlying rounds of debate, and the prompts that elicit such behavior
across a set of six different reasoning and factual accuracy tasks.

## 2. Language Generation through Multiagent Debate

We present an approach to generate language responses through multiagent debate. We provide an overview of our approach in Section 2.1. We further discuss convergence to consensus in the debate process in Section 2.2. The overall overview of our approach is shown in Figure 2.

### 2.1. Multiagent Language Generation

Consider your work process when solving the following math question on an exam: "What is the area of a triangle with side lengths of 3, 4, 5?". In one thread of work, you may recognize that the triangle side-lengths directly correspond to a right triangle, and thus directly compute the area as $0.5 \times 3 \times 4=64$. To make sure that you have the right answer, you may then try to solve the problem differently by estimating an angle $\theta$ in the triangle using the Law of Cosines, and then obtain the area by using the formula $0.5 \times 3 \times 4 \times \sin (\theta)$, arriving at another answer to the given exam problem.

When these lines of work give the same answer, your confidence about the answer increases. In contrast, when these answers are different, individual lines of work may engage in a mental "debate" procedure, where you closely crossexamine the reasoning and assumptions of each line of work and refine solutions until a consistent answer.

Similarly, consider writing a biography of a historical figure. To ensure the factuality of the biography, you may consult multiple different sources on each fact. Facts that are consistent in each source increase your confidence about the facts. In contrast, facts that are inconsistent require careful cross-examination between sources to determine the final

User Input: A treasure hunter found a buried treasure chest filled with gems. There were 175 diamonds, 35 fewer rubies than diamonds, and twice the number of emeralds than the rubies. How many of the gems were there in the chest?
Round 1
Agent 1: Let $x$ be the number of rubies. Then the number of diamonds is 175 .
And the number of emeralds is $2(x-35)$, since there are twice as many
emeralds as rubies and there are 35 fewer rubies than diamonds.
To find the total number of gems, we add the number of diamonds, rubies,
and emeralds: Total gems $=175+x+2(x-35)=3 x+105$.
We don't know the exact value of $x$, but we do know that it represents the
number of rubies, which is less than $175($ since there are 35 fewer rubies than
diamonds). Let's try some values of $x:$ If $x=50$, then the total gems $=3(50)+$
$105=255$. If $x=100$, then the total gems $=3(100)+105=405$.
Since we know there were fewer rubies than diamonds, the answer must be
closer to 255 than to 405 . After some trial and error, we can find that $x=40$
gives us: Total gems $=3(40)+105=225$.

Agent 2: If there were 175 diamonds, then the number of rubies is 35 fewer, or 175-35=140 rubies.

And since there were twice as many emeralds as rubies, we can represent the number of emeralds as $2 * 140=280$ emeralds.

To find the total number of gems in the chest, we can add up the number of diamonds, rubies, and emeralds: $175+140+280=595$. Therefore, there were 595 gems in the chest.

User Input: Using the solutions from other agents as additional information, can you give an updated response.... (Copy Agent 2's reply)
(Copy Agent 1's reply)

| Round 2 |  |
| :---: | :---: |
| Agent 1: Given the information provided in other agents' solutions, we have two answers: One agent found that the number of rubies is 40 and the total number of gems is 225 . Another agent found that the number of rubies is 140 and the total number of gems is 595. <br> We know that the number of rubies is 35 fewer than the number of diamonds, which is 175 . So the number of rubies should be 175-35 $=140$, as the second agent found. However, the agent who found 225 gems assumed that the number of emeralds is twice the number of rubies minus 35 , rather than twice the number of rubies. If we use this assumption, we get: Number of rubies $=140$. Number of emeralds $=2 * 140-35=245$. Total number of gems $=175+140+245=560$. Our final answer is 560 . | Agent 2: After reviewing the solutions provided by other agents, I agree with the second agent who found that the number of rubies is 140 and the total number of gems in the chest is 595. Therefore, my answer to the math problem is 595 . |

User Input: Using the solutions from other agents as additional information, can you give an updated response....

| (Copy Agent 2's reply) | (Copy Agent 1's reply) |
| :--- | :--- |
| Agent 1: After reviewing the solutions provided by other agents, I agree <br> with the second agent who found that the number of rubies is 140 and the <br> total number of gems in the chest is 595 . Therefore, my answer to the math <br> problem is 595. | Agent 2: After reviewing the solutions <br> provided by other agents, I also agree <br> that the number of rubies is 140 and the <br> total number of gems in the chest is 595. |

Figure 2. Illustration of Debate. Illustration of the debate procedure. Prompts are shortened, see Table A4 for full prompts.
consistent data.
To mimic the above multi-threaded reasoning process and multi-source factuality checking processes, we propose to generate answers subject to a multiagent debate procedure between multiple instances of large language models. Given a question, multiple agents represented as copies of a large language model, generate answers to the question. Each response serves as a possible thought process or source of information which agents may re-examine to find consistent final answers.
After initial responses are generated from different agents, we initiate a round of debate between agents. Individual responses from other agents are concatenated and given as context to each agent, with each agent instructed to construct a new response based on such responses. Each language agent is thus responsible for both verifying the collection
of responses given by other agents, and refining its own response based on other agents' responses. We iteratively repeat this debate procedure over multiple rounds for improved performance. Such a procedure enables each language model to leverage both its opinions and the opinions of other agents to reflect and improve responses, enabling models to recover even if answers from both agents are incorrect (Figure A19, A20).

Concretely, we first prompt each agent to independently solve the given problem or task. After each agent generates a response, we feed each agent a consensus prompt, illustrated in Figure 3, where each agent is instructed to update their responses based on the responses of other agents. This resultant consensus prompt may then be repeatedly given, using the updated responses of each agent. We illustrate an overview of this multiagent debate procedure in Figure 2.

| Debate Type | Prompt |
| :---: | :---: |
| Short Debate | "These are the solutions to the problem from other agents: [other answers] |
| Based off the opinion of other agents, can you give an updated response . . " |  |

Figure 3. Prompts to induce long and short form debate. Responses of other agents to questions are inserted in the middle of the prompt (indicated with [other answers])


Figure 4. Illustration of Solving Math. Reasoning between agents is omitted.

Note that our proposed approach operates in an orthogonal manner to existing approaches to prompt language models. Given a question, we may apply additional techniques for prompting language models to further improve our debate procedure by eliciting additional more detailed responses from language models. We illustrate the synergy of our approach with existing approaches to prompting language models in Figure 6 and apply zero-shot chain-of-thought reasoning in our evaluations.

### 2.2. Consensus in Debates

Given multiple rounds of debate, how can we ensure that a set of language model agents will converge to a final consensus answer? In general, debate can be seen as a multiagent game, where convergence is not guaranteed. Empirically, however, we find that language models are able to converge on a single shared answer after multiple rounds of debate (Figure 4).

We found that we could control the duration of debates by changing how much a language model trusts its own outputs over those generated by other models through different prompts. We illustrate two prompts below in Figure 3, which we use to induce different debate durations between language models, and illustrate the effect of such prompts in Figure 10. Overall, we observed that language model agents were relatively "agreeable", perhaps as a result of instruction tuning or reinforcement learning based on human feedback (Ouyang et al., 2022). In cases of disagreement, we took the majority answer across agents at the end of debate.

## 3. Experiments

In our experiments, we evaluate our multiagent debate procedure and answer the following questions: (1) To what
extent does multiagent debate improve reasoning? (2) To what extent does multiagent debate improve factual validity? (3) What design choices enable multiagent debate to improve language generation performance?

### 3.1. Improving Reasoning with Multiagent Debate

We first evaluate the extent to which multiagent debate improves the underlying reasoning process in language models.

Tasks. We evaluate our approach on three reasoning tasks of increasing difficulty:

- Arithmetic. We first evaluate the ability of models to correctly evaluate an arithmetic expression (containing addition, multiplication, and subtraction) consisting of six different two-digit numbers. For example: What is the result of $12+15 * 21+0-3 * 27$ ?
- GSM8K. Next, we consider harder mathematical reasoning tasks. Using the GSM8K dataset (Cobbe et al., 2021), the models must correctly solve grade school mathematical reasoning tasks.
- Chess. Finally, we consider the strategic reasoning of the ability of models, and ask models to predict the best next move in a game of chess, given the first 14 moves of a chess game between two chess grand-masters described in PGN notation.

We report the accuracy of final answers in arithmetic and GSM8K tasks and report the pawn score (advantage) of predicted moves, as estimated by Stockfish in the Chess Moves.

Baselines. We compare our approach to three baseline approaches. First, we ask the language model to directly generate responses (single agent). Next, we consider asking language models to generate and then "self-reflect" on the responses generated (Madaan et al., 2023). Finally, we con-

| Model | Arithmetic (\%) $\uparrow$ | Grade School Math $(\%) \uparrow$ | Chess $(\Delta \mathbf{P S}) \uparrow$ |
| :--- | :---: | :---: | :---: |
| Single Agent | $67.0 \pm 4.7$ | $77.0 \pm 4.2$ | $91.4 \pm 10.6$ |
| Single Agent (Reflection) | $7.1 \pm 4.5$ | $75.0 \pm 4.3$ | $102.1 \pm 11.9$ |
| Multiagent (Majority) | $75.0 \pm 3.9$ | $81.0 \pm 3.9$ | $105.2 \pm 5.9$ |
| Multiagent (Debate) | $\mathbf{8 1 . 8} \pm \mathbf{2 . 3}$ | $\mathbf{8 5 . 0} \pm \mathbf{3 . 5}$ | $\mathbf{1 2 2 . 9} \pm \mathbf{7 . 6}$ |

Table 1. Multiagent Debate Improves Reasoning Multiagent debate improves the reasoning abilities of language models. Multiagent results in the table are run with 3 agents and two rounds of debate.


Figure 5. Illustration of Solving Grade School Math. Reasoning between agents omitted.
sider generating responses from multiple instances of model and performing majority voting across responses (Wang et al., 2022; Lewkowycz et al., 2022). In addition, in Table A8 in Appendix A.1, we provided additional comparison with an ensemble of "self-reflect agents". As the focus of our experiments is to verify the effectiveness of multiagent agent debate, we run both baselines and our approach, using the identical starting prompt and language model across all evaluations. We primarily evaluate models in a zero-shot setting, with details in Appendix A. 3 and additional fewshot results in Appendix A.1. Experiments are run using chatGPT-3.5 language model (OpenAI, 2022), with additional results with GPT-4 in Table A1 and Llama-7B in Table A5.

Due to computational expense, we evaluate our approach across benchmarks mainly using three agents with two rounds of debates, although we found further gains with both more agents and rounds of debate (Figure 9). Additional evaluation details are found in the Appendix A.2.

Quantitative Results. In Table 1, we report the results of each approach on arithmetic, grade school math, and chess reasoning tasks. In each task, we observe that utilizing multiple different agents to generate solutions improves performance over using a single language model agent to generate a solution. Simultaneously, we also see that reflection, where a language model is asked to critique its early generation, generally gives a modest boost in performance. Multiagent debate, which may be seen as a combination of both reflection and multiagent generation, gives a substantial boost in reasoning across each of the tasks.

Qualitative Results. In Figure 4 and 5, we provide qualitative illustrations of the debate procedure between models. Interestingly, we find cases in which all models initially give an incorrect response, yet the result of debate still obtains the correct answer as agents critique each others' reasoning. Thus, the purpose of our debate isn't just to amplify a correct answer - all models can initially be wrong but arrive at the correct answer through the debate process.


Figure 6. Synergy with Other Methods. Performance of debate increases with use of zero-shot Chain of Thought prompting on GSM8K. Results with few-shot Chain of Thought prompting can be found in Section A.1.

Compatibility with Other Reasoning Methods. Our multiagent generation procedure operates orthogonally approach to other prompting methods which focus on singleagent generation. In Figure 6, we illustrate the performance of multiagent debate with and without zero-shot chain-ofthought prompting (Kojima et al., 2022) on GSM8K. In addition, we report performance with few-shot chain-ofthought prompting in Section A.1. Across settings, multiagent debate is beneficial.

### 3.2. Extracting Factual Information from Multiagent Debate

We next evaluate how multiagent debate improves the underlying factuality in language models.
Tasks. We evaluate the factuality of language models in three different settings:

- Biographies. To evaluate the factuality of language models, we introduce a new task of accurately generating historical biographies of people, which we found existing language models to have a tendency to hallucinate. We constructed ground truth bullet point biographies of 524 well-known computer scientists. We then asked language models to generate bullet point biographies for each person, and evaluated answers with language model critic detailed in Appendix A.2.
- MMLU. Next, we assess the factuality of language models in responding to different factual knowledge questions typically learned and assessed in different exams. We utilize the existing MMLU dataset (Hendrycks et al., 2020) to benchmark the accuracy of responses.
- Chess Move Validity. Lastly, we study the hallucinations in language models when planning under to the

| Model | Biographies | MMLU | Chess Move Validity |
| :--- | :---: | :---: | :---: |
| Single Agent | $66.0 \pm 2.2$ | $63.9 \pm 4.8$ | $29.3 \pm 2.6$ |
| Single Agent (Reflection) | $68.3 \pm 2.9$ | $57.7 \pm 5.0$ | $38.8 \pm 2.9$ |
| Multiagent (Majority) | - | $67.0 \pm 4.7$ | $36.0 \pm 2.8$ |
| Multiagent (Debate) | $\mathbf{7 3 . 8} \pm \mathbf{2 . 3}$ | $\mathbf{7 1 . 1} \pm \mathbf{4 . 6}$ | $\mathbf{4 5 . 2} \pm \mathbf{2 . 9}$ |

Table 2. Multiagent Debate Improves Factual Accuracy. Multiagent debate improves the factual accuracy.

|  | Question: Six positive integers are written on the <br> faces of a cube. Each vertex is labeled with the <br> product of the three numbers on the faces adjacent to <br> the vertex. If the sum of the numbers on the vertices <br> is equal to 1001, then what is the sum of the numbers <br> written on the faces? A) 18. B) 13. C) 1001. D) 31. |  |
| :--- | :--- | :--- |
| Round 1 | Agent 1: A $\boxtimes$ | Agent 2: C $\triangle$ Agent 3: D |

[^1]Figure 7. Illustration of MMLU. Illustration of debate when answering factual tasks. Reasoning omitted.


Figure 8. Debate with Uncertain Answers. For uncertain facts, multiple language agents generate different facts. Debate always causes agents to change opinion and converge to one fact (which is not necessarily correct).
given rules of an existing environment or game. Specifically, we measure the validity of possible moves in a game of Chess given by BIG-Bench Chess-State Tracking Benchmark (Srivastava et al., 2022) task of chess-move prediction. In this task, an agent is given a set of next moves, and must make a valid next move of a piece on a board.

Baselines. We use the same baselines as in Section 3.1. Multiagent (majority) is omitted for the Biographies task as there is no method to take a majority vote over responses.
Results. We analyze the performance of each method in Table 2. We found that approaches based on reflection led to poor performance in the factuality setting. In contrast, debate gives the best performance in this setting also, and significantly outperforms each baseline. We illustrate a debate between agents on the biography task in Figure 12 and on MMLU in Figure 7. We find each agent after debate settled on bullets that were more consistent between agents.

We found that language agents tended to give different answers when the language model was uncertain about the question. However, directly asking each agent about their confidence of the answer (Kadavath et al., 2022) led to high confidence assessments on each answer. When these different language agents were asked instead to communicate with each other, each agent would quickly change their opinion to a consensus answer which was more accurate.

We illustrate this in Figure 8. Interestingly, we found that on facts that the language model was confident in (i.e. many instances of the same model all gave the same answer), it was very difficult to convince an agent to change their opinion, suggesting that "ease of persuasion" may be a method to assess factual confidence.

### 3.3. Analysis: Understanding Multiagent Debate

Finally, we analyze performance gains from debate.
Number of Agents. First, we analyze the impact of agent number on debate. In Figure 9(a), we increase agents used in debate, while fixing the debate length to be two. On Arithmetic, performance monotonically increases with the increased number of agents. We provide further analysis of performance with a large number of agents in Table A4.

Rounds of Debate. Next, we analyze the impact of the number of rounds of debate in multiagent debate. In Figure 9 (b), we increase the debate length between agents, while fixing the number of agents to three. We find that on the Arithmetic task, the performance also monotonically increases with debate length. Interestingly, we found that the confidence of language models in their responses also monotonically increased with debate (Figure A2), with the per token perplexity of generations decreasing 0.280 in round 1 to 0.203 in round 4 . This suggests debate helps guide the language model to confident answers.

Effect of Debate Length on Accuracy. As discussed in Section 2.2, the underlying convergence time in the debate between agents can be controlled by the extent to which agents are encouraged to maintain their opinions. In Figure 10, we consider the effect of short and long-form prompts discussed in Figure 3 on GSM8K. We find that debates using longer prompts lead to slower convergence to correct answers, but also lead to a better final consensus


Figure 9. (a) Performance with Increased Agents. Performance improves as the number of underlying agents involved in debate increases. (b) Performance with Increased Rounds. Performance rises as the number of rounds of underlying debate increases. Analysis in both settings on Arithmetic.


Figure 10. Performance vs Debate Length. Prompts which induce longer debate improve performance. Analysis on GSM8K.
on the correct answer. We provide an analysis of consensus between agents in Figure A1.

Using Different Initialization Prompts. In our experiments, we use the same prompts for all agents. We also consider the effect of using different questions, where we first instruct each language model to behave like a different persona (professor, doctor, mathematician) on the MMLU dataset. We found that improved performance on MMLU from 71.1 to 74.2 with different agents, suggesting further gains can be obtained with specialized agents.

Summarization. While in the majority of experiments in the paper we directly concatenate the responses of other agents as context for an agent to generate a new response, this is expensive when the number of agents involved in debate gets large. We may alternatively first summarize the responses from all other agents into a single response that we provide to agent at each round for more efficient debate. In Figure 11, we analyze the effect compared to directly concatenating the responses of other agents on the Arithmetic. We find this improves the performance of debate, suggesting that summarization is another tool that can further improve multiagent debate. We provide additional analysis of summarization on other tasks in Section A.1.

Using Different Language Models. Our existing debate results are reported using multiple instances of a chatGPT


Figure 11. Effect of Summarization. When there are many agents in a debate, responses from other agents may be first summarized and then given as context, reducing context length. This operation improves performance. Analysis on Arithmetic.
language model. We further assess the impact of using two different language models, where we ask chatGPT and Bard (Pichai, 2023) language models to debate with each other on a set of 20 GSM8K math problems. In this set, we find that multiagent debate improves the performance of both agents, with Bard solving 11 problems, chatGPT solving 14 problems, and joint multiagent debate solving 17 problems. We qualitatively illustrate a debate between agents in Figure 13. While both agents initially provide incorrect answers to the problem, chatGPT can use the response by Bard to generate the final correct answer. We further provide experiments using multiple instances of GPT-4 in Table A1 and multiple instances of Llama-7B in Table A5.

## 4. Related Work

Reasoning and Factuality in Language Models. A wide range of work has explored how to enable reasoning and factuality in language models. To improve reasoning, approaches have relied on techniques such as scratchpads (Nye et al., 2021), verification (Cobbe et al., 2021), chain-of-thought demonstrations (Wei et al., 2022; Kojima et al., 2022; Reynolds \& McDonell, 2021), and intermediate selfreflection (Shinn et al., 2023; Madaan et al., 2023) and finetuning (Lewkowycz et al., 2022; Rajani et al., 2019;

Prompt: Give a bullet point biography of David S. Johnson highlighting their contributions and achievements as a computer scientist, with each fact separated with a new line character.

| Round 1 |  |
| :--- | :--- | :--- |
| Agent 1: -- He earned his bachelor's degree from <br> Dartmouth College in 1967 and his Ph.D. from <br> MIT in 1973. | Agent 2: -- David S. Johnson is an American <br> computer scientist and mathematician. |
| -- Johnson has received numerous awards for his <br> research and contributions to computer science, <br> including the ACM Turing Award in 2013. | -- Johnson was elected to the National Academy of <br> Engineering in 2006. |
| -- Johnson's research interests include <br> optimization, algorithms, and complexity theory. | -- has been awarded numerous honors and <br> awards for his contributions to computer science, <br> including the ACM SIGMOD Innovations Award <br> and the Knuth Prize. |

\(\left.$$
\begin{array}{|l|l|l|}\hline & \text { Round 2 } \\
\hline \begin{array}{l}\text { Agent 1: -- David S. Johnson is a computer } \\
\text { scientist and mathematician. }\end{array} & \begin{array}{l}\text { Agent 2: -- David S. Johnson is an American } \\
\text { computer scientist and mathematician. }\end{array}
$$ <br>

\hline -- Johnson was elected to the National \& -- He has worked at several prestigious\end{array}\right\}\)| universities, including Columbia University, AT\&T |
| :--- |
| Bell Laboratories, and Cornell University. |

Figure 12. Illustration of Generating Biographies. Illustration of generating bullet biographies of computer scientists. For brevity, only the first 3 generated bullets are shown.

Zelikman et al., 2022). To improve factuality, approaches have relied on training techniques such as RLHF (Ziegler et al., 2019; Liu et al., 2022a; Christiano et al., 2017), pruning truthful datasets (Lee et al., 2022), external knowledge retrieval (Guu et al., 2020) and training-free methods based off likelihood estimation (Kadavath et al., 2022).

Our work provides an alternative way to obtain reasoning and factuality in language models using multiagent debate, which only requires black-box access to a language generator. Prior work also has explored how to take the majority vote across different models (Li et al., 2022b; Cobbe et al., 2021; Wang et al., 2022; Thoppilan et al., 2022; Lewkowycz et al., 2022) while in this work, we use rounds of debate between language model to combine answers. Most similar to our work, (Irving et al., 2018) also proposes a debate procedure to verify the accuracy and safety of powerful AI agents. In contrast to our approach, in their work, agents are asked to alternatively provide proof of input, and humans are tasked with assessing these debates.

Compositional Generation. Our work is also related to existing works that focus on text generation by combining different models (Du et al., 2020; Liu et al., 2022b; Zeng et al., 2022; Alayrac et al., 2022; Du et al., 2023). Most similar to our work, (Li et al., 2022a; Zeng et al., 2022) pro-
pose to combine multiple different large pretrained models together for multimodal reasoning. In contrast, in our work, we aim to use communication between different language models to enable more effective reasoning and factuality in language models.

Multiagent Language Models. Largely concurrent to the work in this paper, a body of recent work has explored how multiagent interaction can be used with large language models (Zhang et al., 2023; Li et al., 2024; Hong et al., 2023). Existing works have explored how language models can enable applications such as embodied cooperation (Zhang et al., 2023; Mandi et al., 2023; Chen et al., 2023), AI societies (Li et al., 2024), and software engineering (Hong et al., 2023; Wu et al., 2023). Most similar to our work, recent works have explored how debate can be applied to language models (Parrish et al., 2022; Michael et al., 2023; Khan et al., 2024). In Parrish et al. (2022) and Michael et al. (2023), debate between two unreliable human experts is used to facilitate a less-informed human judge to answer questions about a passage. While Parrish et al. (2022) finds limited gains from debate, Michael et al. (2023) finds that with more free-form debate humans judges are able answer more accurately. Subsequently, Khan et al. (2024) finds that by making language models more persuasive in such a debate formulation, the improvement in judgment increases,

## Question: Carla is downloading a 200 GB file. Normally she can download $2 \mathrm{~GB} /$ minute, but $40 \%$ of the way through the download, Windows forces a restart to install updates, which takes 20 minutes. Then Carla has to restart the download from the beginning. How load does it take to download the file?

| Round 1 |  |
| :---: | :---: |
| ChatGPT: Before the restart, Carla was able to download $40 \%$ of the file, which is $0.4 * 200 \mathrm{~GB}=80 \mathrm{~GB}$. At a rate of $2 \mathrm{~GB} /$ minute, it would take her: time $=80 \mathrm{~GB} \div 2$ $\mathrm{GB} /$ minute $=40$ minutes. <br> Because of the restart, Carla had to start over. This means she had to download the entire 200 GB file. The time it takes to download the remaining 120 GB is: time $=120 \mathrm{~GB}$ $\div 2 \mathrm{~GB} /$ minute $=60$ minutes . <br> But Carla also had to wait 20 minutes for the Windows update to install, so the total time it took to download the file is: total time $=40$ minutes +60 minutes +20 minutes $=$ 120 minutes. | Bard: Carla is downloading a 200 GB file at a speed of $2 \mathrm{~GB} /$ minute. This means that it will take her 200 / $2=100$ minutes to download the file. <br> However, Windows forces a restart to install updates, which takes 20 minutes. This means that Carla will have to restart the download from the beginning. Since Carla has already downloaded $40 \%$ of the file, she only needs to download 200-40 $=160 \mathrm{~GB}$. At a speed of $2 \mathrm{~GB} /$ minute, it will take Carla $160 / 2=80$ minutes to download the file. <br> Adding the time for the Windows restart, it will take Carla $80+20=\mathbf{1 0 0}$ minutes to download the file. |
| Round 2 |  |
| ChatGPT : Based on the calculations provided by the other agent, the correct answer is: Before the restart, Carla was able to download $40 \%$ of the file, which is $0.4 * 200 \mathrm{~GB}=80 \mathrm{~GB}$. At a rate of $2 \mathrm{~GB} /$ minute, it would take her: time $=$ $80 \mathrm{~GB} \div 2 \mathrm{~GB} /$ minute $=40$ minutes |  |
| However, because of the restart, Carla had to start over. This means she had to download the entire 200 GB file. The time it takes to download the entire 200 GB file is: time $=200 \mathrm{~GB} \div 2 \mathrm{~GB} /$ minute $=100$ minutes |  |
| And since Carla had to wait 20 minutes for the Windows update to install, the total time it took to download the file is: total time $=40$ minutes +100 minutes +20 minutes $=\mathbf{1 6 0}$ minutes. |  |

Figure 13. Debate Between chatGPT and Bard. Illustration of debate between different models.
suggesting that debate is a viable way to provide oversight on increasingly stronger models.

In contrast, our work and work in (Liang et al., 2023; Chan et al., 2023) focus on using debate at inference time to improve the performance of language models. Concurrent to our work, in Liang et al. (2023), debate is formulated by having one language model generate an answer, and a separate language model generate an adversarial variant of the answer. A moderator language model then mediates conversation between agents with debate proceeding until convergence between agents. In contrast, in our debate formulation, we ask a set of agents to generate a set of multiple answers given a question. We then implement debate between agents by directly asks agents to adjust their responses given the responses of all other agents across multiple set of rounds. We illustrate how this formulation improves reasoning and factuality in language models. Subsequently, Chan et al. (2023) adopts our formulation of debate, where a set of agents respond to the responses of other agents and applies it to evaluate the generations of language models.

## 5. Limitations and Conclusion

Limitations. In comparison with other prompting techniques, our approach is more expensive as it requires both multiple agents and a debate procedure - this can be omitted by distilling the final answer after debate to the original base
model. In addition, we found that as debates got longer, current language models struggled to fully process the context. Training language models with longer-context or summarizing previous responses may alleviate this difficulty.

Conclusion. In this paper, we presented an orthogonal approach to improve the performance of language models through the ideas of multiagent societies. We believe the perspective of having modular language agents working in combination to solve different difficult tasks will prove to be a fruitful area of research that is orthogonal to improving performance with larger amounts of computational training.

## Acknowledgements

We acknowledge support from NSF grant 2214177; from AFOSR grant FA9550-22-1-0249; from ONR MURI grant N00014-22-1-2740; and from ARO grant W911NF-23-10034. Yilun Du is supported by a NSF Graduate Fellowship.

## Impact Statement

In this paper, we have demonstrated a set of techniques for improving both the factuality and reasoning abilities of LLMs. As LLMs are increasingly being adopted in industry, such techniques can be helpful to ensure that responses are logical and factually grounding in the settings the models are being used.

## References

Alayrac, J.-B., Donahue, J., Luc, P., Miech, A., Barr, I., Hasson, Y., Lenc, K., Mensch, A., Millican, K., Reynolds, M., et al. Flamingo: A visual language model for fewshot learning. NeurIPS, 2022. URL https://arxiv. org/abs/2204.14198.8

Chan, C.-M., Chen, W., Su, Y., Yu, J., Xue, W., Zhang, S., Fu, J., and Liu, Z. Chateval: Towards better llm-based evaluators through multi-agent debate. arXiv preprint arXiv:2308.07201, 2023. 9

Chen, Y., Arkin, J., Zhang, Y., Roy, N., and Fan, C. Scalable multi-robot collaboration with large language models: Centralized or decentralized systems? arXiv preprint arXiv:2309.15943, 2023. 8

Christiano, P. F., Leike, J., Brown, T., Martic, M., Legg, S., and Amodei, D. Deep reinforcement learning from human preferences. In Neural Information Processing Systems, 2017. 8

Cobbe, K., Kosaraju, V., Bavarian, M., Chen, M., Jun, H., Kaiser, L., Plappert, M., Tworek, J., Hilton, J., Nakano, R., et al. Training verifiers to solve math word problems. arXiv preprint arXiv:2110.14168, 2021. 4, 7, 8

Du, Y., Li, S., and Mordatch, I. Compositional visual generation with energy based models. In Advances in Neural Information Processing Systems, 2020. 8

Du, Y., Durkan, C., Strudel, R., Tenenbaum, J. B., Dieleman, S., Fergus, R., Sohl-Dickstein, J., Doucet, A., and Grathwohl, W. Reduce, reuse, recycle: Compositional generation with energy-based diffusion models and mcmc. arXiv preprint arXiv:2302.11552, 2023. 8

Guu, K., Lee, K., Tung, Z., Pasupat, P., and Chang, M.W. REALM: Retrieval-augmented language model pretraining. arXiv preprint arXiv:2002.08909, 2020. 8

Hendrycks, D., Burns, C., Basart, S., Zou, A., Mazeika, M., Song, D., and Steinhardt, J. Measuring massive multitask language understanding. arXiv preprint arXiv:2009.03300, 2020. 5

Hendrycks, D., Burns, C., Kadavath, S., Arora, A., Basart, S., Tang, E., Song, D., and Steinhardt, J. Measuring mathematical problem solving with the math dataset. arXiv preprint arXiv:2103.03874, 2021. 14

Hong, S., Zheng, X., Chen, J., Cheng, Y., Wang, J., Zhang, C., Wang, Z., Yau, S. K. S., Lin, Z., Zhou, L., et al. Metagpt: Meta programming for multi-agent collaborative framework. arXiv preprint arXiv:2308.00352, 2023. 8

Irving, G., Christiano, P., and Amodei, D. Ai safety via debate. arXiv preprint arXiv:1805.00899, 2018. 8

Kadavath, S., Conerly, T., Askell, A., Henighan, T., Drain, D., Perez, E., Schiefer, N., Dodds, Z. H., DasSarma, N., Tran-Johnson, E., et al. Language models (mostly) know what they know. arXiv preprint arXiv:2207.05221, 2022. 6, 8

Khan, A., Hughes, J., Valentine, D., Ruis, L., Sachan, K., Radhakrishnan, A., Grefenstette, E., Bowman, S. R., Rocktäschel, T., and Perez, E. Debating with more persuasive llms leads to more truthful answers. arXiv preprint arXiv:2402.06782, 2024. 8

Kojima, T., Gu, S. S., Reid, M., Matsuo, Y., and Iwasawa, Y. Large language models are zero-shot reasoners. arXiv preprint arXiv:2205.11916, 2022. 1, 5, 7

Lee, N., Ping, W., Xu, P., Patwary, M., Fung, P. N., Shoeybi, M., and Catanzaro, B. Factuality enhanced language models for open-ended text generation. Advances in Neural Information Processing Systems, 35:34586-34599, 2022. 8

Lewkowycz, A., Andreassen, A., Dohan, D., Dyer, E., Michalewski, H., Ramasesh, V., Slone, A., Anil, C., Schlag, I., Gutman-Solo, T., et al. Solving quantitative reasoning problems with language models. arXiv preprint arXiv:2206.14858, 2022. 5, 7, 8

Li, G., Hammoud, H., Itani, H., Khizbullin, D., and Ghanem, B. Camel: Communicative agents for" mind" exploration of large language model society. Advances in Neural Information Processing Systems, 36, 2024. 8

Li, S., Du, Y., Tenenbaum, J. B., Torralba, A., and Mordatch, I. Composing ensembles of pre-trained models via iterative consensus. arXiv preprint arXiv:2210.11522, 2022a. 8

Li, Y., Choi, D., Chung, J., Kushman, N., Schrittwieser, J., Leblond, R., Eccles, T., Keeling, J., Gimeno, F., Dal Lago, A., et al. Competition-level code generation with alphacode. Science, 378(6624):1092-1097, 2022b. 8

Liang, T., He, Z., Jiao, W., Wang, X., Wang, Y., Wang, R., Yang, Y., Tu, Z., and Shi, S. Encouraging divergent thinking in large language models through multi-agent debate. arXiv preprint arXiv:2305.19118, 2023. 9, 13, 14

Liu, H., Lee, L., Lee, K., and Abbeel, P. Instructionfollowing agents with jointly pre-trained vision-language models. arXiv preprint arXiv:2210.13431, 2022a. 8

Liu, N., Li, S., Du, Y., Torralba, A., and Tenenbaum, J. B. Compositional visual generation with composable diffusion models. arXiv preprint arXiv:2206.01714, 2022b. 8

Madaan, A., Tandon, N., Gupta, P., Hallinan, S., Gao, L., Wiegreffe, S., Alon, U., Dziri, N., Prabhumoye, S., Yang, Y., et al. Self-refine: Iterative refinement with selffeedback. arXiv preprint arXiv:2303.17651, 2023. 1, 4, 7

Mandi, Z., Jain, S., and Song, S. Roco: Dialectic multirobot collaboration with large language models. arXiv preprint arXiv:2307.04738, 2023. 8

Michael, J., Mahdi, S., Rein, D., Petty, J., Dirani, J., Padmakumar, V., and Bowman, S. R. Debate helps supervise unreliable experts. arXiv preprint arXiv:2311.08702, 2023. 8

Minsky, M. Society of mind. Simon and Schuster, 1988. 1
Nye, M., Andreassen, A. J., Gur-Ari, G., Michalewski, H., Austin, J., Bieber, D., Dohan, D., Lewkowycz, A., Bosma, M., Luan, D., et al. Show your work: Scratchpads for intermediate computation with language models. arXiv preprint arXiv:2112.00114, 2021. 7

OpenAI. Chatgpt: Optimizing language models for dialogue, Dec 2022. URL https://openai.com/ blog/chatgpt/. 1,5

Ouyang, L., Wu, J., Jiang, X., Almeida, D., Wainwright, C. L., Mishkin, P., Zhang, C., Agarwal, S., Slama, K., Ray, A., et al. Training language models to follow instructions with human feedback. arXiv preprint arXiv:2203.02155, 2022. 4

Parrish, A., Trivedi, H., Nangia, N., Padmakumar, V., Phang, J., Saimbhi, A. S., and Bowman, S. R. Two-turn debate doesn't help humans answer hard reading comprehension questions. arXiv preprint arXiv:2210.10860, 2022. 8

Pichai, S. An important next step on our ai journey, Feb 2023. URL https: //blog.google/technology/ai/ bard-google-ai-search-updates/. 1, 7

Rajani, N. F., McCann, B., Xiong, C., and Socher, R. Explain yourself! leveraging language models for commonsense reasoning. arXiv preprint arXiv:1906.02361, 2019. 7

Reynolds, L. and McDonell, K. Prompt programming for large language models: Beyond the few-shot paradigm. In Extended Abstracts of the 2021 CHI Conference on Human Factors in Computing Systems, pp. 1-7, 2021. 7

Shinn, N., Labash, B., and Gopinath, A. Reflexion: an autonomous agent with dynamic memory and self-reflection. arXiv preprint arXiv:2303.11366, 2023. 1, 7

Srivastava, A., Rastogi, A., Rao, A., Shoeb, A. A. M., Abid, A., Fisch, A., Brown, A. R., Santoro, A., Gupta, A., Garriga-Alonso, A., et al. Beyond the imitation game: Quantifying and extrapolating the capabilities of language models. arXiv preprint arXiv:2206.04615, 2022. 6, 14

Thoppilan, R., De Freitas, D., Hall, J., Shazeer, N., Kulshreshtha, A., Cheng, H.-T., Jin, A., Bos, T., Baker, L., Du, Y., et al. Lamda: Language models for dialog applications. arXiv preprint arXiv:2201.08239, 2022. 8

Wang, X., Wei, J., Schuurmans, D., Le, Q., Chi, E., and Zhou, D. Self-consistency improves chain of thought reasoning in language models. arXiv preprint arXiv:2203.11171, 2022. 5, 8

Wei, J., Wang, X., Schuurmans, D., Bosma, M., Chi, E., Le, Q., and Zhou, D. Chain of thought prompting elicits reasoning in large language models. arXiv preprint arXiv:2201.11903, 2022. 7, 12

Wu, Q., Bansal, G., Zhang, J., Wu, Y., Zhang, S., Zhu, E., Li, B., Jiang, L., Zhang, X., and Wang, C. Autogen: Enabling next-gen 1 lm applications via multi-agent conversation framework. arXiv preprint arXiv:2308.08155, 2023. 8

Zelikman, E., Wu, Y., Mu, J., and Goodman, N. Star: Bootstrapping reasoning with reasoning. Advances in Neural Information Processing Systems, 35:15476-15488, 2022. 8

Zeng, A., Wong, A., Welker, S., Choromanski, K., Tombari, F., Purohit, A., Ryoo, M., Sindhwani, V., Lee, J., Vanhoucke, V., et al. Socratic models: Composing zero-shot multimodal reasoning with language. arXiv preprint arXiv:2204.00598, 2022. URL https:// arxiv.org/abs/2204.00598.8

Zhang, H., Du, W., Shan, J., Zhou, Q., Du, Y., Tenenbaum, J. B., Shu, T., and Gan, C. Building cooperative embodied agents modularly with large language models. arXiv preprint arXiv:2307.02485, 2023. 8

Ziegler, D. M., Stiennon, N., Wu, J., Brown, T. B., Radford, A., Amodei, D., Christiano, P., and Irving, G. Fine-tuning language models from human preferences. arXiv preprint arXiv:1909.08593, 2019. 8

## A. Appendix

In this appendix, we provide additional analysis and visualizations of the debates used in the main paper in Section A.1. We further provide detailed experimental details on each dataset in Section A.2. An anonymous repo with the code of the paper can be found in https://anonymous. 4open.science/r/llm_multiagent_debate_ anonymous-BE27/README.md.

## A.1. Additional Results

Consensus Between Agents. In Figure A1, we illustrate the consensus between agents using either short or long consensus prompts discussed in Figure 3 on the arithmetic dataset. The use of debate prompts that encourage agents to adapt more to the opinions of other agents improves consensus at the expense of worse performance (Figure 10). In general, we found high rates of consensus at the end of debate among agents across datasets with a consensus rate of $89.3 \pm 1.7$ on arithmetic, $88.6 \pm 1.8$ on GSM8K, $89.2 \pm 1.0$ on MMLU and $96.0 \pm 1.1$ on Chess Validity

Results with GPT-4. We further report results of our approach on the GPT-4 model in Table A1. Similar to the GPT-3 model, our approach also similarly improves the performance of the GPT-4 model.
Results with Few-Shot Chain-of-Thought Prompting. On GSM, we further consider applying few-shot chain-ofthought prompting to both our approach on baselines, using 8 -shot reasoning examples from (Wei et al., 2022). In this setting, we find a Single Agent obtains performance $81.0 \pm$ 3.9, Single Agent (Reflection) obtains performance $78.0 \pm$ 4.1, Multiagent (Majority) obtains performance $84.0 \pm 3.6$ and Multiagent Debate obtains performance $89.0 \pm 3.1$. We find that multiagent debate similarly improves performance with few-shot prompting.

Effect Short and Long Debate Prompts. We found that short debate prompts in Figure 3 consistently led to faster consensus and thus shorter debates than long debate prompts across GPT 3.5 and 4. After 3 rounds of debate on GSM8K problems, we find that GPT 3.5 with a short debate prompt obtained $96.1 \%$ consensus across agents while a long debate prompt led to $92.3 \%$ consensus across agents. We found the same trend in GPT 4, with a short debate prompt inducing $99.7 \%$ consensus across agents while a long debate prompt induced a consensus of $97.0 \%$ across agents.

Replication Across Multiple Splits of Data. To verify the reproducibility of our approach, we evaluate our approach across different splits of data on GSM8K in Table A2. Across each split, we find that our approach outperforms the performance of baselines.

Sensitivity to Selected Prompts. To evaluate the sensitivity of our approach to selected debate prompts in Table A4,


Figure A1. Effect of Prompts on Consensus. Using a short debate prompt induces faster consensus between agents. Analysis on GSM8K.


Figure A2. Effect of Debate on Perplexity. Using multiple rounds of debate decreases the per token perplexity of generations. This indicates that more rounds of debate increase the confidence of language model responses. Analysis on Arithmetic.
we paraphrase prompts each domain first using chatGPT-3.5 and report results in Table A3. Our approach is not sensitive to the precise prompts used for the debate procedure and retains good performance.

More Agents. To evaluate the performance of our approach with a large number of agents, we compare using 10 agents across 2 rounds of debate to a majority vote over 50 agents in Table A4. Our approach outperforms majority voting even under a very large number of agents.

Open Source LLMs. We further illustrate the applicability of our approach to the chat-Llama 27B in Table A5. Our approach can also be applied to existing opensource LLMs.

Summarization. We further evaluate the performance of summarization on 3 agents and 2 rounds of debate on the GSM8K and MMLU tasks. We find that summarization on GSM leads to a performance of 87.0 compared to the original performance of 85.0. We found that summarization on MMLU leads to a performance boost of 73.0 compared to the original performance of 71.1 .

Agreeableness. We analyze the effect of agreeableness on debate on the RLHF-aligned chat-Llama 7B model compared to the unaligned Llama 7B model. On the GSM8K
task, we find that LLama 7B model achieves a consensus of 51.3 after debate compared to consensus of 62.7 with aligned chat-LLama 2 model. On the MMLU task, we find that LLama 7B model achieves a consensus of 74.3 after debate compared to consensus of 100.0 with aligned chat-LLama 2 model. Our results indicate that RLHF significantly improves consensus in multiagent debate.

Comparison with Other Debate Methods. We compare our approach to debate to a concurrent approach to multiagent debate in (Liang et al., 2023) which uses an affirmative, negative and moderator agent. In Table A6, we find that our approach substantially outperforms Liang et al. (2023), as it more effectively encourages language models to focus on reasoning.

Effect of Reasoning Chain on Debate. In our implementation of multiagent debate, we concatenate both the reasoning chain of other language model agents as well as their final answer. In Table A7, we explore the ablative effect of only providing models with the final answer from other agents. We find that providing the reasoning chain of other agents substantially improves performance.

Reflection Ensemble. We further compare our approach to constructing an ensemble of self-reflection agents in Table A8. We find that our approach outperforms this baseline, indicating the efficacy of reflection across the responses of multiple agents.

Additional Qualitative Visualizations. We added additional qualitative visualizations of the debate process. In Figure A6, Figure A7, Figure A8, Figure A9, Figure A10, we illustrate debates between agents in the GSM8K dataset which result in the correct answer. In Figure A11, Figure A12, Figure A13, we further illustrate debates in GSM8K which lead to the incorrect answer. We further provide an example illustration of debate in arithmetic in Figure A14, arithmetic with summarization of individual responses of agents in Figure A15, MMLU in Figure A16, a debate with the full contents biographies in Figure A17, and debate in chess in Figure A18. We show examples where debate is able to correct two incorrectly generated answers in Figure A19 and Figure A20. In general, we found that debate improved the performance of final generated answers, though sometimes answers would converge to the incorrect value.

## A.2. Evaluation Details

We provide detailed evaluation details for each setting in the paper. We run all experiments using the gpt-3.5-turbo-0301 model. We provide a table listing the prompts used to prompt models and initialize debate in Table A4. All baselines and models were evaluated on the same set of problems.

Arithmetic. To evaluate the arithmetic task, we generated six random integers for each task between 0 and 30 . We then evaluated the extent to which the correct integer answer was correctly obtained. We evaluated models on one hundred randomly generated arithmetic tasks.

Grade School Math. To evaluate the GSM8K task, we evaluated the accuracy at which models were able to obtain the final correct answer, as extracted from a box. We evaluated models on one hundred randomly selected grade school math problems.

Chess. To evaluate the chess reasoning task, we used chess games from https://www.pgnmentor.com/ players/Adams.zip. We asked chatGPT to predict the next move for white to move at turn 14 and reported the relative Stockfish pawn score with search depth 20 after executing the suggested move from chatGPT. We evaluated models on three hundred randomly selected chess games.

Biographies. To evaluate the biographies task, we use a chatGPT critic to automatically evaluate the accuracy of generated biographies. We first extract a set of ground truth facts about a person from Wikipedia. For each ground truth fact, we ask chatGPT to assess if the fact is consistent with a full generated biography. We prompting chatGPT with the prompt: Consider the following biography of <person>: <generated biography> Is the above biography above consistent with the fact below? <ground truth bullet> Give a single-word answer, yes, no, or uncertain. We then evaluate and report the percentage of ground bullets that chatGPT returns either yes or no on. We provide two example judgements in Figure A5.

We conducted a human study on a total of 100 judgements in the biograph task using this automated evaluation from chatGPT. We found that chatGPT was very acccurate on this task, and provided the correct judgement in 93 out of the 100 cases. In the 7 incorrect judgements, we found in 6 of them, chatGPT was overly conservative and would assert that a fact was inconsistent with a generated biography when in fact it was consistent (but some part of the fact was missing from the generated biography). In the last incorrect judgement, we found that chatGPT incorrectly assessed that a fact stating a person was dead was consistent with a generated biography that said the person was in a current working occupation. As a result, we believe that this evaluation metric is a fast and relatively accurate way to assess the the accuracy of fact generation.

MMLU. To evaluate MMLU, we measured the accuracy in which models were able to select the correct multiplechoice answer in each problem. We evaluated models on one hundred randomly selected MMLU questions randomly distributed across each of the subject areas.

Chess Validity. To evaluate chess validity, we consider the BIG-Bench Chess-State Tracking Benchmark (Srivastava et al., 2022), where we used the hardest reported task in the benchmark synthetic_short. Each generated answer was deemed correct as long as it was one of the valid answers in the sequence. We evaluated models on one hundred randomly selected chess validity tasks.

## A.3. Baseline Details

Single Agent. For the single agent baseline, we use the same starting prompt as used in debate (illustrated in Figure A4).

Self-Reflection. For the self-reflection baseline, we use the same starting prompt as used in debate (illustrated in Figure A4). We then use the self-reflection prompt in Figure A3. We use one round of self-reflection.

Multiagent Majority. For the multiagent majority, we prompt each model with the same starting prompt as used in debate (illustrated in Figure A4). We then take the majority vote across all generations, where we use a total of 6 agents to match the cost of using 3 agents with 2 rounds of debate for multiagent debate.

## A.4. Computational Cost

We report per method and per dataset generation token costs of single agent, self-reflection, multiagent majority methods to answer a question in Table A9. Our reported generated tokens count is summed across generations from all agents and across all rounds of debate.

| Model | MMLU | MATH |
| :--- | :---: | :---: |
| Single Agent | $82.0 \pm 3.8$ | $47.0 \pm 4.9$ |
| Single Agent (Reflection) | $83.0 \pm 3.7$ | $50.0 \pm 5.0$ |
| Multiagent (Majority) | $84.0 \pm 3.7$ | $45.0 \pm 4.9$ |
| Multiagent (Debate) | $\mathbf{8 8 . 0} \pm \mathbf{3 . 2}$ | $\mathbf{5 6 . 0} \pm \mathbf{5 . 0}$ |

Table A1. Multiagent Debate on GPT-4. Our approach also improves the performance of a GPT-4 model on the MMLU and MATH dataset (Hendrycks et al., 2021).

| Model | GSM8K 1 | GSM8K 2 | GSM8K 3 |
| :--- | :---: | :---: | :---: |
| Single Agent | $77.0 \pm 4.2$ | $71.0 \pm 4.5$ | $73.0 \pm 4.4$ |
| Single Agent (Reflection) | $75.0 \pm 4.3$ | $70.0 \pm 4.6$ | $75.0 \pm 4.3$ |
| Multiagent (Majority) | $81.0 \pm 3.9$ | $76.0 \pm 4.3$ | $76.0 \pm 4.3$ |
| Multiagent (Debate) | $\mathbf{8 5 . 0} \pm \mathbf{3 . 5}$ | $\mathbf{8 1 . 0} \pm \mathbf{3 . 9}$ | $\mathbf{8 4 . 0} \pm \mathbf{3 . 6}$ |

Table A2. Evaluation on Different Splits of GSM8K. Our approach consistently outperforms each baseline across different splits of GSM8K problems

| Model | Biographies | Arithmetic | GSM8K |
| :--- | :---: | :---: | :---: |
| Single Agent | $66.0 \pm 2.2$ | $67.0 \pm 4.7$ | $77.0 \pm 4.2$ |
| Single Agent (Reflection) | $68.3 \pm 2.9$ | $72.1 \pm 4.5$ | $75.0 \pm 4.3$ |
| Multiagent (Majority) | - | $74.0 \pm 4.3$ | $81.0 \pm 3.9$ |
| Multiagent (Debate) | $\mathbf{7 3 . 8} \pm \mathbf{2 . 3}$ | $81.8 \pm 2.3$ | $85.0 \pm 3.5$ |
| Multiagent (Debate/Paraphrase) | $72.1 \pm 2.4$ | $\mathbf{8 4 . 0} \pm \mathbf{2 . 1}$ | $\mathbf{8 6 . 0} \pm \mathbf{3 . 4}$ |

Table A3. Debate Performance with Paraphrased Prompts. Our approach performs well even when debate prompts are paraphrased.

| Model | Arithmetic | GSM8K | MMLU |
| :--- | :---: | :---: | :---: |
| Majority Vote (50 Agents) | $92.0 \pm 2.7$ | $85.0 \pm 3.6$ | $67.0 \pm 4.7$ |
| Debate (10 Agents 2 Rounds) | $\mathbf{9 6 . 0} \pm \mathbf{1 . 3}$ | $\mathbf{8 9 . 0} \pm \mathbf{3 . 1}$ | $\mathbf{7 1 . 0} \pm \mathbf{4 . 5}$ |

Table A4. Multiagent Debate with Many Agents. Our approach also improves the performance with a very large number of agents.

| Model | Arithmetic | GSM8K | MMLU |
| :--- | :---: | :---: | :---: |
| Single Agent | $9.0 \pm 1.6$ | $20.7 \pm 2.3$ | $41.0 \pm 2.8$ |
| Single Agent (Reflection) | $10.7 \pm 1.7$ | $21.0 \pm 2.3$ | $39.7 \pm 2.8$ |
| Multiagent (Majority) | $11.0 \pm 1.8$ | $25.7 \pm 2.5$ | $43.3 \pm 2.9$ |
| Multiagent (Debate) | $\mathbf{1 3 . 3} \pm \mathbf{1 . 9}$ | $\mathbf{2 9 . 3} \pm \mathbf{2 . 6}$ | $\mathbf{4 7 . 7} \pm \mathbf{2 . 9}$ |

Table A5. Multiagent Debate on chat-Llama 7B. Our approach also improves the performance of the chat-Llama model.

| Model | Arithmetic | GSM8K | MMLU |
| :--- | :---: | :---: | :---: |
| Debate ( (Liang et al., 2023)) | $51.5 \pm 5.0$ | $61.0 \pm 4.8$ | $46.0 \pm 4.0$ |
| Debate (Ours) | $\mathbf{8 1 . 8} \pm \mathbf{2 . 3}$ | $\mathbf{8 4 . 0} \pm \mathbf{2 . 1}$ | $\mathbf{7 1 . 1} \pm \mathbf{4 . 6}$ |

Table A6. Debate Comparison. Our approach outperforms an alternative debate formulation from (Liang et al., 2023)

| Type | Arithmetic | MMLU | GSM8K |
| :--- | :---: | :---: | :---: |
| Final Result | $77.5 \pm 2.9$ | $63.0 \pm 4.8$ | $81.6 \pm 2.1$ |
| Final Result + Reasoning | $\mathbf{8 1 . 8} \pm \mathbf{2 . 3}$ | $\mathbf{7 1 . 1} \pm \mathbf{4 . 6}$ | $\mathbf{8 5 . 0} \pm \mathbf{3 . 5}$ |

Table A7. Effect of Reasoning Chain on Debate. Providing the reasoning of other agents during debate substantially outperforms only providing the final answer from other agents.

| Model | Arithmetic | GSM8K | MMLU |
| :--- | :---: | :---: | :---: |
| Single Agent | $67.0 \pm 4.7$ | $77.0 \pm 4.2$ | $63.9 \pm 4.8$ |
| Single Agent (Reflection) | $72.1 \pm 2.4$ | $75.0 \pm 4.3$ | $57.7 \pm 5.9$ |
| Multiagent (Majority) | $75.0 \pm 3.9$ | $81.0 \pm 3.9$ | $67.0 \pm 4.7$ |
| Multiagent (Reflection) | $76.0 \pm 4.3$ | $80.0 \pm 4.0$ | $65.0 \pm 4.7$ |
| Multiagent (Debate) | $\mathbf{8 1 . 8} \pm \mathbf{2 . 3}$ | $\mathbf{8 4 . 0} \pm \mathbf{2 . 1}$ | $\mathbf{7 1 . 1} \pm \mathbf{4 . 6}$ |

Table A8. Reflection Ensemble. Our approach also improves over an ensemble of self-reflection models.

| Method | Arithmetic | GSM | Chess Reasoning | Biography | MMLU | Chess Validity |
| :--- | :---: | :---: | :---: | :---: | :---: | :---: |
| Single Agent | $95.6 \pm 5.1$ | $111.5 \pm 5.5$ | $8.3 \pm 0.6$ | $220.5 \pm 3.6$ | $91.7 \pm 5.1$ | $39.0 \pm 1.1$ |
| Single Agent (Reflection) | $170.2 \pm 5.7$ | $155.2 \pm 10.2$ | $64.6 \pm 3.2$ | $297.2 \pm 11.8$ | $97.2 \pm 5.9$ | $92.8 \pm 1.6$ |
| Multiagent (Majority) | $564 \pm 10.7$ | $660.1 \pm 16.2$ | $49.2 \pm 3.0$ | $1295 \pm 7.3$ | $422.31 \pm 12.3$ | $331.8 \pm 2.9$ |
| Multiagent (Debate) | $548.1 \pm 9.4$ | $524.2 \pm 11.7$ | $199.5 \pm 5.3$ | $967.1 \pm 43.7$ | $527.7 \pm 17.1$ | $306.1 \pm 1.9$ |

Table A9. Generation Token Cost of Methods on Each Dataset. Average number of generated tokens (summed across all rounds of debate / convservation) when answering a query per method per dataset.

| Task | Type | Prompt |
| :---: | :---: | :--- |
| Arithmetic | Reflection | Can you verify that your answer is correct. Please reiterate your answer, making sure to <br> state your answer at the end of the response. |
| GSM8K | Reflection | Can you double check that your answer is correct. Please reiterate your answer, <br> with your final answer a single numerical number, in the form Voxed\{\{answer\}\}. |
| Chess Reasoning | Reflection | Given the board state and your previous move suggestion, can you confirm <br> your best next chess move? Give a single move suggestion for white following the form <br> 14. <XXX> and make sure the chess move is valid in the current board state. |
| Biographies | Reflection | Closely examine your biography and provide an updated bullet point biography. |
| MMLU | Reflection | Can you double check that your answer is correct. Put your <br> final answer in the form $(X)$ at the end of your response. |
| Chess Validity | Reflection | Can you double check that your answer is valid? State your final answer <br> in a newline with a 2 letter response following the regex [a-h][1-8]. |

Figure A3. Self Reflection Prompts. List of self-reflection prompts used in each task


#### Abstract

Biography: Darrell Long is a computer scientist and professor at the University of California, Santa Cruz. He received his Ph.D. in Computer Science from the University of Maryland in 1985 and has since made significant contributions in the field of storage systems and distributed systems. Long's research focuses on file systems, storage virtualization, and high-performance computing, and he has published over 150 papers in these areas. He has received numerous awards and recognitions for his work, including the prestigious ACM SIGMOD Test-of-Time Award, the IEEE Reynold B. Johnson Information Storage Systems Award, and the USENIX Lifetime Achievement Award. Long has served as Chief Scientist at Panasas, a computational storage company, and has been a consultant for several other tech companies. He is a Fellow of both the Association for Computing Machinery (ACM) and the Institute of Electrical and Electronics Engineers (IEEE) and a member of the National Academy of Engineering. Long has played leadership roles in several organizations, including serving as the Chair of the USENIX Association Board of Directors, Editor-in-Chief of the ACM Transactions on Storage, and on the editorial boards of several other major computer science journals.


Fact: Darrell Don Earl Long is a computer scientist and computer engineer.
ChatGPT Judgement: True.
Fact: Darrell Don Earl Long did his undergraduate studies at San Diego State University and received his Ph.D. from the University of California, San Diego.
ChatGPT Judgement: False.
Figure A5. Example Biography chatGPT Evaluations. To evaluate the accuracy of biographies, we prompt chatGPT with a generated biography and a ground truth fact about the person. We then ask chatGPT to verify if the generated biography is consistent with the ground truth fact.

| Task | Type | Prompt |
| :---: | :---: | :---: |
| Arithmetic | Starting | What is the result of $\}+\{ \} *\{ \}+\{ \}-\{ \} *\{ \}$ ? Make sure to state your answer at the end of the response. |
|  | Debate | These are the recent/updated opinions from other agents: <other agent responses> Use these opinions carefully as additional advice, can you provide an updated answer? Make sure to state your answer at the end of the response. |
| GSM8K | Starting | Can you solve the following math problem? <Problem> Explain your reasoning. Your final answer should be a single numerical number, in the form \boxed\{\{answer\}\}, at the end of your response. |
|  | Debate | These are the solutions to the problem from other agents: <other agent responses> Using the solutions from other agents as additional information, can you provide your answer to the math problem? The original math problem is <Problem>. Your final answer should be a single numerical number, in the form \boxed\{\{answer\}\}, at the end of your response. |
| Chess | Starting | Here is the current sequence of moves in a chess game: <moves>. What is the best chess move I should execute next? Give a single move suggestion of the form 14. $<X X X>$ and make sure the chess move is valid in the current board state. |
|  | Debate | Here are other chess move suggestions from other agents: <other agent responses> Using the chess suggestions from other agents as additional advice and your earlier generated solution, can you give me your updated thoughts on the best next chess move I should play given the chess sequence? Give a single move suggestion of the form 14. $\langle X X X>$ and make sure the chess move is valid in the current board state. |
| Biographies | Starting | Give a bullet point biography of highlighting their contributions and achievements as a computer scientist, with each fact separated with a new line character. |
|  | Debate | Here are some bullet point biographies of <person> given by other agents: <other agent response> Closely examine your biography and the biography of other agents and provide an updated bullet point biography. |
| MMLU | Starting | Can you answer the following question as accurately as possible? : A) , B), C), D) Explain your answer, putting the answer in the form $(X)$ at the end of your response. |
|  | Debate | These are the solutions to the problem from other agents: <other agent responses> Using the reasoning from other agents as additional advice, can you give an updated answer? Examine your solution and that other agents. Put your answer in the form $(X)$ at the end of your response. |
| Chess Validity | Starting | Given the chess game, give one valid destination square for the chess piece at. State the destination square in the form $(X)$, where $X$ follows the regex [a-h][1-8], for example (e5). Give a one line explanation of why your destination square is a valid move. |
|  | Debate | Here are destination square suggestions from other agents: <other agent responses> Can you double check that your destination square is a valid move? Check the valid move justifications from other agents. State your final answer in a newline with a 2 letter response following the regex [a-h][1-8]. |

Figure A4. Multiagent Debate Prompts. List of prompts used in each task.

User Input: Can you solve the following math problem? John arm wrestles 20 people. He beats $80 \%$. How many people did he lose to? Explain your reasoning. Your final answer should be a single numerical number, in the form \boxed \{answer\}, at the end of your response.

| Round 1 |  |
| :---: | :---: |
| Agent 1: If John beat $80 \%$ of the people he arm wrestled, he must have lost to the remaining $20 \%$. To find how many people that is, we can set up a proportion: <br> $\$ \backslash$ frac $\{\backslash$ text $\{$ number of people John beat $\}\}\{\backslash$ text $\{$ total number of people John arm wrestled $\}\}=80 \upharpoonleft \%=\backslash \operatorname{frac}\{4\}\{5\} \$$ | Agent 2: If John beats $80 \%$ of the people he arm wrestles, then he must lose to $20 \%$ of them. Since $20 \%$ of 20 people is $\$(0.20)(20)=4 \$$, John must have lost to $\$$ lboxed $\{4\} \$$ people. |
| Let x be the number of people John lost to. Then we know: |  |
| $\$ \backslash$ frac $\{\backslash$ text $\{$ number of people John beat $\}\}\{\backslash$ text $\{$ total number of people John arm wrestled $\}\}=\backslash$ frac $\{\backslash$ text $\{$ number of people John beat $\}\}\{\backslash$ text $\{$ number of people John beat $\}+\backslash \operatorname{text}\{x\}\}=\backslash \operatorname{frac}\{4\}\{5\} \$$ |  |
| Cross-multiplying yields: |  |
| ```$5\text{(number of people John beat)} = 4(\text {number of people John beat} +\text{x})$``` |  |
| Simplifying this equation gives: |  |
| \$ $\backslash$ text $\{\mathrm{x}\}=\backslash$ boxed $\{25\} \$$ |  |
| Therefore, John must have lost to 25 people. |  |

User Input: These are the solutions to the problem from other agents: One agent solution:
(Copy Agent 2's reply)
(Copy Agent 1's reply)
Using the solutions from other agents as additional information, can you provide your answer to the math problem? Please update and respond to other agents. Your final answer should be a single numerical number, in the form lboxed \{answer\}, at the end of your response.

## Round 2

Agent 1: Both solutions are correct in understanding that since John beat $80 \%$ of the people he arm wrestled, he lost to the remaining $20 \%$. However, the first solution made an error in the calculation, where they found that John won against 16 people out of $20(80 \%$ of 20$)$, leading them to think that he lost to 4.
Since John arm-wrestled 20 people in total, and won against $80 \%$ of them, he must have lost to the remaining $20 \%$. Therefore, the number of people John lost to is:
$\$ \backslash$ text $\{$ Number of people John lost to $\}=20 \backslash$ cdot $0.2=\backslash$ boxed $\{4\} \$$
So the correct answer is indeed 4, as the second agent already provided.

Agent 2: I agree with the other agents that John lost to $20 \%$ of the people he arm wrestled, which is equal to 4 people since he arm wrestled a total of 20 people. Therefore the answer is \$lboxed $\{4\} \$$.

User Input: These are the solutions to the problem from other agents: One agent solution:
(Copy Agent 2's reply)

Using the solutions from other agents as additional information, can you provide your answer to the math problem? Please update and respond to other agents. Your final answer should be a single numerical number, in the form lboxed \{answer\}, at the end of your response.

| Round 3 |  |
| :---: | :---: |
| Agent 1: I agree with both previous agents that since John beat $80 \%$ of the people he arm wrestled, he lost to the remaining $20 \%$. Therefore, the number of people John lost to is: <br> $\$ \backslash$ text $\{$ Number of people John lost to $\}=20 \backslash$ cdot $0.2=\backslash$ boxed $\{4\} \$$ So the answer is indeed 4 , as the first agent incorrectly calculated but the second agent correctly calculated. | Agent 2: I agree with the other agents that John lost to $20 \%$ of the people he arm wrestled, which is equal to 4 people since he arm wrestled a total of 20 people. Therefore the answer is \$lboxed $\{4\}$ \$. |

Figure A6. Example of a correct GSM8K Debate.

User Input: Can you solve the following math problem? Shondra has 7 fewer plants than Toni. Toni has $60 \%$ more plants than Frederick. If Frederick has 10 plants, how many plants does Shondra have? Explain your reasoning. Your final answer should be a single numerical number, in the form \boxed\{answer\}, at the end of your response.

| Round 1 |  |
| :---: | :---: |
| Agent 1: If Frederick has 10 plants, then Toni has $60 \%$ more plants, which is equal to $10+(60 / 100)^{*} 10=16$ more plants than Frederick. Therefore, Toni has a total of $10+16=26$ plants. <br> Shondra has 7 fewer plants than Toni, which means she has 26-7 $=19$ plants. | Agent 2: First, we need to determine how many plants Toni has. We know that Toni has $60 \%$ more plants than Frederick, who has 10 plants. To find $60 \%$ of 10 , we can multiply 10 by 0.6 , which gives us 6 . Therefore, Toni has $10+6=16$ plants. |
| Final answer: \boxed $\{19\}$. |  |
|  | Next, we know that Shondra has 7 fewer plants than Toni. So we can subtract 7 from Toni's 16 plants to get Shondra's number of plants: 16-7=9 plants. |
| X | Therefore, Shondra has $\backslash$ boxed $\{9\}$ plants. |

User Input: These are the solutions to the problem from other agents: One agent solution: (Copy Agent 2's reply)
(Copy Agent 1's reply)
Using the solutions from other agents as additional information, can you provide your answer to the math problem? Please update and respond to other agents. Your final answer should be a single numerical number, in the form lboxed \{answer\}, at the end of your response.

| Round 2 |  |
| :--- | :--- | :--- |
| Agent 1: I confirm that the solution provided by the other agent is correct. | Agent 2: I agree with the previous <br> agents that Toni has 16 plants and |
| Toni has 16 plants, and since Shondra has 7 fewer plants than Toni, she has |  |
| $16-7=9$ plants. | Shondra has 7 fewer plants than Toni, <br> which is 9 plants. Therefore, my final |
| answer is $\backslash$ boxed $\{9\}$. |  |

User Input: These are the solutions to the problem from other agents: One agent solution:

$$
\begin{aligned}
& \text { (Copy Agent 2's reply) } \\
& \text { (Copy Agent 1's reply) }
\end{aligned}
$$

Using the solutions from other agents as additional information, can you provide your answer to the math problem?
Please update and respond to other agents. Your final answer should be a single numerical number, in the form lboxed\{answer\}, at the end of your response.

| Round 3 |  |
| :---: | :---: |
| Agent 1: I agree with the previous agents that Toni has 16 plants and Shondra has 7 fewer plants than Toni, which means Shondra has 16-7=9 plants. | Agent 2: I concur with the other agents that Toni has 16 plants and Shondra has 7 fewer plants than Toni, which means Shondra has 16-7=9 plants. |
| Therefore, my final answer is \boxed $\{9\}$. | Therefore, my final answer is \boxed $\{9\}$. |
| $\checkmark$ | $\downarrow$ |

Figure A7. Example of Correct GSM8K Debate.

User Input: Can you solve the following math problem? Bob is in charge of doing laundry for a large hotel. Each room has two sheets, one comforter, twice as many pillow cases as sheets and twice as many towels as pillow cases. How many pieces of laundry are there in 80 rooms? Explain your reasoning. Your final answer should be a single numerical number, in the form \boxed \{answer\}, at the end of your response.

| Round 1 |  |
| :---: | :---: |
| Agent 1: For each room, Bob needs to wash: | Agent 2: In each room, there are: |
| - 2 sheets |  |
| - 1 comforter | - 2 sheets |
| $-2 \times 2=4$ pillow cases | - 1 comforter |
| -4 $\times 2=8$ towels | - 2 pillow cases (twice as many as sheets) |
|  | - 4 towels (twice as many as pillow cases) |
| So for each room, Bob needs to wash a total of $2+1+4+8=15$ pieces of laundry. | So in total, each room has $2+1+2+4=9$ pieces of laundry. |
| For 80 rooms, Bob needs to wash a total of $80 \times 15=1200$ pieces of laundry. | Therefore, in 80 rooms, there are |
| Therefore, the answer is $\backslash \operatorname{boxed}\{1200\}$. | $80 * 9=$ boxed $\{720\}$ piec |

User Input: These are the solutions to the problem from other agents: One agent solution:
(Copy Agent 2's reply) (Copy Agent 1's reply)
Using the solutions from other agents as additional information, can you provide your answer to the math problem? Please update and respond to other agents. Your final answer should be a single numerical number, in the form lboxed \{answer\}, at the end of your response.

| Round 2 |  |
| :--- | :--- |
| Agent 1: Based on the given information, each room has: <br> -2 sheets | Agent 2: Based on the information <br> provided, each room has: |
| -1 comforter |  |
| $-2 \times 2=4$ pillow cases | -2 sheets |
| $-4 \times 2=8$ towels | -1 comforter |
| This gives a total of $2+1+4+8=15$ pieces of laundry per room. | $-2 \times 2=4$ pillow cases |
| Therefore, in 80 rooms, there are $80 \times 15=$ boxed $\{1200\}$ pieces of | This gives a total of $2+1+4+8=15$ pieces of <br> laundry per room. |
| laundry in total. | Therefore, in 80 rooms, there are $80 \times 15=$ <br> Note: While the other agent's solution is correct in terms of the number <br> of pieces of laundry per room, they made a mistake in their final <br> calculation by multiplying 80 by 9 instead of 15. |

User Input: These are the solutions to the problem from other agents: One agent solution:
(Copy Agent 2's reply)

Using the solutions from other agents as additional information, can you provide your answer to the math problem? Please update and respond to other agents. Your final answer should be a single numerical number, in the form lboxed \{answer\}, at the end of your response.

| Round 3 |  |
| :---: | :---: |
| Agent 1: Based on the information provided, each room has: <br> - 2 sheets <br> - 1 comforter <br> $-2 \times 2=4$ pillow cases <br> $-4 \times 2=8$ towels | Agent 2: Based on the information provided, each room has: <br> - 2 sheets <br> - 1 comforter <br> $-2 \times 2=4$ pillow cases <br> $-4 \times 2=8$ towels |
| This gives a total of $2+1+4+8=15$ pieces of laundry per room. <br> Therefore, in 80 rooms, there are $80 \times 15=\backslash$ boxed $\{1200\}$ pieces of laundry in total. | This gives a total of $2+1+4+8=15$ pieces of laundry per room. <br> Therefore, in 80 rooms, there are $80 \times 15=$ |
| I agree with the other agent's corrected calculation. | lboxed $\{1200\}$ pieces of laundry in total. <br> I agree with the other agents' answers and calculations. |

Figure A8. Example of Correct GSM8K Debate.

User Input: Can you solve the following math problem? Lori wants to buy a $\$ 320.00$ pair of shoes and a matching belt that is $\$ 32.00$. Her part-time job pays her $\$ 8.00$ an hour. How many hours will she have to work before she can make her purchase? Explain your reasoning. Your final answer should be a single numerical number, in the form \boxed \{answer\}, at the end of your response.
Round 1

Agent 1: To find out how many hours Lori will need to work to make her purchase, we need to first calculate the total cost of both the shoes and belt:

$$
\$ 320.00 \text { (shoes) }+\$ 32.00(\text { belt })=\$ 352.00
$$

Next, we need to figure out how much money Lori will earn from her part-time job. We can set up a proportion:
$\$ 8.00 /$ hour $=\mathrm{X} /$ total hours
We can solve for X by cross-multiplying:
\$8.00/hour $=$ X/total hours
$\$ 8.00 *$ total hours $=\mathrm{X}$
$\mathrm{X}=\$ 352.00$
This means that Lori will need to work a total of:
$\$ 352.00 \div \$ 8.00 /$ hour $=44$ hours
Therefore, Lori will need to work 44 hours before she can make her purchase. Answer: \boxed $\{44\}$.

Agent 2: Lori needs to earn a total of $\$ 352.00$ ( $\$ 320.00$ for shoes $+\$ 32.00$ for belt) to make her purchase. She earns $\$ 8.00$ per hour at her job.

To find out how many hours she needs to work, we can set up an equation:
$\$ 8.00 /$ hour $\backslash$ times $\times$ hours $=\$ 352.00$
Solving for x , we can divide both sides by \$8.00/hour:
$x$ hours $=\$ 352.00 / \$ 8.00 /$ hour
$x$ hours $=44$ hours
Therefore, Lori needs to work 44 hours before she can make her purchase.

$$
\backslash \text { boxed }\{44\}
$$

User Input: These are the solutions to the problem from other agents: One agent solution:
(Copy Agent 2's reply)
(Copy Agent 1's reply)
Using the solutions from other agents as additional information, can you provide your answer to the math problem? Please update and respond to other agents. Your final answer should be a single numerical number, in the form lboxed \{answer\}, at the end of your response.

| Round 2 |
| :--- |
| Agent 1: I agree with the other agents that Lori needs to work a total of |
| $\$ 352.00 \div \$ 8.00 /$ hour $=44$ hours to make her purchase. The final |
| answer is: $\backslash$ boxed $\{44\}$. |

Agent 2: I agree with the other agents that Lori needs to work for 44 hours to make her purchase. The final answer is $\operatorname{lboxed}\{44\}$

User Input: These are the solutions to the problem from other agents: One agent solution:
(Copy Agent 2's reply) (Copy Agent 1's reply)

Using the solutions from other agents as additional information, can you provide your answer to the math problem? Please update and respond to other agents. Your final answer should be a single numerical number, in the form lboxed \{answer\}, at the end of your response.

## Round 3

Agent 1: I agree with the other agents that Lori needs to work 44 hours to make her purchase. Therefore, the final answer is $\backslash$ boxed $\{44\}$.

Agent 2: After reviewing the solutions provided by other agents, I agree that Lori needs to work 44 hours to make her purchase, so the final answer is $\mathrm{lboxed}\{44\}$.

Figure A9. Example of Correct GSM8K Debate.

User Input: Can you solve the following math problem? Jennifer bought 12 oranges from the market, she gave her three daughters 2 oranges each, and her only boy got 3 oranges. How many oranges did she remain with? Explain your reasoning. Your final answer should be a single numerical number, in the form \boxed \{answer\}, at the end of your response.

## Round 1

Agent 1: Jennifer gave a total of 2 oranges/daughter x 3 daughters +3 oranges/son $=9$ oranges to her children. Thus, she remained with 12 oranges -9 oranges $=3$ oranges.
lboxed $\{3\}$.

Agent 2: Jennifer gave away a total of \$2 \cdot $3+3=9 \$$ oranges. Therefore, she remained with $\$ 12-9=\backslash \operatorname{boxed}\{3\} \$$ oranges.

User Input: These are the solutions to the problem from other agents: One agent solution:
(Copy Agent 2's reply) (Copy Agent 1's reply)

Using the solutions from other agents as additional information, can you provide your answer to the math problem? Please update and respond to other agents. Your final answer should be a single numerical number, in the form lboxed \{answer\}, at the end of your response.

## Round 2

Agent 1: I agree with the previous solutions provided. Jennifer gave a total of 2 oranges/daughter x 3 daughters +3 oranges/son $=9$ oranges to her children. Thus, she remained with 12 oranges -9 oranges $=3$ oranges.

```
lboxed {3}.
```

Agent 2: I agree with the other agents' solutions. Jennifer gave away a total of 2 oranges/daughter x 3 daughters +3 oranges/son $=9$ oranges. Therefore, she remains with 12 oranges -9 oranges $=$ $\$$ bboxed $\{3\} \$$ oranges.

User Input: These are the solutions to the problem from other agents: One agent solution:

> (Copy Agent 2's reply)
(Copy Agent 1's reply)
Using the solutions from other agents as additional information, can you provide your answer to the math problem? Please update and respond to other agents. Your final answer should be a single numerical number, in the form lboxed\{answer\}, at the end of your response.

## Round 3

Agent 1: I confirm the previous solutions provided. Jennifer gave a total of 2 oranges/daughter $\times 3$ daughters +3 oranges/son $=9$ oranges to her children. Therefore, she remains with 12 oranges -9 oranges $=$ $\$$ lboxed $\{3\} \$$ oranges.

Agent 2: I agree with the other agents. Jennifer gave away a total of 2 oranges/daughter x 3 daughters +3 oranges/son $=9$ oranges to her children, therefore she remains with 12 oranges - 9 oranges $=\$ \backslash$ boxed $\{3\} \$$ oranges.

User Input: Can you solve the following math problem? A toy manufacturer receives an order for 400 toys. 5 workers are available to work on the order. 2 of the workers produce 6 toys an hour, and another 2 workers produce 4 toys an hour. They all work on the order during their 10 -hour shift, and by the end of their shift the manufacturer still needs another 20 toys to be able to ship the order. How many toys per hour does the fifth worker produce? Explain your reasoning. Your final answer should be a single numerical number, in the form \boxed \{answer\}, at the end of your response.
Agent 1: Let's first find the total number of toys that the four workers
produce in 10 hours.
The two workers who produce 6 toys an hour make a total of $6 \times 2 \times 10=$
120 toys.
The other two workers who produce 4 toys an hour make a total of $4 \times 2 \times 10$
$=80$ toys.
Together, they produce $120+80=200$ toys.

Since the manufacturer still needs another 20 toys to complete the order, we know that all 5 workers together need to produce a total of $400+20=420$ toys.

Subtracting the 200 toys that have already been produced, we're left with 420 $-200=220$ toys that need to be produced by the fifth worker.

Since the workers have worked for 10 hours, we can set up the equation:
Toys per hour for fifth worker $\times 10=220$
Solving for the toys per hour for the fifth worker, we get:
Toys per hour for fifth worker $=220 / 10=22$
Therefore, the fifth worker must produce $\backslash$ boxed $\{22\}$ toys per hour.

Agent 2: Let the number of hours each group of workers (6-toy producers and 4-toy producers) work be $x$ and $y$, respectively. Then we have the following system of equations:
$\begin{array}{ll}2(6 x)+2(4 y)=400+20 \\ 12 x+8 y=210 & (\text { Equation 1) } \\ x+y=10 & \text { (Equation 2) }\end{array}$
Simplifying Equation 2, we get:
$y=10-x$
Substituting this into Equation 1, we get:
$12 \mathrm{x}+8(10-\mathrm{x})=210$
$4 \mathrm{x}=50$
$\mathrm{x}=12.5$
Thus, the first 4 workers produce a total of $2(6)(12.5)+2(4)(10-$ $12.5)=150$ toys. Since they still need another 20 toys to meet the order, the fifth worker must produce $20 / 10=2$ toys per hour.

Therefore, the fifth worker produces $\backslash \operatorname{boxed}\{2\}$ toys per hour.

User Input: These are the solutions to the problem from other agents: One agent solution:
(Copy Agent 2's reply)

Using the solutions from other agents as additional information, can you provide your answer to the math problem?
Please update and respond to other agents. Your final answer should be a single numerical number, in the form lboxed \{answer\}, at the end of your response.
Round 2
produced by the four workers in 10 hours is incorrect. The correct calculation should be:

2 (6 toys per hour) $\times 10$ hours $\times 2$ workers $=240$ toys
$2(4$ toys per hour) $\times 10$ hours $\times 2$ workers $=160$ toys
Total toys produced by the four workers $=240+160=400$ toys
So, the manufacturer still needs to produce another 400-20=380 toys.
To produce the remaining 380 toys in 10 hours, all five workers need to produce $380 / 10=38$ toys per hour.

Since four workers are already producing 6 and 4 toys per hour, we can set up the equation:
$2(6)+2(4)+x=38$
Simplifying this, we get: $x=7$
Therefore the fifth worker must produce $\backslash$ boxed $\{7\}$ toys per hour.

Agent 2: I agree with the updated solution provided by the other agent. The total toys produced by the four workers in 10 hours is 400 , and since the manufacturer still needs to produce another 380 toys, all five workers need to produce 38 toys per hour.

Setting up the equation, with $x$ being the number of toys produced by the fifth worker per hour, we get:
$2(6)+2(4)+x=38$
$12+8+\mathrm{x}=38$
$\mathrm{x}=18$
Therefore, the fifth worker must produce $\backslash$ boxed $\{18\}$ toys per hour.

User Input: These are the solutions to the problem from other agents: One agent solution:
(Copy Agent 2's reply)

Using the solutions from other agents as additional information, can you provide your answer to the math problem?
Please update and respond to other agents. Your final answer should be a single numerical number, in the form \boxed \{answer\}, at the end of your response.

| Round 3 |  |
| :---: | :---: |
| Agent 1: I believe there is an error in the updated solution provided by the other agents. The total number of toys produced by the four workers in 10 hours is indeed 400, and since the manufacturer still needs to produce another 380 toys, all five workers need to produce 38 toys per hour. | Agent 2: I agree with the latest solution provided by the other agent. The total number of toys produced by the four workers is 400 , and the manufacturer still needs to produce another 380 toys. To produce the remaining 380 toys in 10 hours, all five workers need to produce 38 toys per hour. |
| Setting up the equation, with $x$ being the number of toys produced by the fifth worker per hour, we get: | Setting up the equation, with $x$ being the number of toys produced by the fifth worker per hour, we get: |
| $2(6)+2(4)+\mathrm{x}=38$ $12+8+\mathrm{x}=38$ | $2(6)+2(4)+x=38$ |
| $\begin{aligned} & 12+8+x=38 \\ & x=18 / 5 \end{aligned}$ | $12+8+\mathrm{x}=38$ |
|  | $\mathrm{x}=18 / 5$ |
| Therefore, the fifth worker must produce $\backslash$ boxed $\{3.6\}$ toys per hour. | Therefore, the fifth worker must produce $\backslash$ boxed $\{3.6\}$ toys per hour. |

Figure A11. Example of Incorrect GSM8K Debate.

User Input: Can you solve the following math problem? Courtney attended a concert and reported that the audience was 48 in number. However, Kelly went to the same concert and said that Courtney had made the mistake of overstating the number of people in attendance by $20 \%$. If Kelly was right, how many people really attended the concert? Explain your reasoning. Your final answer should be a single numerical number, in the form \boxed \{answer\}, at the end of your response.

## Round 1

Agent 1: If Courtney overstated the number of people by $20 \%$, that means she reported $120 \%$ of the actual number of people. Let's call the actual number of people $x$. Then we can set up the following equation:
$1.2 x=48$
Solving for x , we get:
$x=40$
Therefore, the actual number of people who attended the concert was $\backslash$ boxed $\{40\}$.

Agent 2: Let's call the actual number of people at the concert "x." Then, according to Kelly's statement, we know that:
$\mathrm{x}=48-0.20(48)$
$\mathrm{x}=48-9.6$
$\mathrm{x}=38.4$
Therefore, the actual number of people at the concert was $\$$ lboxed $\{38\} \$$.

User Input: These are the solutions to the problem from other agents: One agent solution: (Copy Agent 2's reply)

Copy Agent 1's reply)
Using the solutions from other agents as additional information, can you provide your answer to the math problem? Please update and respond to other agents. Your final answer should be a single numerical number, in the form boxed \{answer\}, at the end of your response.

| Round 2 |  |
| :--- | :--- |
| Agent 1: The other agent's solution assumes that Courtney <br> overstated the number of people by 20 percentage points, not 20 <br> percent of the original number. This means their answer is <br> incorrect. | Agent 2: I agree with the first agent's solution that <br> we need to consider the 20\% decrease from the <br> original number of people. Let x be the actual <br> number of people at the concert, then we have: |
| Instead, we can use the method in the first solution. Let x be the <br> actual number of people at the concert. Then, we know that: | $\mathrm{x}=48-0.20(48)$ <br> $\mathrm{x}=48-9.6$ <br> $\mathrm{x}=38.4$ |
| $\mathrm{x}=48 *(100 \%-20 \%)=48 * 0.8=38.4$ | Rounding to the nearest whole number, we get <br> that the actual number of people at the concert <br> was $\$$ boxed $\{38\} . \$$ |
| Therefore, the actual number of people at the concert was <br> $\$ 1$ boxed $\{38\} \$$. | I also agree with the second agent's point that we <br> should be careful in interpreting the 20\%, as it <br> could be presented in percentage points or as a <br> percent of the original number. |

User Input: These are the solutions to the problem from other agents: One agent solution:
(Copy Agent 2's reply)
(Copy Agent 1's reply)
Using the solutions from other agents as additional information, can you provide your answer to the math problem? Please update and respond to other agents. Your final answer should be a single numerical number, in the form lboxed \{answer\}, at the end of your response.

| Round 3 |
| :--- | :--- |\(\left.\quad \begin{array}{l}Agent 2: I agree with the previous agents that the <br>

Agent 1: I agree with the previous agents that we need to <br>
consider the 20 \% decrease from the original number of people. <br>
Let \mathrm{x} be the actual number of people at the concert, then we have:\end{array} \quad $$
\begin{array}{l}\text { 20\% decrease should be calculated as a percentage } \\
\text { of the original number of people. Let } \mathrm{x} \text { be the } \\
\text { actual number of people at the concert, then we } \\
\text { have: }\end{array}
$$\right]\)

Figure A12. Example of Incorrect GSM8K Debate.

User Input: Can you solve the following math problem? Peter has twice as many socks as Jack and half times as many dishes as jack. Jack collected twice as many dishes as socks in the store. If jack collected 60 dishes, calculate the total number of socks and dishes they have together? Explain your reasoning. Your final answer should be a single numerical number, in the form \boxed \{answer\}, at the end of your response.

| Round 1 |  |
| :---: | :---: |
| Agent 1: If Jack collected twice as many dishes as socks, then he collected $60 / 2=\ll 60 / 2=30 \gg 30$ socks. | Agent 2: If Jack has collected twice as many dishes as socks and he collected 60 dishes, then he has $60 / 2=30$ socks. |
| Peter has half times as many dishes as Jack, meaning he has | Peter has twice as many socks as Jack, which means he has $2 * 30=60$ socks. |
| $60 / 2=\ll 60 / 2=30 \gg 30$ dishes. | Peter also has half as many dishes as Jack, which |
| If Peter has twice as many socks as Jack, then he has | means he has 60/2 $=30$ |
| $2 * 30=\ll 2 * 30=60 \gg 60$ socks. | dishes. |
| Together, Peter and Jack have |  |
| $60+30+30=\ll 60+30+30=120 \gg 120$ socks and dishes. Answer: |  |
| lboxed $\{120\}$. | The total number of dishes and socks they have together is $60+30+30=\backslash$ boxed $\{120\}$. |

User Input: These are the solutions to the problem from other agents: One agent solution: (Copy Agent 2's reply)
(Copy Agent 1's reply)
Using the solutions from other agents as additional information, can you provide your answer to the math problem? Please update and respond to other agents. Your final answer should be a single numerical number, in the form lboxed \{answer\}, at the end of your response.

| Round 2 |  |
| :---: | :---: |
| Agent 1: I agree with the solutions provided by the other agents. <br> If Jack collected 60 dishes and he collected twice as many dishes as socks, then he collected $60 / 2=30$ socks. | Agent 2: I agree with the solutions provided by the other agents. The total number of socks and dishes they have together is 120 . Therefore, my answer is \boxed $\{120\}$. |
| If Peter has twice as many socks as Jack, then he has $2 * 30=60$ socks. |  |
| If Peter has half as many dishes as Jack, then he has $60 / 2=30$ dishes. |  |
| Therefore, the total number of dishes and socks they have together is $60+30+30=\backslash$ boxed $\{120\}$. |  |

User Input: These are the solutions to the problem from other agents: One agent solution:

> (Copy Agent 2's reply) (Copy Agent 1's reply)

Using the solutions from other agents as additional information, can you provide your answer to the math problem? Please update and respond to other agents. Your final answer should be a single numerical number, in the form lboxed \{answer\}, at the end of your response.

## Round 3

Agent 1: I agree with the solutions provided by the other agents.

If Jack collected 60 dishes and he collected twice as many dishes as socks, then he collected $60 / 2=30$ socks.

If Peter has twice as many socks as Jack, then he has $2 * 30=60$ socks.

If Peter has half as many dishes as Jack, then he has $60 / 2=30$ dishes.

Therefore, the total number of dishes and socks they have together is $60+30+30=\langle\operatorname{boxed}\{120\}$.

Agent 2: I agree with the solutions provided by the other agents. Therefore, the total number of socks and dishes they have together is 120 . My final answer is \boxed $\{120\}$.

Figure A13. Example of Incorrect GSM8K Debate.

Instruction: What is the result of $4+23^{*} 6+24-24^{*} 12$ ? Make sure to state your answer at the end of the response.

## Round 1

Agent 1: The result of $4+23^{*} 6+24-24^{*} 12$ is -244 .

Instruction: These are the recent/updated opinions from other agents: (Copy Agent 2's reply) Use these opinions carefully as additional advice, can you provide an updated answer? Make sure to state your answer at the end of the response.

Agent 2: Using the order of operations (PEMDAS), we will first multiply 23 by 6 to get 138 . Then we will multiply 24 by 12 to get 288 . Finally, we will add up all the terms: $4+138+24$ -$288=-122$.

Instruction: These are the recent/updated opinions from other agents: (Copy Agent 1's reply) Use these opinions carefully as additional advice, can you provide an updated answer? Make sure to state your answer at the end of the response.

## Round 2

Agent 1: Based on the order of operations (PEMDAS), the correct way to solve the expression $4+23 * 6+24-24 * 12$ is:
$4+138+24-288=-122$
Therefore, the updated answer is -122 .

Instruction: These are the recent/updated opinions from other agents: (Copy Agent 2's reply) Use these opinions carefully as additional advice, can you provide an updated answer? Make sure to state your answer at the end of the response.

Agent 2: One common way to accurately solve mathematical expressions is by using the order of operations (PEMDAS) which means Parentheses, Exponents, Multiplication and Division, and finally Addition and Subtraction from left to right. To apply this rule in solving the given mathematical expression, we will first multiply 23 by 6 then multiply 24 by 12 before performing any addition or subtraction. Using PEMDAS,
$4+23 * 6+24-24 * 12=4+138+24-288=-$ 122.

Therefore, based on the order of operations, the result of the expression is -122 , contrary to the opinion provided by another agent.

Instruction: These are the recent/updated opinions from other agents: (Copy Agent 1's reply) Use these opinions carefully as additional advice, can you provide an updated answer? Make sure to state your answer at the end of the response.

## Round 3

Agent 1: I agree with the analysis and solution provided by the other agent. Using the order of operations (PEMDAS), we perform multiplication before addition or subtraction. Therefore the expression is calculated as follows:

$$
4+23 * 6+24-24 * 12=4+138+24-288=-122
$$

So, the updated answer is -122 .

Agent 2: Yes, I agree with the other agents that based on the order of operations (PEMDAS), the correct way to solve the expression $4+23 * 6+24$ $-24 * 12$ is:
$4+138+24-288=-122$
Therefore, the updated answer is -122 .

Figure A14. Example of Arithmetic Debate.

Instruction: What is the result of $4+23 * 6+24-24 * 12$ ? Make sure to state your answer at the end of the response.

## Round 1

Agent 1: The result of $4+23 * 6+24-24 * 12$ is -244 .

Instruction: These are the recent/updated opinions from other agents: (Copy Agent 2's reply) Use these opinions carefully as additional advice, can you provide an updated answer? Make sure to state your answer at the end of the response.

Agent 2: Using the order of operations (PEMDAS), we will first multiply 23 by 6 to get 138. Then we will multiply 24 by 12 to get 288 . Finally, we will add up all the terms: $4+138+24$ -$288=-122$.

Instruction: These are the recent/updated opinions from other agents: (Copy Agent 1's reply) Use these opinions carefully as additional advice, can you provide an updated answer? Make sure to state your answer at the end of the response.

## Round 2

Agent 1: Based on the order of operations (PEMDAS), the correct way to solve the expression $4+23 * 6+24-24 * 12$ is:
$4+138+24-288=-122$
Therefore, the updated answer is -122 .

Instruction: These are the recent/updated opinions from other agents: (Copy Agent 2's reply) Use these opinions carefully as additional advice, can you provide an updated answer? Make sure to state your answer at the end of the response.

Agent 2: One common way to accurately solve mathematical expressions is by using the order of operations (PEMDAS) which means Parentheses, Exponents, Multiplication and Division, and finally Addition and Subtraction from left to right. To apply this rule in solving the given mathematical expression, we will first multiply 23 by 6 then multiply 24 by 12 before performing any addition or subtraction. Using PEMDAS,
$4+23 * 6+24-24 * 12=4+138+24-288=-$ 122.

Therefore, based on the order of operations, the result of the expression is -122 , contrary to the opinion provided by another agent.

Instruction: These are the recent/updated opinions from other agents: (Copy Agent 1's reply) Use these opinions carefully as additional advice, can you provide an updated answer? Make sure to state your answer at the end of the response.

## Round 3

Agent 1: I agree with the analysis and solution provided by the other agent. Using the order of operations (PEMDAS), we perform multiplication before addition or subtraction. Therefore the expression is calculated as follows:
$4+23 * 6+24-24 * 12=4+138+24-288=-122$
So, the updated answer is -122 .

Agent 2: Yes, I agree with the other agents that based on the order of operations (PEMDAS), the correct way to solve the expression $4+23 * 6+24$ $-24^{*} 12$ is:
$4+138+24-288=-122$
Therefore, the updated answer is -122 .

Figure A15. Example of Arithmetic Debate with Summarization. Four separate agents participate in debate, with two illustrated above. Instruction contains the summarized responses across agents.

Instruction: Can you answer the following question as accurately as possible? You suspect that your patient has an enlarged submandibular salivary gland. You expect the enlarged gland: A) to be palpable intraorally. B) to be palpable extraorally. C) to be palpable both intra- and extraorally. D) only to be detectable by radiographical examination. $\backslash n$ Explain your answer, putting the answer in the form (X) at the end of your response.
Agent 1: An enlarged submandibular salivary
gland is typically palpable extraorally. This is
because the submandibular gland is located
underneath the mandible bone, and therefore is
more easily palpable from the outside of the
mouth rather than inside. However, in some cases,
the gland may also be palpable intraorally if it is
significantly enlarged. Therefore, the answer is
(C) to be palpable both intra- and extraorally.

Instruction: Here are responses from other agents: (Copy Agent 2's reply) Using the solutions from other agents as additional information, can you give an updated response....

Agent 2: Based on the location of the submandibular salivary gland, it is more likely to be palpable extraorally than intraorally. This is because the gland is located beneath the mandible and extends posteriorly, making it more accessible for palpation outside of the oral cavity. Therefore, the expected answer is (B) to be palpable extraorally

Instruction: Here are responses from other agents: (Copy Agent l's reply) Using the solutions from other agents as additional information, can you give an updated response....

## Round 2

Agent 1: After considering the reasoning from other agents, I would update my answer to (C) to be palpable both intra- and extraorally. The submandibular salivary gland is located in the submandibular triangle, extending from the hyoid bone to the mandible. This location allows the gland to be palpable both intra- and extraorally. However, it is primarily palpable extraorally due to its location beneath the mandible, making it more accessible for examination outside of the oral cavity. Additionally, in some cases, if the gland is significantly enlarged, it may also be palpable intraorally. $\mathrm{n} \backslash n \mathrm{It}$ is essential to perform a thorough examination by palpating both intra- and extraorally to assess the gland properly. By combining the inputs from other agents, I conclude that it is crucial to perform a complete examination, including extraoral and intraoral palpation, to evaluate the gland's size and location. Therefore, the updated answer is (C) to be palpable both intra- and extraorally.

Agent 2: After carefully considering the reasoning from other agents, my updated answer would be (C) to be palpable both intra- and extraorally. The submandibular salivary gland is located in the submandibular triangle, extending from the hyoid bone to the mandible. While the gland is primarily palpable extraorally due to its location beneath the mandible, it can also be palpable intraorally in some cases if significantly enlarged. $\operatorname{nn} \backslash n T h e r e f o r e, ~ i t ~ i s ~$ essential to perform a thorough examination by palpating both intra- and extraorally to assess the gland properly, as stated by the other agents. This will allow for a more accurate and comprehensive evaluation of the gland's size and location. Overall, the input from other agents supports the conclusion that the gland can be palpable both intra- and extraorally, and a complete examination is necessary to evaluate its condition.

Figure A16. Example of MMLU Debate.

Instruction: Give a bullet point biography of Stephan R. Bourne highlighting their contributions and achievements as a computer scientist, with each fact separated by a new line character.

## Round 1

Agent 1:

- Stephen R. Bourne is a British-born American computer scientist.
- He is best known for creating the Bourne shell (sh) for the Unix operating system in 1979. \n- The Bourne shell was the default Unix shell for many years and is still in use today.
- Bourne also worked on the development of the UNIX System III and IV operating systems.
- He later joined AT\&T Bell Laboratories, where he worked on the UNIX System V operating system and its derivatives.
- Bourne was awarded the ACM Software System Award in 1983 for his work on the Bourne shell.
- He is also a recipient of the USENIX Lifetime Achievement Award (2001) and the IEEE Computer Pioneer Award (2010).
- Bourne is now retired and continues to be involved in the computer science community.

Instruction: Here are some bullet point biographies from other agents: (Copy Agent 2's reply) Using the biographies as additional information, can you give an updated bullet point biography.

Agent 2:

- Stephen R. Bourne is a computer scientist known for his development of the Unix shell, or command-line interface, known as the Bourne shell.
- He earned a bachelor's degree in mathematics and a doctorate in computer science from Cambridge University.
- Bourne worked at Bell Labs, where he helped to create the Unix operating system and its commandline interface.
- In addition to the Bourne shell, he also created the Bourne-Again shell (bash) and was involved in the creation of the Korn shell.
- Bourne was awarded the USENIX Lifetime Achievement Award in 1999 and was elected as a Fellow of the Association for Computing Machinery in 2013.
- He has authored and co-authored numerous technical papers and books, including "The Unix System," a widely used textbook on Unix programming.

Instruction: Here are some bullet point biographies from other agents: (Copy Agent 1's reply) Using the biographies as additional information, can you give an updated bullet point biography.

## Round 2

Agent 1:
Stephen R. Bourne is a computer scientist known for his development of the Unix shell, or command-line interface, known as the Bourne shell.

- He earned a bachelor's degree in mathematics and a doctorate in computer science from Cambridge University.
- Bourne worked at Bell Labs, where he played a key role in the creation of the Unix operating system, and was involved in the development of its command-line interfaces.
- In addition to the Bourne shell, he also created the Bourne-Again shell (bash) and contributed to the development of the Korn shell.
- Bourne was awarded the ACM Software System Award in 1983 for his work on the Bourne shell, and the USENIX Lifetime Achievement Award in 1999 for his contributions to the field of computer science.
- He was also elected as a Fellow of the Association for Computing Machinery in 2013 in recognition of his significant contributions to the field.
- Bourne has authored and co-authored numerous technical papers and books, including "The Unix System," a widely used textbook on Unix programming, and continues to be involved in the computer science community after his retirement.

Agent 2:

- Stephen R. Bourne is a highly respected computer scientist and mathematician who revolutionized the computing field with his development of the Unix shell, popularly known as the Bourne shell.
- He earned his bachelor's degree in mathematics and a doctorate in computer science from Cambridge University, and began his career at Bell Labs where he played an integral role in the creation of the Unix operating system.
During his tenure, Bourne developed the Bourne shell, which became the foundation for the Unix commandline interface.
- Bourne was not only a prolific developer, but he also contributed to the development of several other command-line interfaces, including the Bourne-Again shell (bash) and the Korn shell.
- As a testament to his significant contributions to the field of computer science, Bourne was honored with the ACM Software System Award in 1983, and later the USENIX Lifetime Achievement Award in 1999.
- In recognition of his substantial contributions, he was elected as a Fellow of the Association for Computing Machinery in 2013.
- Bourne is also an accomplished author and co-author of several technical papers and books, including "The Unix System," a widely used textbook on Unix programming.
Despite his retirement, Bourne continues to be actively involved in the computer science community, sharing his expertise and experiences with younger generations of aspiring computer scientists.

Figure A17. Example of Biography Debate. While we found that generated biographies after debate to be more accurate, many facts remain incorrect.

Instruction: Here is the current sequence of moves in a chess game: 1. e4 c6 2. d4 d5 3. Nd2 Qc7 4. Bd3 dxe4 5. Nxe4 Bf5 6. Nf3 Nd7 7. O-O e6 8. c4 Bg6 9. d5 Bxe4 10. Bxe4 Ndf6 11. Re1 Nxe4 12. Rxe4 Nf6 13. Re1 O-O-O . What is the best chess move I should execute next? Give a single move suggestion of the form $14 .<\mathrm{XXX}>$ and make sure the chess move is valid in the current board state.

## Round 1

Agent 1: 14. dxe6
Instruction: Using the solutions from other agents as additional information. (Copy Agent 2's reply)

Agent 2: 14. Be3
Instruction: Using the solutions from other agents as additional information. (Copy Agent 1's reply)

## Round 2

Agent 1: Based on the current board state and the suggested moves, I would suggest the move: 14. Qa4. This move attacks the undefended bishop on f5 and puts pressure on black's position.

Instruction: Using the solutions from other agents as additional information. (Copy Agent 2's reply)

Agent 2: I would still suggest the move 14. Be3. It develops a piece and prepares to put pressure on black's position. Additionally, it reinforces the e 4 pawn and prepares to castle kingside.

Instruction: Using the solutions from other agents as additional information. (Copy Agent 1's reply)

## Round 3

Agent 1: After considering the suggestions of other agents, I still believe that the move 14. Qa4 is the best move in the current board state. It attacks the bishop on $\mathrm{f5}$, prepares to develop the queen, and puts pressure on Black's position.

Agent 2: After considering the suggestions of other agents, I would suggest the move 14. Qa4. It attacks the bishop on $f 5$, prepares to develop the queen and puts pressure on Black's position.

Figure A18. Example of Chess Debate.

Instruction: What is the result of $23+19 * 22+22-7 * 9$ ? Make sure to state your answer at the end of the response.

## Round 1

Agent 1: First, we multiply 19 by 22 to get 418.
Then, we multiply 7 by 9 to get 63 .
We can use the order of operations, also known as PEMDAS, to simplify the expression:
$23+418+22-63$
$=460$
Therefore, the result of $23+19 * 22+22-7 * 9$ is 460 .

Instruction: Using the solutions from other agents as additional information. (Copy Agent 2's reply)

Agent 2: Using the order of operations (PEMDAS), we first multiply 19 and 22 , resulting in 418 . Then, we multiply 7 and 9 , resulting in 63 . Next, we subtract 63 from 22, which gives us -41 . Then, we add together 23, 418, -41 , and 22, which results in 422.

Therefore, the answer is 422 .

Instruction: Using the solutions from other agents as additional information. (Copy Agent 1's reply)

## Round 2

Agent 1: It seems that the first agent made an error by subtracting 63 from 22 instead of subtracting 63 from the previous result of the multiplication of 19 and 22.

Using the correct order of operations, the expression should be simplified as follows:
$19 * 22=418$
$7 * 9=63$
$23+418+22-63=400$

Therefore, the updated answer is 400 .

Agent 2: The order of operations (PEMDAS) must be followed to correctly simplify the expression.

First, we multiply 19 and 22 to get 418.
Next, we multiply 7 and 9 to get 63 .

Then, we add 23 to 418 to get 441 .
After that, we add 22 to get 463.

Finally, we subtract the product of 7 and 9 (63) from 463 , resulting in 400.

Therefore, the updated answer for $23+19 * 22+22-$ $7 * 9$ is 400 .

Figure A19. Example of Arithmetic Debate.

Instruction: What is the result of $29+28 * 15+8-2 * 27$ ? Make sure to state your answer at the end of the response.

## Round 1

Agent 1: The answer is 441.
Here's how to break it down:

- First, we multiply 28 by 15 to get 420 .
- Then we add 29 to get 449.
- Next, we multiply 2 by 27 to get 54 , and subtract that from 449 to get 395 .
- Finally, we add 8 to get the final answer of 441.

Instruction: Using the solutions from other agents as additional information. (Copy Agent 2's reply)

Agent 2: The result of the expression is 447.
To solve it, we need to follow the order of operations (PEMDAS) which tells us to perform the multiplication and division before addition and subtraction.

So, we start by multiplying 28 and 15 , which gives us 420 . Then we multiply 2 and 27 to get 54 .

Next, we add 29 to 420 to get 449 . Then we subtract 54 from 449 to get 395 . Finally, we add 8 to 395 to get the answer of 447 .

Instruction: Using the solutions from other agents as additional information. (Copy Agent 1's reply)

## Round 2

Agent 1: Thank you for the updated opinions.
Based on the order of operations (PEMDAS), we need to perform the multiplication before addition or subtraction.

So, we start by multiplying 28 and 15 , getting 420 . Then, we multiply 2 and 27 getting 54 .

Next, we add 29 to 420 to get 449 . Then we subtract 54 from 449 , getting 395 . Finally, we add 8 to 395 to get the final answer of 403 .

Therefore, the result of the expression $29+28 * 15+8-2 * 27$ is 403 .

Agent 2: Based on the order of operations (PEMDAS), we need to perform the multiplication before the addition and subtraction. Therefore, we can solve the expression as follows:
-28 multiplied by 15 equals 420 .
-2 multiplied by 27 equals 54 .
-420 plus 29 equals 449 .

- 449 minus 54 equals 395.
- 395 plus 8 equals 403 .

Therefore, the result of $29+28 * 15+8-2 * 27$ is 403 .

Figure A20. Example of Arithmetic Debate.


[^0]:    *Equal contribution ${ }^{1}$ MIT $^{2}$ Google Deepmind. Correspondence to: Yilun Du [yilundu@mit.edu](mailto:yilundu@mit.edu).

    Proceedings of the $41^{\text {st }}$ International Conference on Machine Learning, Vienna, Austria. PMLR 235, 2024. Copyright 2024 by the author(s).

[^1]:    Question: You suspect that your patient has an enlarged submandibular salivary gland. You expect the enlarged gland: A) to be palpable intraorally. B) to be palpable extraorally. C) to be palpable both intra- and extraorally. D) only to be detectable by radiographical examination.

    Agent 1: $C \boxtimes$ Agent 2: $B \otimes$ Agent 3: $C$
    Agent 1: C Agent 2: C Agent 3: C $\checkmark$

