
Conformal Predictions under Markovian Data

Frédéric Zheng¹ Alexandre Proutiere¹

Abstract

We study the split Conformal Prediction method when applied to Markovian data. We quantify the gap in terms of coverage induced by the correlations in the data (compared to exchangeable data). This gap strongly depends on the mixing properties of the underlying Markov chain, and we prove that it typically scales as $\sqrt{t_{\text{mix}} \ln(n)/n}$ (where t_{mix} is the mixing time of the chain). We also derive upper bounds on the impact of the correlations on the size of the prediction set. Finally we present K -split CP, a method that consists in thinning the calibration dataset and that adapts to the mixing properties of the chain. Its coverage gap is reduced to $t_{\text{mix}}/(n \ln(n))$ without really affecting the size of the prediction set. We finally test our algorithms on synthetic and real-world datasets.

1. Introduction

Machine Learning (ML) algorithms and in particular those based on deep neural networks are being increasingly used by practitioners in critical decision processes. If the resulting models often achieve unprecedented levels of prediction performance, they most often come without any guarantees, which can be too hazardous in application fields where safety is crucial (including health care, self driving vehicles, etc.). Introduced and developed by Vovk, see (Vovk et al., 2005; Gammernan & Vovk, 2007; Shafer & Vovk, 2008), Conformal Prediction (CP) is a robust and flexible tool to quantify and handle the inherent uncertainty of these models. The CP framework allows us to build, from any black-box prediction model, a new model whose output is a prediction set with guaranteed level of certainty. CP has become very popular recently, see (Lei & Wasserman, 2014; Lei et al., 2018; Romano et al., 2019; Angelopoulos & Bates, 2023) and references therein, and applied to various learning tasks

¹Division of Decision and Control Systems, EECS KTH and Digital Futures, Stockholm, Sweden. Correspondence to: Frédéric Zheng <fzheng@kth.se>.

including regression and classification.

CP typically works as follows. Consider a model $\hat{\mu}$ fitted to a training dataset and with input (or covariate) X and output $\hat{\mu}(X)$. Based on this model and on a *calibration* dataset of n (covariate, response) pairs $\{(X_1, Y_t)\}_{t=1}^n$, CP builds, for a given certainty level $1 - \alpha$, $C_n(X_{n+1})$ a set of most likely responses. This set is chosen with size as small as possible while ensuring the following coverage guarantee $\mathbb{P}[Y_{n+1} \in C_n(X_{n+1})] \geq 1 - \alpha$.

The coverage guarantees achieved using CP crucially require the calibration data to be sampled independently and identically (i.i.d.) or at least to be exchangeable. This assumption does not hold in many real-world datasets, where indeed samples may be highly correlated. Such correlations arise naturally in learning tasks pertaining to time series (e.g. predicting the evolution of market behavior) (Gibbs & Candes, 2021; Zaffran et al., 2022) or to more general dynamical systems (e.g. in Reinforcement Learning or adaptive control) (Foffano et al., 2023; Dixit et al., 2023). In this paper, we aim at studying CP methods when the correlations in the data are modelled as a Markov chain ($\{(X_1, Y_t)\}_{t \geq 1}$ is a Markov chain). This model is general and includes the classical scenario where the successive covariates form a Markov chain and where the response Y_t remains independent of $\{(X_s, Y_s)\}_{s \neq t}$ (Bresler et al., 2020; Oliveira et al., 2022). For this model, we address the following questions. (a) *How are the coverage and the prediction set size affected by the correlations in the data?* (b) *How can we adapt the CP methods to minimize the impact of these correlations?*

Contributions.

1. We first provide a theoretical analysis of marginal coverage and of the prediction set size under the classical split CP method in the case of Markovian data. The additional coverage gap due to the correlations depends on the mixing properties of the underlying Markov chain, and we establish that under generic ergodicity assumptions, this gap scales at most as $\sqrt{t_{\text{mix}} \frac{\ln(n)}{n}}$ where n is the size of the calibration dataset and t_{mix} denotes the mixing time of the Markov chain. We further show that typically, the increase in the size of the prediction interval due to the correlations does not exceed $\sqrt{t_{\text{mix}}/n}$.
2. We then investigate the idea of thinning the calibration dataset so as to alleviate the impact of correlations. We

refer to as the K -split CP, the resulting method where one in K samples of the calibration dataset is kept. We optimize the value of K to achieve an efficient trade-off between coverage and size of the prediction set. The optimal value depends on the mixing time of the Markov chain, and can be estimated in an online manner without really impacting the coverage. K -split CP improves the coverage gap that now scales $\frac{t_{\text{mix}}}{n \ln(n)}$ and has very little impact on the size of the prediction interval.

3. Our theoretical results are confirmed using numerical experiments, both on synthetic and real-world data (e.g., for the prediction on the EUR/SEK exchange rate).

2. Related Work

Extending the CP framework to non-exchangeable data has been considered in the literature, e.g. to model distribution shifts (Tibshirani et al., 2019; Barber et al., 2023). The work (Oliveira et al., 2022) is close to ours. There, the authors investigate the Split CP for dependent data verifying a set of concentration-type assumptions, such as stationary β -mixing chains. They apply concentration inequalities on an independent calibration dataset and come back to the original trajectory by adding an additional β coefficient. Our analysis of split CP is inspired by their work. However, we leverage the Markovian nature of the data and in particular concentration results available for Markov chains. The key differences between (Oliveira et al., 2022) and our work are: 1) our analysis can be conducted without the stationarity assumption (which we believe is very restrictive); 2) we get more explicit, more generic and simpler upper bounds on the coverage gap for Markovian data; 3) we present an analysis of the size of the prediction interval; 4) we present and investigate the idea of thinning the calibration dataset.

A different line of research consists in adjusting the empirical quantile level to account for the possible undercoverage due to the correlations. For example the authors of (Gibbs & Candes, 2021; Zaffran et al., 2022) propose that for each new observed covariate X_n , the quantile level α_n used to compute the next CP interval is modified (α_n is updated following a classic stochastic approximation scheme).

3. Preliminaries

In this section, we first introduce our models, and provide a few existing results on Markov chains and β -mixing stochastic processes that will be useful in our analysis¹. We then recall the classical Conformal Prediction framework for i.i.d. or exchangeable data.

¹Refer to (Meyn & Tweedie, 1993) for a more extensive exposition on Markov chains with general state space, and to (Levin & Peres, 2017) for mixing properties of Markov chains.

3.1. Markovian data

We assume that $\{(X_t, Y_t)\}_{t \geq 1-N}$ ² is an homogeneous Markov chain with kernel P . The covariates $\{X_t\}_{t \geq 1-N}$ take values in $\mathcal{X} \subset \mathbb{R}^d$, and the responses $\{Y_t\}_{t \geq 1-N}$ in $\mathcal{Y} \subset \mathbb{R}^r$. We assume that the Markov chain is ϕ -irreducible, aperiodic, and admits a stationary distribution π . We define $Z_t = (X_t, Y_t)$ and $\mathcal{Z} = \mathcal{X} \times \mathcal{Y}$.

Example 1. An important example where $Z = \{Z_t\}_{t \geq 1-N}$ is indeed a Markov chain is the case studied in (Bresler et al., 2020) where $X = \{X_t\}_{t \geq 1-N}$ is an homogeneous Markov chain of kernel P_X , and where given $X_t = x$, the response Y_t is independent of $\{Z_s\}_{s \neq t}$ and of distribution $b(x, \cdot)$. For instance, one may assume that $Y_t = \mu(X_t) + \varepsilon_t$ where μ is an unknown function and the noise process $\{\varepsilon_t\}_{t \geq 1-N}$ is i.i.d. and independent of $\{X_t\}_{t \geq 1-N}$. The kernel of Z is simply $P_X(x, dx')b(x', dy)$ and its stationary distribution $\pi_X(dx)b(x, dy)$ where π_X is that of X . Our results apply to this example but to the more general case where Z is a Markov chain.

Mixing time. The mixing time of the Markov chain quantifies the time it takes to approach its steady state. When its transition kernel is P , the mixing time is defined as $t_{\text{mix}} = \tau(\frac{1}{4})$ where for any $\varepsilon > 0$, $\tau(\varepsilon) = \min\{t \geq 1 : \sup_z \|P^t(z, \cdot) - \pi\|_{TV} \leq \varepsilon\}$.

Geometric ergodicity. Consider a ϕ -irreducible and aperiodic Markov chain with stationary distribution π . The chain is geometrically ergodic (Roberts & Rosenthal, 1997) if there exists a constant $0 \leq \rho < 1$, referred to as the rate of the chain, such that for π -a.e. $z \in \mathcal{Z}$, there exists $Q(z) < \infty$ with for all $n \geq 1$, $\|P^n(z, \cdot) - \pi\|_{TV} \leq Q(z)\rho^n$. We know (Nummelin & Tuominen, 1982) that the function Q can be chosen so that it is π -integrable ($\int_{\mathcal{Z}} Q(z)\pi(dz) < \infty$). The chain is *uniformly* geometrically ergodic if the constant $Q(z)$ is independent of z .

Note that if the state space \mathcal{Z} is finite, an aperiodic and irreducible Markov chain is uniformly geometrically ergodic (Levin & Peres, 2017). Also observe that in the case of Example 1, if X is geometrically ergodic with rate ρ , then so is Z with the same rate.

Connection between ρ and t_{mix} . These two quantities, involved in the statement of our results, are closely linked in the case of uniformly geometric Markov chains, as shown in (Paulin, 2015) (see Proposition 3.4). Indeed, the rate ρ of a uniformly geometric ergodic Markov chain can be chosen so that $\rho \leq \sqrt{1 - \gamma_{\text{ps}}}$ where γ_{ps} denotes its pseudo-spectral gap. Combined with the fact that $\gamma_{\text{ps}} \geq \frac{1}{2t_{\text{mix}}}$, we obtain $\rho \leq \sqrt{1 - \frac{1}{2t_{\text{mix}}}}$. When the chain is reversible, it

²Starting the process at time $-N + 1$ will turn convenient notation-wise. The first samples up to time 0 will constitute the training data, and the remaining samples the calibration data.

admits an absolute spectral gap $\gamma \geq \gamma_{\text{ps}}$, leading to a tighter bound of $\rho \leq 1 - \gamma$. For a more exhaustive discussion on the connection between the mixing time, the rate and the spectral gap or its variants for non-reversible Markov chain, refer to (Chatterjee, 2023).

3.2. β -mixing processes

As in (Oliveira et al., 2022), we are interested in β -mixing processes as they can be divided into approximately independent blocks. This property will be instrumental the analysis of split CP. Let $\{Z_t\}_{t \geq 1}$ be a stochastic process. For any $s > t \geq 1$, we denote by $Z_s^t = (Z_s, Z_{s+1}, \dots, Z_t)$, by \mathbb{P}_s^t its distribution, and by $\sigma(Z_s^t)$ the σ -algebra generated by Z_s^t . The β -coefficients of the process are defined as $\beta(a) = \sup_{t \geq 1} \mathbb{E} \sup\{|\mathbb{P}(B|\sigma(Z_1^t)) - \mathbb{P}(B)|, B \in \sigma(Z_{t+a}^\infty)\}$ (Davydov, 1973; Yu, 1994). The process is called β -mixing if $\beta(a) \rightarrow 0$ as $a \rightarrow \infty$. When the process admits a stationary distribution π , we will also use a slightly different definition of the β -coefficients: $\beta'(a) = \sup_{t \geq 1} \mathbb{E}[|\mathbb{P}_{t+a}[\cdot|Z_1^t] - \pi|_{TV}]$, where $\mathbb{P}_{t+a}[\cdot|Z_1^t]$ is the conditional distribution of Z_{t+a} given Z_1^t .

The connection between the rate of convergence of Markov chains and their β -mixing coefficients have been extensively studied. We cite here results from (Davydov, 1973; Nummelin & Tuominen, 1982; Liescher, 2005) that will turn instrumental in our analysis.

Stationary Geometrically Ergodic Markov chains. Let $\{Z_t\}_{t \geq 1}$ be a geometrically ergodic Markov chain with rate $\rho < 1$, and with stationary distribution π . Assume that it is in steady-state ($Z_1 \sim \pi$), then its β -mixing coefficients satisfy (Davydov, 1973): $\beta(a) \leq C\rho^a$, with $C = \int_{\mathcal{Z}} Q(z)\pi(dz)$.

Non-stationary Geometrically Ergodic Markov chains. Let $\{Z_t\}_{t \geq 1}$ be a geometrically ergodic Markov chain with rate ρ , and with stationary distribution π . Let ν_1 be its initial distribution and suppose that $\int_{\mathcal{Z}} Q(z)\nu_1(dz) < \infty$. Then we can show as in the proof of Proposition 3 in (Liescher, 2005) that for all $a \geq 1$, $\beta'(a) \leq 3C'(\sqrt{\rho})^a$ where $C' = \max(\int_{\mathcal{Z}} Q(z)\pi(dz), \int_{\mathcal{Z}} Q(z)\nu_1(dz))$.

3.3. Split Conformal Prediction

We outline below the classical Split CP framework and its basic coverage guarantees in the case of i.i.d. or exchangeable data. Assume that we have access to a dataset $\{(X_t, Y_t)\}_{t=1-N}$. Given a new data point X_{n+1} , we wish to create a conformal prediction set $C_n(X_{n+1})$ with minimal size such that Y_{n+1} lies in $C_n(X_{n+1})$ with probability at least $1 - \alpha$.

In the original split CP framework, the dataset is divided into two parts in a random split as the data is exchangeable. However to handle temporal dependencies in our

analysis, we use a sequential split: the training dataset is $D_{\text{tr}} = \{(X_t, Y_t)\}_{t=1-N}^0$ and the calibration dataset $D_{\text{cal}} = \{(X_t, Y_t)\}_{t=1}^n$. The first dataset is used to fit model $\hat{f}_N : \mathcal{X} \rightarrow \mathbb{R}^r$. Based on this model, a score function $s : \mathcal{X} \times \mathbb{R}^r \rightarrow \mathbb{R}$ is designed (e.g., this function could simply be $s(x, y) = \|\hat{f}_N(x) - y\|$). Next, we compute the scores of the samples in the calibration dataset $s_1 = s(X_1, Y_1), \dots, s_n = s(X_n, Y_n)$, as well $\hat{q}_{\alpha, n}$ as the $\frac{[(n+1)(1-\alpha)]}{n}$ quantile of these scores. Given the new data point X_{n+1} , the conformal prediction set $C_n(X_{n+1})$ is finally defined as $\{y : s(X_{n+1}, y) \leq \hat{q}_{\alpha, n}\}$.

In the case of i.i.d. (or exchangeable) calibration data (i.e., $\{(X_i, Y_i)\}_{i=1}^{n+1}$ are i.i.d.), the above construction enjoys the following marginal coverage guarantee (Vovk et al., 2005):

$$1 - \alpha \leq \mathbb{P}[Y_{n+1} \in C_n(X_{n+1})] \leq 1 - \alpha + \frac{1}{1+n}. \quad (1)$$

We investigate whether this guarantee holds in the case of Markovian data and in the case it does not, how we can adapt the above framework.

3.4. Notation

The notation $v_{N,m} = \mathcal{O}_{N,n}(w(N, n))$ means that there exist $M > 0$ and two integers N_0, n_0 such that for all $N \geq N_0 \vee \log^2(n)$, $n \geq n_0$, $|v_{N,m}| \leq Mw(N, n)$. The notation \mathcal{O}_n, o_n are the usual big-O and little-o notation as $n \rightarrow \infty$. $Z(N, n) = o_{\mathbb{P}}(1)$ means that $\lim_{n \rightarrow \infty} \lim_{N \rightarrow \infty} \mathbb{P}[|Z(N, n)| \geq \epsilon] = \lim_{N \rightarrow \infty} \lim_{n \rightarrow \infty} \mathbb{P}[|Z(N, n)| \geq \epsilon] = 0$ for all $\epsilon > 0$. Let $\mathcal{L}(A)$ be the Lebesgue measure of the set A and $A \Delta B$ be the symmetric difference between sets A, B .

4. Split CP for Markov Chains

In this section, we apply the classical Split CP framework to Markovian data, and evaluate the coverage gap induced by the non-exchangeability of the data. We also study the size of the conformal prediction set.

4.1. Marginal coverage

We formally define the notion of *coverage gap* as follows. Let $C_n(X_{n+1})$ be a prediction set based on the n last observed samples taken from the calibration dataset. We say that it has a coverage gap smaller than $\gamma \geq 0$ if:

$$\mathbb{P}[Y_{n+1} \in C_n(X_{n+1})] \in \left[1 - \alpha - \gamma, 1 - \alpha + \gamma + \frac{1}{n}\right]. \quad (2)$$

Consider the Markov chain $\{(X_t, Y_t)\}_{t \geq -N+1}$ with initial distribution ν_0 (i.e., $(X_{1-N}, Y_{1-N}) \sim \nu_0$) and stationary distribution π . In the following, to simplify the notation,

we define $\delta(a) = \|\nu_0 P^a - \pi\|_{TV}$ for any $a \in \mathbb{N}$. When applying the standard Split CP to this Markov chain, the induced coverage gap will necessarily depends on its transient behavior, partly described by t_{mix} and by the sequences $\{\delta(a)\}_{a \geq 0}$ and $\{\beta(a)\}_{a \geq 0}$. To analyze the coverage gap, we consider two scenarios:

1. *With restart.* In this scenario, we assume that the calibration and the training datasets are independent, and that the Markov chain is restarted with distribution ν_1 at time 1. Assuming such independence makes the conditioning on the training dataset easy and simplifies the analysis. When applying a restart, we will use the notation: for $a \geq 0$, $\delta_1(a) = \|\nu_1 P^a - \pi\|_{TV}$.

2. *Without restart.* Here, the samples composing the training and calibration datasets come from a single trajectory of the Markov chain. The datasets are hence not independent and the distribution of (X_1, Y_1) is $\nu_0 P^N$. We tackle this case by creating a separation of $r \in [n]$ time-steps between the training and calibration datasets so as to reduce their stochastic dependence.

The following proposition provides upper bounds on the coverage gap in both scenarios. These upper bounds will be optimized in various cases depending on the nature of the Markov chain.

Proposition 4.1. (1. *With restart*) Applying Split CP with restart yields a coverage gap γ satisfying, for any $u > \frac{1}{n} \sum_{a=1}^n \delta_1(a)$, $\gamma \leq \gamma(u)$ where

$$\gamma(u) = u + e^{-\frac{2n}{9t_{\text{mix}}}} \left(u - \frac{1}{n} \sum_{a=1}^n \delta_1(a)\right)^2 + \delta_1(n+1). \quad (3)$$

(2. *Without restart*) Applying Split CP yields a coverage gap γ satisfying, for any $u > \delta(N)$ and for any $r \in [n]$, $\gamma \leq \gamma(u, r)$ where

$$\begin{aligned} \gamma(u, r) = & u + e^{-\frac{2(n-r)}{9t_{\text{mix}}}} \left(u - \delta(N)\right)^2 + \delta(n+N+1) \\ & + 2\beta(r) + \frac{1 + \alpha r}{n+1}. \end{aligned} \quad (4)$$

We may optimize the above coverage gap upper bounds over u and r . The following theorem summarizes the outcomes of this optimization.

Theorem 4.1. (1. *With restart*) Applying Split CP with restart yields a coverage gap γ satisfying

$$\gamma = \mathcal{O}_n \left(\sqrt{t_{\text{mix}} \frac{\ln(n)}{n}} + \delta_1(n+1) + \frac{1}{n} \sum_{a=1}^n \delta_1(a) \right). \quad (5)$$

We deduce that, if the Markov chain is geometrically ergodic, then $\gamma = \mathcal{O}_n \left(\sqrt{t_{\text{mix}} \frac{\ln(n)}{n}} \right)$.

(2. *Without restart*) Applying Split CP to a geometrically

ergodic Markov chain yields a coverage gap γ satisfying

$$\gamma = \mathcal{O}_{N,n} \left(\sqrt{t_{\text{mix}} \frac{\ln(n)}{n}} \right). \quad (6)$$

Proposition 4.1 and Theorem 4.1 quantify the coverage gap due to the non-exchangeability of the data. The gap depends (in both cases with or without restart) on the mixing properties of the underlying Markov chain. As expected, the gap becomes negligible only if the number of samples in the calibration dataset significantly exceeds the mixing time.

4.2. Size of the conformal prediction set

To investigate the impact of the non-exchangeability of the data on the size of the conformal prediction set, we use the same setting as that used in (Lei et al., 2018). We consider real valued responses, and we make the following assumptions (similar to those made in Example 1).

Assumption 1. For any $t \geq -N + 1$ and x , given $X_t = x$, the response Y_t can be written as $\mu(x) + \varepsilon$ where the density of the noise ε is symmetric around 0 and non-increasing on \mathbb{R}_+ . The training dataset is used to design an estimator $\hat{\mu}_N$ of the function μ , and the score function used in the CP method is defined through the residuals: $s_t = |Y_t - \hat{\mu}_N(X_t)|$.

The next assumption states that the estimator $\hat{\mu}_N$ becomes accurate as the size of the training dataset increases. This assumption is not critical, and results for the size of the prediction set can also be obtained by just assuming stability (Lei et al., 2018), i.e., convergence of $\hat{\mu}_N$ to a given $\tilde{\mu}$ possibly different than μ .

Assumption 2. There exist two sequences $\{c_N\}_{N \geq 1}$ and $\{d_N\}_{N \geq 1}$ such that $c_N = o_N(1)$, $d_N = o_N(1)$ and $\mathbb{P}(\|\hat{\mu}_N - \mu\|_\infty \geq c_N) \leq d_N$.

Assumption 3. The noise $|Y - \mu(X)|$ admits a density f that is lower bounded by $\kappa > 0$ on a neighborhood \mathcal{N} of its $(1 - \alpha)$ quantile q_α .

Under the above assumptions and should the data be i.i.d., we know (Lei & Wasserman, 2014) that the optimal conformal prediction set (that with conditional coverage $1 - \alpha$ and minimal size) is given $X_{n+1} = x$, $C^*(x) = [\mu(x) - q_\alpha, \mu(x) + q_\alpha]$, where q_α is the $(1 - \alpha)$ quantile of the law of $|\varepsilon|$. We use the set $C^*(x)$ as a benchmark to study the impact of (i) non-i.i.d. data and (ii) of the fact that μ and q_α are unknown. We compare $C^*(x)$ to the returned conformal prediction set $C_n(x) = [\hat{\mu}_N(x) - \hat{q}_{\alpha,n}, \hat{\mu}_N(x) + \hat{q}_{\alpha,n}]$.

4.2.1. CONCENTRATION OF EMPIRICAL QUANTILES

We start quantifying the deviation between the empirical quantile $\hat{q}_{\alpha,n}$ used to build the prediction set and its true

counterpart \hat{q}_α , the $(1 - \alpha)$ quantile of $|Y - \hat{\mu}_N(X)|$ when $(X, Y) \sim \pi$. We introduce the set $\mathcal{U} = \{t > 0 : [q_{\alpha-t}, q_{\alpha+t}] \subset \mathcal{N}\}$.

Proposition 4.2. *Suppose that Assumptions 1, 2, 3 hold. (1. With restart) Let $u \geq 0$ and let $u' = u - (2\kappa c_N + d_N)$. Assume that u' lie in \mathcal{U} , and that $u' > \frac{1}{n} \sum_{a=1}^n \delta_1(a)$. Then we have:*

$$\mathbb{P}(|\hat{q}_{\alpha,n} - \hat{q}_\alpha| > \frac{u}{\kappa}) \leq 2e^{-\frac{2n}{9t_{\text{mix}}}} \left(u' - \frac{1}{n} \sum_{a=1}^n \delta_1(a)\right)^2 + d_N.$$

Therefore, for any $\delta > 0$ with probability at least $1 - \delta - d_N$, we have $|\hat{q}_{\alpha,n} - \hat{q}_\alpha| \leq \frac{u^*}{\kappa}$ with

$$u^* = d_N + 2\kappa c_N + \frac{1}{n} \sum_{i=1}^n \delta_1(a) + \sqrt{\frac{9t_{\text{mix}} \ln(\frac{2}{\delta})}{2n}},$$

provided that $u^* \in \mathcal{U}$.

(2. Without restart) For any $r \in [n]$ and $u \in \mathbb{R}$, let $u' = u + \frac{1+r\alpha}{n+1} - (2\kappa c_N + d_N)$. Assume that both $\frac{1+r\alpha}{n+1}$ and u' lie in \mathcal{U} and that $u' > \delta(N)$. Then we have:

$$\mathbb{P}(|\hat{q}_{\alpha,n} - \hat{q}_\alpha| > \frac{u}{\kappa}) \leq 2e^{-\frac{2(n-r)}{9t_{\text{mix}}}} \left(u' - \delta(N)\right)^2 + d_N + \beta(r).$$

Further assuming that the Markov chain is geometrically ergodic with rate ρ , we deduce that for any $\delta > 0$ with probability at least $1 - \delta - d_N - \frac{2}{n}$, we have $|\hat{q}_{\alpha,n} - \hat{q}_\alpha| \leq \frac{u^*}{\kappa}$ with

$$u^* = d_N + 2\kappa c_N + \delta(N) + \alpha \frac{\ln(n)}{n \ln(\frac{1}{\rho})} + \sqrt{\frac{9t_{\text{mix}} \ln(\frac{2}{\delta})}{2n}},$$

provided that $u^* \in \mathcal{U}$.

Note that when the Markov chain is geometrically ergodic, with or without restart, we indeed obtain that $\hat{q}_{\alpha,n}$ concentrates around \hat{q}_α as n and N grow large. Further observe that as a direct application of the above proposition, we recover the result from (Kolla et al., 2019) stating that in the case of i.i.d. data, with probability at least $1 - \delta$, $|\hat{q}_{\alpha,n} - \hat{q}_\alpha| \leq c\sqrt{\ln(2/\delta)/2n}$ where c is a constant that depends on κ . Handling Markovian data induces an additional multiplicative factor $\sqrt{t_{\text{mix}}}$ in the difference between the true and the estimated quantile of the residuals.

4.2.2. ASYMPTOTIC OPTIMALITY

To compare $C_n(x)$ to $C^*(x)$, we further need to make an assumption about the consistency of the estimated model $\hat{\mu}_N$. We make the same assumption as in (Lei et al., 2018):

Assumption 4. *The estimated model $\hat{\mu}_N$ satisfies: $\Delta_N = \mathbb{E}_{X \sim \pi, \text{tr}}[(\hat{\mu}_N(X) - \mu(X))^2] = o_N(1)$.*

As shown in (Lei et al., 2018), this assumption allows us to control the gap between the true quantile q_α and that

of the estimated model \hat{q}_α . Indeed, if we assume that the density function of $|\varepsilon|$ has continuous derivate upper bounded by $M > 0$ and is lower bounded by $\kappa > 0$ on $(q_\alpha - \eta, q_\alpha + \eta)$ for some $\eta > (M/2\kappa)\Delta_N$, then $|\hat{q}_\alpha - q_\alpha| \leq (M/2\kappa)\Delta_N$. Combining this result to that of Proposition 4.2, we can finally quantify the difference between $C_n(X_{n+1})$ and $C^*(X_{n+1})$.

Proposition 4.3. *Assume that the Markov chain $Z = (X, Y)$ is β -mixing. We have:*

$$\mathcal{L}(C^*(X_{n+1}) \Delta C_n(X_{n+1})) = o_{\mathbb{P}}(1),$$

if either (i) the Markov chain is in steady-state (i.e., $\nu_0 = \pi$ without restart or $\nu_1 = \pi$ with restart), and Assumptions 1, 3, 4 hold, or (ii) Assumptions 1, 2, 3, 4 hold.

The proposition states that asymptotically, when the calibration dataset grows large, the prediction set becomes very close to the optimal conformal prediction set. It is worth noting that we do not assume that the Markov chain is geometrically ergodic. We do not quantify the speed at which $C_n(X_{n+1})$ converges to $C^*(X_{n+1})$, and hence assuming β -mixing is enough.

5. K -Split Conformal Prediction

The coverage gap γ identified when applying the classical CP approach to Markovian data stems from the correlations in the data. A natural way to decrease the impact of these correlations consists in thinning the calibration data. Specifically, we may build the prediction set based on calibration samples taken every K steps. We refer to this approach as the *K -split conformal prediction*. It applies the classical CP method to the calibration data $D_{\text{cal},K} = \{(X_{iK+1}, Y_{iK+1})\}_{i=0}^{\lfloor \frac{n}{K} \rfloor - 1}$. In the following, for simplicity, we assume that $n = Km$ for some $m \in \mathbb{N}$ so that $D_{\text{cal},K}$ consists of m samples. We define $\hat{q}_{\alpha,K,n}$ as the $\frac{\lfloor (m+1)(1-\alpha) \rfloor}{m}$ quantile of the scores $s_1 = s(X_1, Y_1), \dots, s_m = s(X_{K(m-1)+1}, Y_{K(m-1)+1})$. Given the new data point X_{n+1} , the conformal prediction set $C_K(X_{n+1})$ is defined as $\{y : s(X_{n+1}, y) \leq \hat{q}_{\alpha,K,n}\}$.

The analysis made in the previous section extends here since the thinned process remains a Markov chain but with transition kernel P^K . However, applying Theorem 4.1 does not provide interesting results. Indeed, observe that the mixing time of thinned process is $t_{\text{mix},K} = \lceil \frac{t_{\text{mix}}}{K} \rceil$. This implies that the coverage gap γ obtained in Theorem 4.1 would scale as

$$\sqrt{\lceil \frac{t_{\text{mix}}}{K} \rceil K \frac{\ln(n/K)}{n}} \geq \sqrt{t_{\text{mix}} \frac{\ln(n/K)}{n}}.$$

Hence, from this analysis, we do not see any improvement in thinning the process. We propose below an alternative analysis based on the so-called *blocking technique* (that can be traced back to (Bernstein, 1927)).

5.1. Marginal coverage guarantees via the blocking technique

The blocking technique applies not only to Markov chains but also to β -mixing processes. Hence we start by stating general results for these processes, and then specialize the results in the case of Markov chains.

5.1.1. β -MIXING PROCESSES

Let $\{X_t, Y_t\}_{1-N}^n$ a stochastic process with initial distribution ν_0 and stationary distribution π . For any $r \in [n]$, denote by $\beta'_r(a)$ the mixing coefficient calculated on the trajectory $\{(X_t, Y_t)\}_{t=r}^n$, i.e., $\beta'_r(a) = \sup_{t \geq r} \mathbb{E}[\|\mathbb{P}^{t+a}[\cdot | (X_i, Y_i)_{i=r}^t] - \pi\|_{TV}]$.

Proposition 5.1. (1. *With restart*) Under the K -Split CP with restart, $\mathbb{P}[Y_{n+1} \in C_K(X_{n+1})]$ belongs to the interval $[1 - \alpha - \gamma(K), 1 - \alpha + \gamma(K) + \frac{K}{n}]$ for any $n \geq 1$ and any $K \in [n]$ where

$$\gamma(K) = \begin{cases} 2\frac{n}{K}\beta(K) & \text{if } \nu_1 = \pi, \\ 2\frac{n}{K}\beta'_1(K) & \text{otherwise.} \end{cases} \quad (7)$$

(2. *Without restart*) Under the K -Split CP without restart, $\mathbb{P}[Y_{n+1} \in C_K(X_{n+1})]$ belongs to the interval $[1 - \alpha - \gamma(K, r), 1 - \alpha + \gamma(K, r) + \frac{K}{n-r}]$ for any $n \geq 1, r \in [n], K \in [n-r]$ where

$$\gamma(K, r) = \begin{cases} \frac{1+\alpha r}{n+1} + 2\frac{n-r}{K}\beta(K) + \beta(r) & \text{if } \nu_0 = \pi, \\ \frac{1+\alpha r}{n+1} + 2\frac{n-r}{K}\beta'_r(K) + \beta(r) & \text{otherwise.} \end{cases} \quad (8)$$

The above bounds on the coverage gap are rather simpler than those obtained in the case without thinning (see Proposition 4.1) and only depend on the β coefficients of the process. Note that these bounds do make sense only when these coefficients decrease rapidly with K . They are also useless in the case without thinning ($K = 1$).

The term $(1 - \alpha)$ appearing in the lower bounds of the coverage stems from the analysis of split CP in the i.i.d setting. In this setting, due to the finiteness of the calibration dataset, the coverage is actually lower bounded by :

$$\text{coverage} \geq \frac{[(n_{\text{cal}} + 1)(1 - \alpha)]}{n_{\text{cal}} + 1} \geq 1 - \alpha, \quad (9)$$

where n_{cal} is the size of the calibration dataset which is $\frac{n}{K}$ for K -split CP. When n_{cal} (resp. $\frac{n}{K}$) is small, the difference between the two bounds can be non negligible and lead to an over-coverage of split CP (resp. K -split CP). We illustrate this phenomenon and propose a slight refinement of the analysis of K -split CP in Section 6.

5.1.2. GEOMETRICALLY ERGODIC MARKOV CHAINS

Next we optimize the value of K to achieve a good trade-off between coverage and size of the prediction set. Observe that indeed such a trade-off exists. If we just want to maximize coverage, then we can select values of K maximizing the functions $\gamma(K)$ or $\gamma(K, r)$ defined in Proposition 5.1. This would lead to choosing K very large, but would be at the expense of enlarging the prediction set. This trade-off is confirmed by the terms K/n and $K/(n-r)$ in the upper limit of the intervals where $\mathbb{P}[Y_{n+1} \in C_K(X_{n+1})]$ lies. To address the coverage-size trade-off, we choose to minimize the size of these intervals. Specifically, for the case with restart, we wish to minimize $\gamma(K) + K/n$ over K ; whereas for the case without restart, the function to minimize is $\gamma(K, r) + K/(n-r)$ over K and r .

To this aim, we assume that the process is a geometrically ergodic Markov chain with rate ρ . In the case without restart, to simplify the optimization in r , we observe that $\beta'_r(a) \leq \beta'_1(a) \leq C'\rho^{\frac{a}{2}}$ for $C' = 3 \max(\int_{\mathcal{X}} Q(x)\pi(dx), \int_{\mathcal{X}} Q(x)\nu_1(dx)) < \infty$ (refer to subsection 3.2).

Theorem 5.1. (1. *With restart*) Assume $\int_{\mathcal{X}} Q(x)\nu_1(dx) < \infty$. Under the K -Split CP with restart and with³ $K^* = \frac{W_0(n^2(\ln \rho)^2)}{\ln(\frac{1}{\rho})}$, $\mathbb{P}(Y_{n+1} \in C_{K^*}(X_{n+1}))$ belongs to $[1 - \alpha - \gamma_{\text{low}}, 1 - \alpha + \gamma_{\text{up}}]$ with $\gamma_{\text{low}} = \mathcal{O}_n(\frac{1}{n \ln(n) \ln(\frac{1}{\rho})})$ and $\gamma_{\text{up}} = \mathcal{O}_n(\frac{\ln(n)}{n \ln(\frac{1}{\rho})})$.

(2. *Without restart*) The same result as in the case of restart holds if we assume $\int_{\mathcal{X}} Q(x)\nu_0 P^N(dx) < \infty$.

This theorem states that for geometrically ergodic Markov chains, the coverage gap of the optimal K -split CP scales as $\frac{1}{n \ln(n) \ln(1/\rho)}$. For clarity, we restate the result with t_{mix} only as in Section 4. As discussed in Section 3, we can choose ρ satisfying $\rho \leq (1 - \frac{1}{2t_{\text{mix}}})^{1/2}$. Therefore, the coverage gap of optimal K -split CP scales as $\frac{t_{\text{mix}}}{n \ln(n)}$. Recall that the gap of the split CP without thinning was scaling as $\sqrt{t_{\text{mix}} \frac{\ln(n)}{n}}$, therefore K -split CP divides the coverage gap by a factor $\sqrt{n \ln(n)}^{3/2}$.

5.2. Size of the conformal prediction set

As mentioned earlier, by thinning the initial Markov chain, we obtain another Markov chain with kernel P^K . Hence our results pertaining to the size of the prediction set and derived in Subsection 4.2 remain valid. More precisely, all statements made in Proposition 4.2 hold provided that we replace n by n/K , t_{mix} by $\lceil \frac{t_{\text{mix}}}{K} \rceil$, and $\frac{1}{n} \sum_{a=1}^n \delta_1(a)$ by $\frac{K}{n} \sum_{a=1}^{\lfloor n/K \rfloor} \delta_1(Ka)$ (note that all terms related to the training data, e.g. c_N, d_N , remain unchanged). When applying

³ W_0 is the Lambert function of order 0.

the results to geometrically ergodic chains, we also have to replace the rate ρ by ρ^K . The following proposition summarizes the above observations.

Proposition 5.2. *Suppose that Assumptions 1, 2, 3 hold. Assume that the Markov chain is geometrically ergodic with rate ρ . Applying the K -Split CP with or without restart yields that: for any $\delta > 0$ with probability at least $1 - \delta - d_N - \frac{2}{n}$, we have $|\hat{q}_{\alpha,n} - \hat{q}_\alpha| \leq \frac{u^*}{\kappa}$ with*

$$u^* = d_N + 2\kappa c_N + \delta(N) + \mathcal{O}_n \left(\sqrt{\frac{(t_{\text{mix}} \vee K) \ln(\frac{2}{\delta})}{n}} \right),$$

provided that $u^* \in \mathcal{U}$.

We can plug K^* , the optimal value of K identified in Theorem 5.1, in the above result. Using the fact that $W_0(x) = O(\ln x)$, we have $K^* = \mathcal{O}_n(\frac{\ln(n)}{\ln(\frac{1}{\rho})}) = \mathcal{O}_n(t_{\text{mix}} \ln(n))$. Hence with this choice of K , the high probability upper bound of $|\hat{q}_{\alpha,n} - \hat{q}_\alpha|$ goes

$$\begin{aligned} & \text{from } \mathcal{O}_n\left(\sqrt{\frac{t_{\text{mix}} \ln(\frac{2}{\delta})}{n}}\right) \quad \text{for the classical Split CP} \\ & \text{to } \mathcal{O}_n\left(\sqrt{\frac{t_{\text{mix}} \ln(n) \ln(\frac{2}{\delta})}{n}}\right) \quad \text{for the } K^*\text{-Split CP.} \end{aligned}$$

This result suggests that the impact of thinning the calibration dataset using the optimal K^* does not impact much the size of conformal prediction set. Finally observe that with thinning, the results of Proposition 4.3 hold.

5.3. Adaptive K -Split CP

As demonstrated in Theorem 5.1, the choice of the thinning parameter K in the K -split CP is critical to achieve an efficient trade-off between coverage and prediction set size. However, K^* is a function of the rate ρ of the Markov chain, and the latter is initially unknown. We assume below that the state space is finite so that the Markov chain is uniformly geometric ergodic. In this case, its rate can be selected as $\rho = \sqrt{1 - \gamma_{\text{ps}}}$ or $\rho = 1 - \gamma$ if the chain is reversible. Both quantities γ_{ps} , γ can be estimated (Hsu et al., 2015; 2019; Combes & Touati, 2019; Wolfer & Kontorovich, 2019). For simplicity, we state the results in the reversible case and assuming that we estimate γ using the training dataset. Our results can however be easily extended to the non-reversible case by estimating γ_{ps} (Wolfer & Kontorovich, 2019).

We use the estimator $\hat{\gamma}_N$ proposed in (Hsu et al., 2015) to construct $\hat{\rho}_N = 1 - \hat{\gamma}_N$. This estimator enjoys the following guarantees: for any $N \geq 1$, for any $\delta \in (0, 1)$,

$$\mathbb{P}_{\text{tr}} \left[|\hat{\rho}_N - \rho| \geq C \sqrt{\frac{\ln(\frac{|\mathcal{X}|}{\delta}) \ln(N)}{\pi_* (1 - \lambda) N}} \right] \leq \delta,$$

where C is a universal constant and $\pi_* := \min_{z \in \mathcal{Z}} \pi(z)$ (here π is the stationary distribution of the Markov chain).

Based on $\hat{\rho}_N$, we build the following estimator for K^* : $\hat{K}_N = \ln(n) / \ln(1/\hat{\rho}_N)$. When plugging this value in our K -Split CP method, we get the following guarantees.

Proposition 5.3. (1. *With restart*) *Assume that $\int_{\mathcal{X}} Q(x) \nu_1(dx) < \infty$. Under the K -Split CP with restart and thinning parameter \hat{K}_N , $\mathbb{P}(Y_{n+1} \in C_{\hat{K}_N}(X_{n+1}))$ belongs to the interval $[1 - \alpha - \gamma_{\text{low}} - l_{N,n}, 1 - \alpha + \gamma_{\text{up}} + u_{N,n}]$ with $\gamma_{\text{low}}, \gamma_{\text{up}}$ as in Theorem 5.1 and*

$$l_{N,n} = \mathcal{O}_{N,n} \left(\frac{\ln(N)}{\sqrt{N}} \frac{1}{\rho(\ln \rho)^2 n \ln(n)} \right),$$

$$u_{N,n} = \mathcal{O}_{N,n} \left(\frac{\ln(N)}{\sqrt{N}} \frac{1}{n} \right)$$

(2. *Without restart*) *The same result as in the case of restart holds if we assume $\int_{\mathcal{X}} Q(x) \nu_0 P^N(dx) < \infty$.*

From the above results, we conclude that the estimation of K^* does not impact much the coverage gap, and one can easily enjoy the benefits of the optimal K -split CP without the apriori knowledge of K^* .

6. Experiments

We illustrate experimentally the coverage gap and set length of our two described conformal prediction methods, namely original split conformal and K -split conformal. The score considered is the usual residual $|Y - \hat{\mu}(X)|$ (all responses are real-valued) therefore the conformal set length is simply $2\hat{q}_{\alpha,n}$ (resp $2\hat{q}_{\alpha,K,n}$). When the model is specified, the latter is compared with the optimal length $2q_\alpha$ (as defined in Section 4.2) through a relative difference.

The objective of our experiments is to mainly assess the performance of the original split CP method in the case of Markovian data, and that of our proposed algorithms, namely K -split CP and corrected K -split CP (as defined below). Note that we do not compare our algorithms to those designed to cope with non-stationary dependent data such as ACI (Gibbs & Candes, 2021; Zaffran et al., 2022) (ACI extends the original CP algorithm by allowing the confidence level α to change in an active and online manner depending on the observed coverage).

6.1. K -split conformal prediction in practice

As discussed in Section 5, K -split CP may exhibit an over-coverage when there are not enough values in the reduced calibration dataset. To circumvent this issue, we can adjust the quantile level used by K -split CP. This level is set to α' by taking into account the difference between the two bounds presented (9). Specifically, we select α' such that

$$\frac{[(\frac{n}{K} + 1)(1 - \alpha')]}{\frac{n}{K} + 1} = 1 - \alpha$$

6.2. Synthetic data

In this subsection, we apply all split CP methods to synthetic examples, one with a finite state space and one with a continuous state space. An experiment consists in generating one trajectory of length $N + n + 1$ and in applying CP to the last point. We repeat the experiment $N_{\text{trials}} = 1000$ times and report the average coverage rate. We fix $N = 10000$.

6.2.1. THE LAZY RANDOM WALK ON $\mathbb{Z}/w\mathbb{Z}$

Consider the stationary lazy random walk. It is a simple example of a finite state space, irreducible, aperiodic and reversible Markov Chain, for which $\rho = \lambda_2 = \frac{1 + \cos(\frac{2\pi}{w})}{2}$. For $w = 20$, this example already exhibits strong temporal correlations as $\rho = 0.97$. The discrete state space setting allows to use the estimator from (Hsu et al., 2015). For a given true model $\mu(x)$ and independent symmetric noise ε_t , we generate $Y_t = \mu(X_t) + \varepsilon_t$ (refer to Appendix B.1 for details).




Figure 1. Coverage for the lazy walk ($w = 20$) as a function of n . The second values on the x -axis represent the (optimal) values of $K^*(n)$.

In Figure 1, we observe that all three methods achieve almost $1 - \alpha$ coverage gap as n increases. Note that K -split CP always has stronger coverage than full split CP, and over-covering steadily diminishes when the number of samples in the reduced dataset increase. The correction proposed allows to remain closer to $1 - \alpha$ in the regime where n is small as expected. Figure 2 shows that all CP methods output an interval whose length approaches the optimal length as n increases. Again, we observe the overcoverage phenomenon when n is low for K -split CP, but asymptotically vanishes as analysed in Section 5.

6.2.2. THE GAUSSIAN AUTOREGRESSIVE (AR) MODEL OF ORDER 1

The classical AR(1) models are reversible Markov chains defined by the following recursive equation: $\forall n, X_{n+1} = \theta X_n + \varepsilon_{n+1}$ with $\varepsilon_n \sim N(0, \omega^2)$ and for some $\theta \in [0, 1[$ and $\omega > 0$. ε_n is a Gaussian noise independent of X .




Figure 2. $\frac{|q_{n,\alpha} - q_\alpha|}{q_\alpha}$ for the lazy walk as a function of n . The second values on the x -axis represent the (optimal) values of $K^*(n)$.

For this stochastic process, we predict the AR model itself, i.e., $Y_t = X_{t+1}$. (X, Y) is a Markov chain with the same kernel as X and it is geometrically ergodic (Bhattacharya & Lee, 1995). Since the state space is continuous, we can not use the estimator defined in Section 5.3, however in this example, we can explicitly compute ρ (refer to Appendix B.2 for details).

In Figure 3, we compare our different CP methods. Again, we observe that K -split CP outperforms full split CP. But all methods achieve a very good coverage, close to $1 - \alpha$ as n grows large.




Figure 3. Coverage for Gaussian AR ($\theta = 0.9, \omega = 1$) as a function of n . The second values on the x -axis represent the (optimal) values of $K^*(n)$.

6.3. Real-world application

In this subsection, we apply split CP to real-world datasets. Experiments to assess the performance of split CP on dependent data have already been conducted in (Wisniewski et al., 2020) (without guarantees) and in (Oliveira et al., 2022) (with guarantees for stationary β -mixing data). We further consider non-stationary data and investigate the performance of K -split CP.

6.3.1. EXCHANGE RATE EUR/SEK

The objective here is to predict the exchange rate EUR/SEK⁴ in 2022. Let X_t be the exchange rate, reported every minute. As done in (Oliveira et al., 2022), we assume that the series of returns $r_t = \frac{X_{t+1}}{X_t} - 1$ is β -mixing with exponentially fast convergence to 0. We estimate the corresponding value ρ by computing the auto-correlations as they exhibit a similar decay as that of the β coefficients, and applying a simple linear regression to their logarithms. This gives a value of $\rho = 0.57$.

At each timestep, we apply conformal prediction to the next return r_{t+1} with a rolling window of fixed size divided into training and calibration datasets (1 month = 30x24x60 data points for each in this example). We compute whether the given CP method covered r_{t+1} in 2022.




Figure 4. Daily coverage for EUR/SEK exchange rate ($n = 43200, N = 43200$) during one month

In Figure 4, we plotted the empirical averages of the coverage over a month. We observe that all three methods achieve $1 - \alpha$ coverage.

6.3.2. ELECTRICITY PRICE FORECASTING

In this second example, we consider the same dataset as (Zaffran et al., 2022), which contains the French electricity price between 2016 to 2019, reported every hour. We consider again the prediction of the one-step return r_{t+1} with a rolling window of fixed size (18 months = 18x30x24 data points for both training/calibration). In this example, we obtained a value of $\rho = 0.78$. We calculate the empirical coverage for the year 2019.

In Figure 7, we can observe that all methods are quite similar in performance, with a coverage close to $1 - \alpha$.

7. Conclusion

In this paper, we extended the analysis of the original split CP method to the case of Markovian data. We established upper bounds on the impact of the correlations in the data on the coverage guarantees and size of the prediction set. When

⁴The dataset can be found at <https://www.histdata.com/>




Figure 5. Monthly coverage for French electricity price forecast ($n = 12960, N = 12960$)

the underlying Markov chain mixes rapidly, this impact is negligible. When this is not the case, handling correlations remains challenging and an interesting topic for future research. We could try for example to identify fundamental limits on the coverage vs. size of the prediction set trade-off satisfied by any conformal method; such limits would indicate the incompressible price one has to pay when dealing with Markovian data. These limits would also provide insights into the design of CP algorithms.

Impact Statement

This paper presents work whose goal is to advance the field of Machine Learning. There are many potential societal consequences of our work, none which we feel must be specifically highlighted here.

Acknowledgements

This research was supported by the Wallenberg AI, Autonomous Systems and Software Program (WASP) funded by the Knut and Alice Wallenberg Foundation, the Swedish Research Council (VR), and Digital Futures.

References

- Angelopoulos, A. N. and Bates, S. Conformal prediction: A gentle introduction. *Found. Trends Mach. Learn.*, 16(4):494–591, mar 2023. ISSN 1935-8237. doi: 10.1561/2200000101. URL <https://doi.org/10.1561/2200000101>.
- Barber, R. F., Candes, E. J., Ramdas, A., and Tibshirani, R. J. Conformal prediction beyond exchangeability. *The Annals of Statistics*, 51(2):816–845, 2023.
- Bernstein, S. Sur l’extension du théorème limite du calcul des probabilités aux sommes de quantités dépendantes. *Mathematische Annalen*, 97:1–59, 1927. URL <https://api.semanticscholar.org/CorpusID:122172457>.
- Bhattacharya, R. N. and Lee, C. Ergodicity of nonlinear first order autoregressive models. *Journal of Theoretical Probability*, 8:207–219, 1995.
- Bresler, G., Jain, P., Nagaraj, D., Netrapalli, P., and Wu, X. Least squares regression with markovian data: fundamental limits and algorithms. In *Proceedings of the 34th International Conference on Neural Information Processing Systems, NIPS’20*, Red Hook, NY, USA, 2020. Curran Associates Inc. ISBN 9781713829546.
- Chatterjee, S. Spectral gap of nonreversible markov chains, 2023.
- Chatzigeorgiou, I. Bounds on the lambert function and their application to the outage analysis of user cooperation. *IEEE Communications Letters*, 17(8):1505–1508, 2013.
- Combes, R. and Touati, M. Computationally efficient estimation of the spectral gap of a markov chain. *Proc. ACM Meas. Anal. Comput. Syst.*, 3(1), mar 2019.
- Davydov, Y. A. Mixing conditions for markov chains. *Teoriya Veroyatnostei i ee Primeneniya*, 18(2):321–338, 1973.
- Dixit, A., Lindemann, L., Wei, S. X., Cleaveland, M., Pappas, G. J., and Burdick, J. W. Adaptive conformal prediction for motion planning among dynamic agents. In Matni, N., Morari, M., and Pappas, G. J. (eds.), *Proceedings of The 5th Annual Learning for Dynamics and Control Conference*, volume 211 of *Proceedings of Machine Learning Research*, pp. 300–314. PMLR, 15–16 Jun 2023. URL <https://proceedings.mlr.press/v211/dixit23a.html>.
- Doukhan, P. *Mixing: Properties and Examples*. Lecture Notes in Statistics. Springer New York, 2012. ISBN 9781461226420.
- Foffano, D., Russo, A., and Proutière, A. Conformal off-policy evaluation in markov decision processes. In *62nd IEEE Conference on Decision and Control, CDC 2023, Singapore, December 13-15, 2023*, pp. 3087–3094. IEEE, 2023. doi: 10.1109/CDC49753.2023.10383469. URL <https://doi.org/10.1109/CDC49753.2023.10383469>.
- Foygel Barber, R., Candes, E. J., Ramdas, A., and Tibshirani, R. J. The limits of distribution-free conditional predictive inference. *Information and Inference: A Journal of the IMA*, 10(2):455–482, 2021.
- Gallegos-Herrada, M. A., Ledvinka, D., and Rosenthal, J. S. Equivalences of geometric ergodicity of markov chains. *Journal of Theoretical Probability*, pp. 1–27, 2023.
- Gammerman, A. and Vovk, V. Hedging predictions in machine learning. *Comput. J.*, 50(2):151–163, mar 2007. ISSN 0010-4620. doi: 10.1093/comjnl/bxl065. URL <https://doi.org/10.1093/comjnl/bxl065>.
- Gibbs, I. and Candes, E. Adaptive conformal inference under distribution shift. *Advances in Neural Information Processing Systems*, 34:1660–1672, 2021.
- Hsu, D., Kontorovich, A., Levin, D. A., Peres, Y., Szepesvári, C., and Wolfer, G. Mixing time estimation in reversible markov chains from a single sample path. *The Annals of Applied Probability*, 29(4):pp. 2439–2480, 2019. ISSN 10505164, 21688737.
- Hsu, D. J., Kontorovich, A., and Szepesvári, C. Mixing time estimation in reversible markov chains from a single sample path. *Advances in neural information processing systems*, 28, 2015.
- Kolla, R. K., Prashanth, L., Bhat, S. P., and Jagannathan, K. Concentration bounds for empirical conditional value-at-risk: The unbounded case. *Operations Research Letters*, 47(1):16–20, 2019.
- Kuznetsov, V. and Mohri, M. Generalization bounds for non-stationary mixing processes. *Machine Learning*, 106(1):93–117, 2017.
- Lei, J. and Wasserman, L. Distribution-free prediction bands for non-parametric regression. *Journal of the Royal Statistical Society: Series B: Statistical Methodology*, pp. 71–96, 2014.
- Lei, J., G’Sell, M., Rinaldo, A., Tibshirani, R. J., and Wasserman, L. Distribution-free predictive inference for regression. *Journal of the American Statistical Association*, 113(523):1094–1111, 2018.
- Levin, D. A. and Peres, Y. *Markov chains and mixing times*, volume 107. American Mathematical Soc., 2017.

- Liebscher, E. Towards a unified approach for proving geometric ergodicity and mixing properties of nonlinear autoregressive processes. *Journal of Time Series Analysis*, 26(5):669–689, 2005.
- Meyn, S. and Tweedie, R. *Markov Chains and Stochastic Stability*. Springer-Verlag, London, 1993. URL [/brokenurl#probability.ca/MT](#).
- Nummelin, E. and Tuominen, P. Geometric ergodicity of harris recurrent markov chains with applications to renewal theory. *Stochastic Processes and Their Applications*, 12(2):187–202, 1982.
- Oliveira, R. I., Orenstein, P., Ramos, T., and Romano, J. V. Split conformal prediction for dependent data. *arXiv preprint arXiv:2203.15885*, 2022.
- Paulin, D. Concentration inequalities for markov chains by marton couplings and spectral methods. 2015.
- Roberts, G. and Rosenthal, J. Geometric ergodicity and hybrid markov chains. 1997.
- Romano, Y., Patterson, E., and Candes, E. Conformalized quantile regression. In Wallach, H., Larochelle, H., Beygelzimer, A., d’Alché-Buc, F., Fox, E., and Garnett, R. (eds.), *Advances in Neural Information Processing Systems*, volume 32. Curran Associates, Inc., 2019. URL https://proceedings.neurips.cc/paper_files/paper/2019/file/5103c3584b063c431bd1268e9b5e76fb-Paper.pdf.
- Shafer, G. and Vovk, V. A tutorial on conformal prediction. *J. Mach. Learn. Res.*, 9:371–421, jun 2008. ISSN 1532-4435.
- Tibshirani, R. J., Foygel Barber, R., Candes, E., and Ramdas, A. Conformal prediction under covariate shift. In Wallach, H., Larochelle, H., Beygelzimer, A., d’Alché-Buc, F., Fox, E., and Garnett, R. (eds.), *Advances in Neural Information Processing Systems*, volume 32. Curran Associates, Inc., 2019. URL https://proceedings.neurips.cc/paper_files/paper/2019/file/8fb21ee7a2207526da55a679f0332de2-Paper.pdf.
- Vovk, V., Gammerman, A., and Shafer, G. *Algorithmic learning in a random world*, volume 29. Springer, 2005.
- Wisniewski, W., Lindsay, D., and Lindsay, S. Application of conformal prediction interval estimations to market makers’ net positions. In *Conformal and probabilistic prediction and applications*, pp. 285–301. PMLR, 2020.
- Wolfer, G. and Kontorovich, A. Estimating the mixing time of ergodic markov chains. In *Conference on Learning Theory*, pp. 3120–3159. PMLR, 2019.
- Yu, B. Rates of convergence for empirical processes of stationary mixing sequences. *The Annals of Probability*, pp. 94–116, 1994.
- Zaffran, M., Féron, O., Goude, Y., Josse, J., and Dieuleveut, A. Adaptive conformal predictions for time series. In *International Conference on Machine Learning*, pp. 25834–25866. PMLR, 2022.

Appendix

A. Proofs

A.1. Notation

We will often make use of the following definition of the total variation between two measures P, Q defined on a measurable space $(\mathcal{X}, \mathcal{F})$

$$\|P - Q\|_{TV} = \sup_{A \in \mathcal{F}} |P(A) - Q(A)| \quad (10)$$

We will also make use of the equivalent formulation of

$$\beta(a) = \sup_{t \geq 1} \|\mathbb{P}_1^t \otimes \mathbb{P}_{t+a}^\infty - \mathbb{P}_{t,a}\|_{TV}$$

for $a \geq 1$, where $\mathbb{P}_{t,a}$ is the joint distribution of (X_1^t, X_{t+a}^∞) (Doukhan, 2012).

It can be shown (Gallegos-Herrada et al., 2023) that if a Markov chain is geometrically ergodic, then for any measure $\nu \in L^p(\pi) = \{\mu, \int |\frac{d\mu}{d\pi}|^p d\pi < \infty\}$ for some $p \geq 1$, there exists $0 \leq \rho < 1$ (that depends on ν through p only) and a constant C_ν such that for all $n \geq 1$, $\|\nu P^n - \pi\|_{TV} \leq C_\nu \rho^n$. Without loss of generality, we will identify for conciseness the constants C, C_ν (resp C') and ρ appearing in the total variation term $\|\nu P^n - \pi\|_{TV}$ and in $\beta(a)$ (resp $\beta'(a)$) to be the same (consider the maximum of each value for example).

We define the following probability notations: (i) the probability over the calibration dataset given the training dataset, $\mathbb{P}_{\text{cal}} = \mathbb{P}_{\text{cal}, (X_{n+1}, Y_{n+1})}(\cdot | D_{\text{tr}})$, (ii) the probability \mathbb{P}_π when the test point (X_{n+1}, Y_{n+1}) is taken independently of the chain and follows the stationary distribution π .

A.2. Proof of Proposition 4.1 and Theorem 4.1

Proof. The proof is divided into 4 parts:

- Calculation of the marginal coverage when the training and calibration datasets are independent
- Calculation of the marginal coverage in the general case
- Optimisation of the coverage gap w.r.t u
- Optimisation w.r.t r

Define the random variable $Z = (X, Y)$ and let $Z_{1-N}^{n+1} = ((X_i, Y_i))_{i=1-N}^{n+1}$

Calculation of the marginal coverage when the training and calibration datasets are independent

Suppose in this section only that Z_{1-N}^0 is fixed therefore S is fixed and furthermore suppose that Z_{1-N}^0 is independent of Z_1^n

Similar to (Foygel Barber et al., 2021; Oliveira et al., 2022), we show with high probability over Z_1^n that $\hat{q}_{\alpha, n}$ is lower bounded by a true quantile \hat{q}_{α^+} of the scores $S(X, Y)$ when $(X, Y) \sim \pi$, and where $\alpha^+ = \alpha + u > \alpha$ for $1 - \alpha \geq u > 0$.

Define the function f

$$f(Z_1^n) = \sum_{i=1}^n \mathbb{1}\{S_i > \hat{q}_{\alpha^+}\}$$

Then by definition on the calibration dataset

$$\begin{aligned}
 \mathbb{E}_{\text{cal}} f(Z_1^n) &= \sum_{i=1}^n \mathbb{P}_{\text{cal}}(S_i > \hat{q}_{\alpha^+}) \\
 &= \sum_{i=1}^n \mathbb{P}_{\text{cal}}(S_i > \hat{q}_{\alpha^+}) - \mathbb{P}_{\pi}(S > \hat{q}_{\alpha^+}) + \mathbb{P}_{\pi}(S > \hat{q}_{\alpha^+}) \\
 &\geq \sum_{i=1}^n [\mathbb{P}_{\pi}(S > \hat{q}_{\alpha^+}) - \|\nu_0 P^{N+i} - \pi\|_{TV}] \\
 &= n\alpha^+ - \frac{1}{n} \sum_{i=1}^n \|\nu_0 P^{N+i} - \pi\|_{TV} \\
 &\geq n\alpha^+ - \delta(N)
 \end{aligned} \tag{11}$$

By (Paulin, 2015), we have the following McDiarmid concentration

$$\forall t > 0, \quad \mathbb{P}_{\text{cal}}\left(\frac{1}{n}(f(Z_1^n) - \mathbb{E}_{\text{cal}} f(Z_1^n)) > -t\right) \geq 1 - \exp\left(\frac{-2n^2 t^2}{\|c\|^2 \tau_{\min}}\right)$$

where

$$\tau_{\min} = \inf_{0 \leq \varepsilon < 1} \tau(\varepsilon) \left(\frac{2-\varepsilon}{1-\varepsilon}\right)^2 \quad \text{and} \quad \tau_{\min} \leq 9t_{\text{mix}}$$

and $c \in \mathbb{R}^{n+1} = (1, \dots, 1)^T$, $\|c\|^2 = n$ is a vector chosen such that for any (Z, Z')

$$f(Z_1^n) - f(Z_1'^n) \leq \sum_{i=1}^T c_i \mathbb{1}\{Z_i \neq Z'_i\}$$

From (11), we have

$$\begin{aligned}
 \frac{1}{n}(f(Z_1^n) - \mathbb{E}_{\text{cal}} f(Z_1^n)) \geq -t &\implies \frac{1}{n} \sum_{i=1}^n \mathbb{1}\{S_i > \hat{q}_{\alpha^+}\} \geq \alpha^+ - \delta(N) - t \\
 &\implies \sum_{i=1}^n \mathbb{1}\{S_i > \hat{q}_{\alpha^+}\} \geq n\alpha + n(u - \delta(N) - t) \\
 &\implies \sum_{i=1}^n \mathbb{1}\{S_i > \hat{q}_{\alpha^+}\} \geq n\alpha \quad \text{for } t = u - \delta(N)
 \end{aligned}$$

By definition of $\hat{q}_{\alpha, n}$, this inequality implies that $\hat{q}_{\alpha^+} \leq \hat{q}_{\alpha, n}$.

For any $1 \leq r \leq n$ and $u > \delta(N)$, let $E_r(u) = \{f(Z_r^{n+1}) \geq \mathbb{E}_{\text{cal}}(f(Z_r^{n+1})) - (u - \delta(N))\}$. Therefore, for any $1 \leq r \leq n$, we proved that under the hypothesis that Z_{1-N}^0 is independent of Z_r^{n+1} , then we have over both the training and (reduced) calibration datasets

$$\mathbb{P}(\hat{q}_{\alpha^+} \leq \hat{q}_{\alpha, (n-r)}) \geq \mathbb{P}(E_r(u)) \geq 1 - \exp\left(\frac{-2(n-r)(u - \delta(N))^2}{9t_{\text{mix}}}\right) \tag{12}$$

where $\hat{q}_{\alpha, (n-r)}$ is the empirical quantile over $\{S(Z_r^n)\}$.

Finally, we relate $\hat{q}_{\alpha, (n-r)}$ to $\hat{q}_{\alpha', n}$ with α' well chosen.

For any dataset D and point $Z \in D$, let $\text{rank}(D, Z)$ be the rank of Z in D . Then by definition of the empirical quantile, $\text{rank}(D_{\text{cal},r}, \hat{q}_{\alpha,(n-r)}) = \lceil (n+r-1)(1-\alpha) \rceil$. When we add the remaining S_1, \dots, S_r then $\text{rank}(D_{\text{cal}}, \hat{q}_{\alpha,(n-r)}) \in [\lceil (n+r-1)(1-\alpha) \rceil - r, \lceil (n+r-1)(1-\alpha) \rceil + r]$. The extremities of the interval correspond to the cases where either $\min(S_1, \dots, S_r) \geq \max(S_{r+1}, \dots, S_n)$ or $\max(S_1, \dots, S_r) \leq \min(S_{r+1}, \dots, S_n)$.

Hence, we directly obtain that $\hat{q}_{\alpha + \frac{1+r\alpha}{n+1}, n} \leq \hat{q}_{\alpha, (n-r)} \leq \hat{q}_{\alpha - \frac{1+r\alpha}{n+1}, n}$ and equivalently $\hat{q}_{\alpha + \frac{1+r\alpha}{n+1}, (n-r)} \leq \hat{q}_{\alpha, n} \leq \hat{q}_{\alpha - \frac{1+r\alpha}{n+1}, (n-r)}$

Calculation of the marginal coverage in the general case

We now consider the probability over the entire chain, therefore the training and calibration datasets are not independent. Similar to (Oliveira et al., 2022), for $r \geq 1$, we consider the reduced chain Z_r^n . Let \mathbb{P}_r^* be the joint distribution under which Z_{1-N}^0 and Z_r^n are supposed independent. Let \mathbb{P}_{n+1}^* be the joint distribution under which Z_{1-N}^0 and Z_{n+1} are supposed independent.

For any $u > \delta(N)$, for any $r \in [1, n]$, we have

$$\begin{aligned}
 & \mathbb{P}(S_{n+1} \leq \hat{q}_{\alpha, n}) \\
 & \geq \mathbb{P}(S_{n+1} \leq \hat{q}_{\alpha', (n-r)}) \quad \text{where } \alpha' = \alpha + \frac{1+r\alpha}{n+1} \\
 & \geq \mathbb{P}(S_{n+1} \leq \hat{q}_{\alpha', (n-r)}, Z_r^n \in E_r(u)) \\
 & \geq \mathbb{P}(S_{n+1} \leq \hat{q}_{\alpha'+u}, Z_r^n \in E_r(u)) \quad \text{since on } E_r(u), \hat{q}_{\alpha'+u} \leq \hat{q}_{\alpha', n-r} \text{ and the c.d.f is non decreasing} \\
 & \geq \mathbb{P}(S_{n+1} \leq \hat{q}_{\alpha'+u}) - \mathbb{P}(Z_r^n \in E_r(u)) \\
 & \geq \mathbb{P}_{n+1}^*(S_{n+1} \leq \hat{q}_{\alpha'+u}) - \beta(n+1) - \mathbb{P}_r^*(Z_r^n \in E_r(u)) - \beta(r) \\
 & \geq \mathbb{P}_\pi(S \leq \hat{q}_{\alpha'+u}) + [\mathbb{P}_{n+1}^*(S_{n+1} \leq \hat{q}_{\alpha'+u}) - \mathbb{P}_\pi(S \leq \hat{q}_{\alpha'+u})] - \beta(n+1) \\
 & \quad - \exp\left(\frac{-2(n-r)(u - \delta(N))^2}{9t_{\text{mix}}}\right) - \beta(r) \\
 & \geq 1 - \alpha - u - \frac{1+r\alpha}{n+1} - \delta(n+N+1) - \exp\left(\frac{-2(n-r)(u - \delta(N))^2}{9t_{\text{mix}}}\right) - \beta(n+1) - \beta(r)
 \end{aligned}$$

where the fifth inequality comes from

$$\begin{aligned}
 \mathbb{P}(S_{n+1} \leq \hat{q}_{\alpha'+u}) &= \mathbb{P}(S_{n+1} \leq \hat{q}_{\alpha'+u}) - \mathbb{P}_{n+1}^*(S_{n+1} \leq \hat{q}_{\alpha'+u}) + \mathbb{P}_{n+1}^*(S_{n+1} \leq \hat{q}_{\alpha'+u}) \\
 &\geq \mathbb{P}_{n+1}^*(S_{n+1} \leq \hat{q}_{\alpha'+u}) - \|\mathbb{P}_{0, n+1} - \mathbb{P}_{1-N}^0 \otimes \mathbb{P}_{n+1}^\infty\|_{TV} \\
 &\geq \mathbb{P}_{n+1}^*(S_{n+1} \leq \hat{q}_{\alpha'+u}) - \sup_t \|\mathbb{P}_{t, n+1} - \mathbb{P}_{1-N}^t \otimes \mathbb{P}_{t+n+1}^\infty\|_{TV} \\
 &= \mathbb{P}_{n+1}^*(S_{n+1} \leq \hat{q}_{\alpha'+u}) - \beta(n+1)
 \end{aligned}$$

and similarly for $\mathbb{P}(Z_r^n \in E_r(u))$.

Notice that in the case of a restart of the chain at the calibration with initial distribution ν_1 , we can choose $r = 0$ and ignore the β terms therefore obtaining

$$\gamma(u) = u + \exp\left(\frac{-2n(u - \frac{1}{n} \sum_{i=1}^n \delta_1(i))^2}{9t_{\text{mix}}}\right) + \delta_1(n+1) \quad (13)$$

If X is ϕ -irreducible aperiodic Markov chain with coefficients $\beta(a)$, we have

$$\gamma(u, r) = u + \exp\left(\frac{-2(n-r)(u - \delta(N))^2}{9t_{\text{mix}}}\right) + \delta(n+N+1) + \beta(n+1) + \beta(r) \quad (14)$$

If X is a ϕ -irreducible aperiodic geometric ergodic Markov chain,

$$\gamma(u, r) = u + \frac{1 + r\alpha}{n + 1} + \exp\left(\frac{-2(n - r)(u - C\rho^N)^2}{9t_{\text{mix}}}\right) + C(\rho^r + \rho^{n+1} + \rho^{n+N+1}) \quad (15)$$

Upper bound

We keep the same notations but replace α^+ by $\alpha^- = \alpha + \frac{\alpha-1}{n} - u < \alpha$ where the additional $\frac{\alpha-1}{n}$ takes into account the correction in the empirical quantile calculation.

Again, supposing that Z_{1-N}^0 is independent of Z_1^{n+1} then, we have

$$\forall t > 0, \quad \mathbb{P}_{\text{cal}}\left(\frac{1}{n}(f(Z) - \mathbb{E}_{\text{cal}}f(Z)) \leq t\right) \geq 1 - \exp\left(\frac{-2nt^2}{9\tau_{\text{min}}}\right)$$

and we also have

$$\begin{aligned} \frac{1}{n}(f(Z) - \mathbb{E}_{\text{cal}}f(Z)) \leq t &\implies \frac{1}{n} \sum_{i=1}^n \mathbb{1}\{S_i > \hat{q}_{\alpha^-}\} \leq \alpha^- + \delta(N) + t \\ &\implies \sum_{i=1}^n \mathbb{1}\{S_i > \hat{q}_{\alpha^-}\} \leq (n+1)\alpha - 1 + n(t - (u - \delta(N))) \\ &\implies n - \sum_{i=1}^n \mathbb{1}\{S_i \leq \hat{q}_{\alpha^-}\} \leq (n+1)\alpha - 1 \quad \text{for } t = u - \delta(N) \\ &\implies \sum_{i=1}^n \mathbb{1}\{S_i \leq \hat{q}_{\alpha^-}\} > (n+1)(1 - \alpha) \geq \lceil (n+1)(1 - \alpha) \rceil \quad \text{since the l.h.s is an integer} \end{aligned}$$

By definition of $\hat{q}_{\alpha, n}$, this inequality implies that $\hat{q}_{\alpha^-} \geq \hat{q}_{\alpha, n}$.

For any $1 \leq r \leq n$ and $u > \delta(N)$, let $E_r(u) = \{f(Z_r^{n+1}) \leq \mathbb{E}_{\text{cal}}(f(Z_r^{n+1})) + (u - \delta(N))\}$.

We have

$$\begin{aligned} \mathbb{P}(S_{n+1} \leq \hat{q}_{\alpha, n}) &\leq \mathbb{P}(S_{n+1} \leq \hat{q}_{\alpha', (n-r)}) \quad \text{where } \alpha' = \alpha - \frac{1 + r\alpha}{n + 1} \\ &= \mathbb{P}(S_{n+1} \leq \hat{q}_{\alpha', (n-r)}, Z_r^n \in E_r(u)) + \mathbb{P}(S_{n+1} \leq \hat{q}_{\alpha', (n-r)}, Z_r^n \in E_r(u)^c) \\ &\leq \mathbb{P}(S_{n+1} \leq \hat{q}_{\alpha' - u}, Z_r^n \in E_r(u)) + \mathbb{P}(Z_r^n \in E_r(u)^c) \\ &\leq \mathbb{P}(S_{n+1} \leq \hat{q}_{\alpha' - u}) + \mathbb{P}(Z_r^n \in E_r(u)^c) \\ &\leq \mathbb{P}_{n+1}^*(S_{n+1} \leq \hat{q}_{\alpha' - u}) + \beta(n+1) + \mathbb{P}_r^*(Z_r^n \in E_r(u)) + \beta(r) \\ &\leq \dots \end{aligned}$$

Hence, we obtain the same value of γ with an additional $\frac{1-\alpha}{n-r}$ for all cases as in 13,14 and 15.

Optimisation of $\gamma(u, r)$ w.r.t to u

We study the behaviour of γ as a function of u .

For simplicity, let $\gamma(u) = u + \exp(-A(u - K)^2) + B$ where

$$A = \frac{2(n-r)}{9t_{\text{mix}}} \quad K = \delta(N) \quad B = \frac{1 + r\alpha}{n + 1} + \delta(n + N + 1) + \beta(n + 1) + \beta(r)$$

We study the behaviour of the shifted function $u \rightarrow \gamma(u + K)$.

By writing the first derivative, we have for $u \geq \delta(N)$

$$\gamma'(u) = 1 - 2Au \exp(-Au^2)$$

Hence by equivalence (as $u > 0$)

$$\begin{aligned} \gamma'(u) = 0 &\iff 1 = 2Au \exp(-Au^2) \\ &\iff 1 = 4A^2 u^2 \exp(-2u^2) \\ &\iff -\frac{1}{2A} = -2Au^2 \exp(-2Au^2) \end{aligned}$$

Let $w = -\frac{1}{2A}$ and $z = -2Au^2$ then we need to solve the following Lambert equation $ze^z = w$ with unknown z .

If $w < \frac{-1}{e}$ i.e $n - r < \frac{9e}{4} t_{\text{mix}} \equiv en_0$ then this equation has no solution as $\forall u, \quad ue^u > \frac{-1}{e}$. And by ascending the previous calculus by equivalence, this implies $\gamma'(u) > 0$.

If $n - r \geq en_0$ then there exists two critical points $z_1 = W_0(w)$ and $z_2 = W_{-1}(w)$, which in turn gives

$$u_1 = \sqrt{-W_0\left(-\frac{1}{2A}\right) \frac{1}{2A}} \quad u_2 = \sqrt{-W_{-1}\left(-\frac{1}{2A}\right) \frac{1}{2A}} \quad u_1 \leq u_2$$

To know the ordering between both, we need to study the second order derivative.

$$\gamma''(u) = 2C \exp(-Au^2)(2Au^2 - 1)$$

Hence, $\forall u < \sqrt{\frac{1}{2A}} = u_0, \quad \gamma''(u) < 0$ and $\forall u \geq u_0, \quad \gamma''(u) \geq 0$.

This implies that γ' is decreasing on $[0, u_0]$ and increasing on $[u_0, 1]$. Since $\gamma'(0) = 1, \lim_{\infty} \gamma' = 1$, and $\gamma'(u_0) < 0$ (which naturally comes from the condition $n - r \geq n_0$), then γ' is positive on $[0, u_1] \cup [u_2, \infty]$, and negative on $[u_1, u_2]$.

This can be summed up in the following variation table.

u	0	u_1	u_0	u_2	$+\infty$
$\gamma''(u)$	-	-	0	+	+
$\gamma'(u)$	1	0	< 0	0	1
$\gamma(u)$	+	0	-	-	0
$\gamma(u)$	1	$\gamma(u_1)$	$\gamma(u_2)$	$\gamma(u_2)$	∞

Therefore, we proved that $\gamma(u_2) \leq \gamma(u_1)$. Finally, we add the constants to obtain the optimal expression of $\gamma(u, r)$ which will again be denoted by $\gamma(r)$.

$$\gamma(r) = \sqrt{-W_{-1}\left(-\frac{n_0}{n-r}\right)\frac{n_0}{n-r}} + \exp\left(\frac{1}{2}W_{-1}\left(-\frac{n_0}{n-r}\right)\right) + \frac{1+r\alpha}{n+1} + \delta(N) + \delta(n+N+1) + \beta(n+1) + \beta(r) \quad (16)$$

Remark that in the restart case, the optimisation remains the same and we obtain

$$\gamma = \sqrt{-W_{-1}\left(-\frac{n_0}{n}\right)\frac{n_0}{n}} + \exp\left(\frac{1}{2}W_{-1}\left(-\frac{n_0}{n}\right)\right) + \frac{1}{n} \sum_{i=1}^n \delta_1(i) + \delta_1(n+1) \quad (17)$$

Let us consider a few important points in the optimisation of $\gamma(r)$:

- From 14 or 15 it is clear that for any $u, r \rightarrow \gamma(u, r)$ is non increasing. Therefore it must also hold for the optimal value of $u = u_2$
- Hence $\sqrt{-W_{-1}\left(-\frac{n_0}{n-r}\right)\frac{n_0}{n-r}} + \exp\left(\frac{1}{2}W_{-1}\left(-\frac{n_0}{n-r}\right)\right) \geq \sqrt{-W_{-1}\left(-\frac{n_0}{n}\right)\frac{n_0}{n}} + \exp\left(\frac{1}{2}W_{-1}\left(-\frac{n_0}{n}\right)\right) = \gamma_{\text{as}}$
- For our method to achieve optimal lower bound γ_{as} , we are interested in processes with β coefficient such that for a well chosen $r(n) = o_n(n)$ we have $\beta(r) = \mathcal{O}_n(\gamma_{\text{as}})$ and $\frac{r}{n} = \mathcal{O}_n(\gamma_{\text{as}})$
- Finally, this will yield an optimal asymptotic value (in the geometric ergodic case) $\gamma = \gamma(r(n)) = \mathcal{O}_{n,N}(\gamma_{\text{as}})$

We start by analysing the asymptotic order of γ_{as} . Using the inequality from (Chatzigeorgiou, 2013), we have

$$\forall z > 0 \quad -1 - \sqrt{2z} - z \leq W_{-1}(-e^{-z-1}) \leq -1 - \sqrt{2z} - \frac{2}{3}z$$

Hence, with $z = \ln \frac{n}{n_0}$, we have

$$\gamma_{\min} \leq \gamma_{\text{as}} \leq \gamma_{\max}$$

where

$$\begin{aligned} \gamma_{\min} &= \sqrt{\frac{n_0}{n} \left(1 + \sqrt{2 \ln \frac{n}{n_0}} + \frac{2}{3} \ln \frac{n}{n_0}\right)} + \sqrt{\frac{n_0}{n}} \exp\left(-\sqrt{\frac{1}{2} \ln \frac{n}{n_0}}\right) \\ \gamma_{\max} &= \sqrt{\frac{n_0}{n} \left(1 + \sqrt{2 \ln \frac{n}{n_0}} + \ln \frac{n}{n_0}\right)} + \sqrt{\frac{1}{e} \left(\frac{n_0}{n}\right)^{2/3}} \exp\left(-\sqrt{\frac{1}{2} \ln \frac{n}{n_0}}\right) \end{aligned}$$

It is clear that $\gamma_{\min} = \mathcal{O}_n\left(\sqrt{\frac{n_0}{n} \left(1 + \sqrt{2 \ln \frac{n}{n_0}} + \frac{2}{3} \ln \frac{n}{n_0}\right)}\right) = \mathcal{O}_n\left(\sqrt{\frac{n_0}{n} \ln \frac{n}{n_0}}\right)$

For γ_{\max} , we prove that $\frac{1}{n^{1/3}} \exp\left(-\sqrt{\frac{1}{2} \ln n}\right) < \sqrt{\frac{\ln n}{n}} = \mathcal{O}_n\left(\sqrt{\frac{n_0}{n} \ln \frac{n}{n_0}}\right)$. This is equivalent to proving that for n large

$$h(n) = \frac{1}{3} \ln n - \sqrt{2 \ln n} - \ln \ln n < 0$$

We have for any z large enough

$$h'(z) = \frac{1}{z} \left(\frac{1}{3} - \frac{1}{\ln z} - \frac{1}{\sqrt{2 \ln z}}\right)$$

and the term inside the parenthesis can be seen as a polynomial of degree 2 with unknown $\frac{1}{\sqrt{\ln z}}$. The discriminant is strictly positive hence h' is the sign of the leading coefficient which is -1 here. Therefore h is decreasing and for $h(e) = \frac{1}{3} - \sqrt{2} < 0$ hence h is negative.

This proves that $\gamma_{\max} = \mathcal{O}_n(\sqrt{\frac{n_0}{n} \ln \frac{n}{n_0}})$ and therefore γ_{as} has also the same rate of convergence i.e $\gamma_{\text{as}} = \mathcal{O}_n(\sqrt{\frac{n_0}{n} \ln \frac{n}{n_0}}) = \mathcal{O}_n(\sqrt{t_{\text{mix}} \frac{\ln(n)}{n}})$.

Hence, ignoring the total variation terms, we are looking for stochastic processes such that there exists $r = o_n(n)$ and $\beta(r) = \mathcal{O}_n(\sqrt{t_{\text{mix}} \frac{\ln(n)}{n}})$.

Optimisation of $\gamma(r)$ for geometrically ergodic chains

We show that all conditions are fulfilled for geometric ergodic Markov chains and we can achieve asymptotically optimal rate of convergence γ_{as}

Indeed, since $\beta(r) \leq C\rho^r$, taking $r = \frac{\ln n}{\ln \frac{1}{\rho}}$ gives $r = o_n(n)$, $\frac{r}{n} = \mathcal{O}_n(\gamma_{\text{as}})$, $\beta(r) = \mathcal{O}_n(\frac{1}{n}) = \mathcal{O}_n(\gamma_{\text{as}})$ and the proof is complete. □

A.3. Proof of Proposition 4.2

Proof. Comparison with μ

Let $q_{\alpha,n}$ be the empirical quantile of $\{|Y_i - \mu(X_i)|\}_{i=1}^n$.

Let $E = \{|\hat{\mu}_N - \mu|_\infty \leq c_N\}$ such that $\mathbb{P}(E) \geq 1 - d_N$.

We have

$$|\hat{q}_{\alpha,n} - \hat{q}_\alpha| \leq |\hat{q}_{\alpha,n} - q_{\alpha,n}| + |q_{\alpha,n} - q_\alpha| + |q_\alpha - \hat{q}_\alpha|$$

Following the proof of (Lei et al., 2018), we have for any $D_{\text{tr}} \in E$,

$$|q_\alpha - \hat{q}_\alpha| \leq \frac{d_N}{\kappa} + c_N$$

and for any D_{cal}

$$|\hat{q}_{\alpha,n} - q_{\alpha,n}| \leq c_N$$

Hence for any $u > 2\kappa c_N + d_N$, we have

$$\{|q_{\alpha,n} - q_\alpha| \leq \frac{u}{\kappa} - 2c_N - \frac{d_N}{\kappa}\} \in \{|\hat{q}_{\alpha,n} - \hat{q}_\alpha| \leq \frac{u}{\kappa}\}$$

This gives

$$\begin{aligned} \mathbb{P}(|\hat{q}_{\alpha,n} - \hat{q}_\alpha| \leq \frac{u}{\kappa}) &\geq \mathbb{P}(|\hat{q}_{\alpha,n} - \hat{q}_\alpha| \leq \frac{u}{\kappa}, E) \\ &\geq \mathbb{P}(|q_{\alpha,n} - q_\alpha| \leq \frac{u}{\kappa} - 2c_N - \frac{d_N}{\kappa}, E) \end{aligned}$$

In the restart version, we can directly apply Equation 12 provided that $u' = u - 2\kappa c_N - d_N \in \mathcal{N}$ and obtain that

$$\begin{aligned} \mathbb{P}(|\hat{q}_{\alpha,n} - \hat{q}_\alpha| \leq \frac{u}{\kappa}) &\geq (1 - 2 \exp[-\frac{2n}{9t_{\text{mix}}}(u' - \delta_1(N))^2]) \mathbb{P}_{\text{tr}}(E) \\ &\geq 1 - 2 \exp[-\frac{2n}{9t_{\text{mix}}}(u' - \delta_1(N))^2] - d_N \end{aligned}$$

Application of the concentration inequality over D_{cal} as in Proposition 4.1

We now compute $\mathbb{P}(|q_{\alpha,n} - q_\alpha| \leq \frac{u}{\kappa} - 2c_N - \frac{d_N}{\kappa}, E)$.

Since the concentration inequality was applied on the last $n - r$ samples, we need to replace $|q_{\alpha,n} - q_\alpha|$ by $|q_{\alpha',n-r} - q_\alpha|$ for a certain α' .

For any $r \in [1, n]$, we have

$$q_{\alpha+d\alpha,(n-r)} \leq q_{\alpha,n} \leq q_{\alpha-d\alpha,(n-r)}$$

where $d\alpha = \frac{1+r\alpha}{n+1}$. Hence, we have for $v = \frac{u}{\kappa} - 2c_N - \frac{d_N}{\kappa}$

$$q_\alpha - v \leq q_{\alpha,n} \leq q_\alpha + v \implies q_\alpha + v \geq q_{\alpha+d\alpha,(n-r)} \quad , \quad q_\alpha - v \leq q_{\alpha-d\alpha,(n-r)}$$

Supposing $d\alpha \in \mathcal{N}$, we also have by the bounded hypothesis on the density f ,

$$(1 - \alpha + d\alpha) - (1 - \alpha) = \int_{q_\alpha}^{q_{\alpha-d\alpha}} f \geq (q_{\alpha-d\alpha} - q_\alpha)\kappa$$

hence

$$q_{\alpha-d\alpha} \leq q_\alpha + \frac{d\alpha}{\kappa} \quad \text{and} \quad q_{\alpha+d\alpha} \geq q_\alpha - \frac{d\alpha}{\kappa}$$

therefore

$$q_\alpha + v \geq q_{\alpha+d\alpha,(n-r)} \quad , \quad q_\alpha - v \leq q_{\alpha-d\alpha,(n-r)} \implies q_{\alpha+d\alpha} + \frac{d\alpha}{\kappa} + v \geq q_{\alpha+d\alpha,(n-r)} \quad , \quad q_{\alpha-d\alpha} - \frac{d\alpha}{\kappa} - v \leq q_{\alpha-d\alpha,(n-r)}$$

Finally, we know from Equation 12 that for any $u' > \delta(N)$, we have provided that $u' \leq \alpha + d\alpha, u' \in \mathcal{N}$

$$q_{\alpha+d\alpha,(n-r)} \leq q_{\alpha+d\alpha-u'} \leq q_{\alpha+d\alpha} + \frac{u'}{\kappa}$$

with probability at least $1 - \exp\left(-\frac{2(n-r)(u'-\delta(N))^2}{9t_{\text{mix}}}\right)$ if the (reduced) calibration and training datasets are independent and similarly

$$q_{\alpha-d\alpha,(n-r)} \geq q_{\alpha-d\alpha+u'} \geq q_{\alpha-d\alpha} - \frac{u'}{\kappa}$$

Therefore identifying $\frac{u'}{\kappa} = \frac{d\alpha}{\kappa} + v$ i.e $u' = d\alpha + v\kappa$ we have

$$\begin{aligned}
 & \mathbb{P}(|q_{\alpha,n} - q_{\alpha}| \leq \frac{u}{\kappa} - 2c_N - \frac{d_N}{\kappa}, E) \\
 & \geq \mathbb{P}(q_{\alpha-d\alpha} - q_{\alpha-d\alpha,n-r} \leq \frac{u - 2\kappa c_N - d_N + d\alpha}{\kappa}, q_{\alpha+d\alpha,n-r} - q_{\alpha+d\alpha} \leq \frac{u - 2\kappa c_N - d_N + d\alpha}{\kappa}, E) \\
 & \geq \mathbb{P}_r^*(q_{\alpha-d\alpha} - q_{\alpha-d\alpha,n-r} \leq \frac{u - 2\kappa c_N - d_N + d\alpha}{\kappa}, q_{\alpha+d\alpha,n-r} - q_{\alpha+d\alpha} \leq \frac{u - 2\kappa c_N - d_N + d\alpha}{\kappa}, E) - \beta(r) \\
 & \geq (1 - 2 \exp[-\frac{2(n-r)}{9t_{\text{mix}}}(u' - \delta(N))^2]) \mathbb{P}_{\text{tr}}(E) - \beta(r) \\
 & \geq 1 - 2 \exp[-\frac{2(n-r)}{9t_{\text{mix}}}(u' - \delta(N))^2] - d_N - \beta(r)
 \end{aligned}$$

□

A.4. Proof of Proposition 4.3

The proof in (Lei et al., 2018) consists in proving the two following points:

- $\forall \varepsilon > 0 \quad \lim_{N \rightarrow \infty} \mathbb{P}_{X_{n+1}, \text{tr}}(|\hat{\mu}_N(X_{n+1}) - \mu(X_{n+1})| > \varepsilon) = 0$
- $\forall \varepsilon > 0 \quad \lim_{N \rightarrow \infty} \lim_{n \rightarrow \infty} \mathbb{P}(|\hat{q}_{\alpha,n} - q_{\alpha}| > \varepsilon) = \lim_{n \rightarrow \infty} \lim_{N \rightarrow \infty} \mathbb{P}(|\hat{q}_{\alpha,n} - q_{\alpha}| > \varepsilon) = 0$

We tackle 1) the additional difficulty that N and n are different, 2) the non iid settings.

- For the second point, we know by (Lei et al., 2018) that $|\hat{q}_{\alpha} - q_{\alpha}| \leq (M/2\kappa)\Delta_N$ hence

$$a_{N,n} = \mathbb{P}(|\hat{q}_{\alpha,n} - q_{\alpha}| > \varepsilon) \leq \mathbb{P}(|\hat{q}_{\alpha,n} - \hat{q}_{\alpha}| + (M/2\kappa)\Delta_N > \varepsilon)$$

By the hypothesis that $\Delta_N = o_N(1)$, $c_N = o_N(1)$, $d_N = o_N(1)$,

- Let N_0 such that $\forall N \geq N_0$, $\Delta_N \leq \varepsilon$
- Let $N_1 \geq N_0$ such that $\forall N \geq N_1$, $f(N) = (M/2\kappa)\Delta_N + 2\kappa c_N + d_N \leq \kappa\varepsilon/2$.

Let n_0 such that $\forall n \geq n_0$, $\kappa\varepsilon - \frac{1+\alpha r}{n+1} \geq f(N_1)$. This is justified as the l.h.s tends to $\kappa\varepsilon$ when $n \rightarrow \infty$ therefore it must be larger than $\kappa\varepsilon/2$ after an integer n_0 .

Then by applying Proposition 4.2 we have for any $n \geq n_0$,

$$\mathbb{P}(|\hat{q}_{\alpha,n} - q_{\alpha}| > \varepsilon) \leq 2e^{-\frac{2(n-\ln n)}{9t_{\text{mix}}}(u' - \delta(N_1))^2} + d_{N_0} + \beta(\ln n)$$

where $u' = \kappa\varepsilon - (M/2\kappa)\Delta_{N_1} - 2\kappa c_{N_1} - d_{N_1} - \frac{1+\alpha \ln n}{n+1}$.

Let $N_2 \geq N_1$ such that $\forall N \geq N_2$, $f(N) \leq f(N_1)$

Then $\forall n \geq n_0$ and $\forall N \geq N_2$

$$\mathbb{P}(|\hat{q}_{\alpha,n} - q_{\alpha}| > \varepsilon) \leq 2e^{-\frac{2(n-\ln n)}{9t_{\text{mix}}}(u' - \delta(N))^2} + d_N + \beta(\ln n)$$

where $u' = \kappa\varepsilon - (M/2\kappa)\Delta_N - 2\kappa c_N - d_N - \frac{1+\alpha \ln n}{n+1}$.

Now that we decoupled n and N (since all integers defined beforehand only depend on ε), we can take the limit w.r.t n and N . Since it is a sum of terms that tend to 0 (supposing the chain is β -mixing) then the order does not count and we have

$$\forall \varepsilon > 0 \quad \lim_{N \rightarrow \infty} \lim_{n \rightarrow \infty} \mathbb{P}(|\hat{q}_{\alpha, n} - q_{\alpha}| > \varepsilon) = \lim_{n \rightarrow \infty} \lim_{N \rightarrow \infty} \mathbb{P}(|\hat{q}_{\alpha, n} - q_{\alpha}| > \varepsilon) = 0$$

- To prove the first point, remark that the condition $\mathbb{E}_{X \sim \pi, \text{tr}}[(\hat{\mu}_N(X) - \mu(X))^2] = o_N(1)$ implies that there exists sequences $\eta_N = o(1)$ and $\rho_N = o(1)$ such that

$$\mathbb{P}(\mathbb{E}_{X \sim \pi}[(\hat{\mu}_N(X) - \mu(X))^2 | \hat{\mu}_N] \geq \eta_N) \leq \rho_N \quad (18)$$

Let $I = \{\Delta_N \leq \eta_N\}$ be the event which has probability at least $1 - \rho_N$.

We calculate

$$\begin{aligned} \mathbb{P}(|\hat{\mu}_N(X_{n+1}) - \mu(X_{n+1})| \geq \eta^{1/3}) &\leq \mathbb{P}_{n+1}^*(|\hat{\mu}_N(X_{n+1}) - \mu(X_{n+1})| \geq \eta^{1/3}) + \beta(n+1) \\ &= \mathbb{E}(\mathbb{P}_{n+1}^*(|\hat{\mu}_N(X_{n+1}) - \mu(X_{n+1})| \geq \eta^{1/3} | D_{\text{tr}})) + \beta(n+1) \end{aligned}$$

Therefore by Markov inequality conditioned on I

$$\mathbb{P}_{n+1}^*(|\hat{\mu}_N(X_{n+1}) - \mu(X_{n+1})| \geq \eta_N^{1/3} | D_{\text{tr}}) \leq \frac{\mathbb{E}_{X_{n+1}^*}(|\hat{\mu}_N(X_{n+1}^*) - \mu(X_{n+1}^*)|^2)}{\eta_N^{2/3}}$$

If $X_{n+1} \sim \pi$ then the r.h.s becomes $\eta_N^{1/3}$ and under β mixing assumption, we prove the first point as we have

$$\mathbb{P}(|\hat{\mu}_N(X_{n+1}) - \mu(X_{n+1})| \geq \eta^{1/3}) \leq \eta_N^{1/3} + \rho_N + \beta(n+1)$$

and the r.h.s tends to 0 independently of the order for n and N .

Otherwise under \mathbb{P}^* , $X_{n+1} \sim \nu_0 P^{n+N+1}$ (or $\nu_1 P^n$ in the restart version)

Suppose that we have furthermore Assumption 2 and let $I = \{\Delta_N \leq \eta_N, \|\hat{\mu}_N - \mu\|_{\infty} \leq c_N\}$ be the event which has probability at least $1 - \rho_N - d_N$. From the same calculus as above, we bound the r.h.s

$$\begin{aligned} \mathbb{E}_{X_{n+1}^*}(|\hat{\mu}_N(X_{n+1}^*) - \mu(X_{n+1}^*)|^2) &= \int_x (\hat{\mu}_N(x) - \mu(x))^2 \nu_0 P^{n+N+1}(dx) \\ &\leq c_N^2 \end{aligned}$$

Hence, we obtain

$$\mathbb{P}(|\hat{\mu}_N(X_{n+1}) - \mu(X_{n+1})| \geq \eta_N^{1/3}) \leq \frac{c_N^2}{\eta_N^{2/3}} + \rho_N + d_N + \beta(n+1)$$

Furthermore, remark that Assumption 2 also implies Equation 18 as

$$\mathbb{E}_{X \sim \pi}[(\hat{\mu}_N(X) - \mu(X))^2 | \hat{\mu}_N] \leq \|\hat{\mu} - \mu\|_{\infty} \leq c_N^2$$

where the last inequality holds with probability at least $1 - d_N$ therefore we can take $\eta_N = c_N^2$ and $\rho_N = d_N$. Hence, we can conclude again by taking the limit, independently of the order, when $n, N \rightarrow \infty$.

A.5. Proof of Proposition 5.1

Proof. For simplicity, we suppose that K divides n (otherwise the difference is asymptotically negligible). For any r, s , let $Z_r^s = \{(X_i, Y_i)\}_{i=r}^s$.

Stationary case:

We describe here the original blocking technique for *stationary* β -mixing sequences with fixed block size.

- Divide a sample $Z = (Z_1, \dots, Z_n)$ into $2m$ blocks B_1, \dots, B_{2m} . Even blocks are of size a and odd blocks are of size b i.e $n = 2m(a + b)$ (suppose n is even as there are minor changes otherwise)
- Denote by $B_{\text{odd}} = (B_1, B_3, \dots, B_{2m-1})$ the list of odd blocks and define its independent version $B_{\text{odd}}^* = (B_1^*, B_3^*, \dots, B_{2m-1}^*)$ where

$$B_{2k-1}^* \stackrel{d}{=} B_{2k-1} \quad B_1^* \perp B_3^* \dots$$

Then we have the following theorem

Theorem A.1 (Blocking technique (Yu, 1994)). *For all bounded measurable functions $|h| \leq M$*

$$|\mathbb{E}(h(B_{\text{odd}})) - \mathbb{E}(h(B_{\text{odd}}^*))| \leq 2Mm\beta(a)$$

1) Suppose in the following only that Z_{1-N}^0 and Z_1^n are independent and let Z_{1-N}^0 fixed such that S is well defined and constant.

Let $B_{\text{odd}} = \{Z_{iK+1}\}_{i=0}^{\frac{n}{K}} = \{Z_{iK+1}\}_{i=0}^{\frac{n}{K}-1} \cup Z_{n+1}$ and $h(B_{\text{odd}}) = 1\{S_{n+1} \leq \hat{q}_{\alpha, K, n}\}$.

Conditionally on Z_{1-N}^0 , Z_1^n still follows the stationary distribution as they are independent. Therefore applying Theorem A.1, we obtain

$$|\mathbb{P}_{\text{cal}, Z_{n+1}}(S_{n+1} \leq \hat{q}_{\alpha, K, n}) - \mathbb{P}_{\text{cal}^*, Z_{n+1}^*}(S_{n+1} \leq \hat{q}_{\alpha, K, n})| \leq 2\frac{n}{K}\beta_{\text{cal}}(K)$$

where under $\mathbb{P}_{\text{cal}^*, Z_{n+1}^*}$, the samples $\{Z_{iK+1}\}_{i=0, \dots, \frac{n}{K}}$ are i.i.d. Since the calibration dataset is independent of the training dataset then $\beta_{\text{cal}}(K)$ can be seen simply as the original $\beta(K)$ coefficient but on a smaller chain. Therefore, $\beta_{\text{cal}}(K) \leq \beta(K)$.

It is well known that

$$1 - \alpha \leq \mathbb{P}_{\text{cal}^*, Z_{n+1}^*}(S_{n+1} \leq \hat{q}_{\alpha, K, n}) \leq 1 - \alpha + \frac{1}{1 + \frac{n}{K}}$$

From there, we can conclude for the restart part in the stationary case.

2) Suppose now that Z_{1-N}^0 and Z_1^n are no longer independent. Denote by $\mathbb{P}_{r, \text{cal}, Z_{n+1}}^*$ the conditional probability under which Z_r^{n+1} is independent of Z_{1-N}^0 . Similarly, denote by $\mathbb{P}_{r, \text{cal}^*, Z_{n+1}^*}^*$ the same probability and with $\{Z_i\}_{i=r}^{n+1}$ i.i.d.

$$\begin{aligned} \mathbb{P}(S_{n+1} \leq \hat{q}_{\alpha, K, n}) &\geq \mathbb{P}(S_{n+1} \leq \hat{q}_{\alpha', K, n-r}) \quad \text{where } \alpha' = \alpha + \frac{1 + \alpha r}{n + 1} \\ &\geq \mathbb{P}_r^*(S_{n+1} \leq \hat{q}_{\alpha', K, n-r}) - \beta(r) \\ &= \mathbb{E}[\mathbb{P}_{r, \text{cal}, Z_{n+1}}^*(S_{n+1} \leq \hat{q}_{\alpha', K, n-r})] - \beta(r) \\ &\geq \mathbb{E}[\mathbb{P}_{r, \text{cal}^*, Z_{n+1}^*}^*(S_{n+1} \leq \hat{q}_{\alpha', K, n-r})] - 2\frac{n-r}{K}\beta(K) - \beta(r) \\ &= 1 - \alpha - \frac{1 + \alpha r}{n + 1} - 2\frac{n-r}{K}\beta(K) - \beta(r) \end{aligned}$$

and similarly,

$$\mathbb{P}(S_{n+1} \leq \hat{q}_{\alpha, K, n-r}) \leq 1 - \alpha + \frac{K}{n-r} + \frac{1 + \alpha r}{n+1} + 2 \frac{n-r}{K} \beta(K) + \beta(r)$$

Non stationary case

1) Suppose again in the following only that Z_{1-N}^0 and Z_1^n are independent and let Z_1^n fixed such that S is well defined and constant. However, Z_1^n is not necessarily stationary anymore as for all $i \in [1, n]$, $\mathbb{P}_i = \mathbb{P}(Z_i) = \nu_0 P^{N+i}$.

We will use an extended result of the original "Blocking Technique" for *non stationary* β -mixing sequences which can be found in (Kuznetsov & Mohri, 2017). More precisely, suppose $n = ma$ for $a \geq 1, m \geq 1$, and define for $j = 0, \dots, a-1$, the list $B_{\text{odd}}^j = (Z_{1+j}, Z_{1+a+j}, \dots, Z_{1+(m-1)a+j})$ where each sample are separated by a . Then we have the following proposition

Proposition A.1 (Proposition 2 of (Kuznetsov & Mohri, 2017)). *For all bounded measurable functions $|h| \leq M$ we have*

$$|\mathbb{E}(h(B_{\text{odd}}^j)) - \mathbb{E}(h(B_{\pi}^*))| \leq 2Mm\beta'(a)$$

where $\beta'(a) = \sup_{t \geq 1} \mathbb{E}(\|P^{t+a}(\cdot | (Z_1, \dots, Z_t)) - \pi\|_{TV})$ and B_{π}^* is an i.i.d sample of size m which follows the stationary distribution π .

Let $B_{\text{odd}} = \{Z_{iK+1}\}_{i=0}^{\frac{n}{K}} = \{Z_{iK+1}\}_{i=0}^{\frac{n}{K}-1} \cup Z_{n+1}$ and $h(B_{\text{odd}}) = 1\{S_{n+1} \leq \hat{q}_{\alpha, K, n}\}$. Then applying Proposition A.1 to B_{odd} , we obtain

$$|\mathbb{P}_{\text{cal}, Z_{n+1}}(S_{n+1} \leq \hat{q}_{\alpha, K, n}) - \mathbb{P}_{\text{cal}^*, Z_{n+1}^*}(S_{n+1} \leq \hat{q}_{\alpha, K, n})| \leq 2 \frac{n}{K} \beta'_1(K)$$

where $\forall i, 0 \leq i \leq \frac{n}{K}$, Z_{iK+1}^* are independent and follow the stationary distribution π .

Therefore, again we fall back under the i.i.d settings and it is well-known that

$$1 - \alpha \leq \mathbb{P}_{\text{cal}^*, Z_{n+1}^*}(S_{n+1} \leq \hat{q}_{\alpha, K, n}) \leq 1 - \alpha + \frac{1}{1 + \frac{n}{K}}$$

We can conclude for the restart version in the non stationary case.

2) Suppose now that Z_{1-N}^0 and Z_1^n are no longer independent then

$$\begin{aligned} \mathbb{P}(S_{n+1} \leq \hat{q}_{\alpha, K, n}) &\geq \mathbb{P}(S_{n+1} \leq \hat{q}_{\alpha', K, n-r}) \quad \text{where } \alpha' = \alpha + \frac{1 + \alpha r}{n+1} \\ &\geq \mathbb{P}_r^*(S_{n+1} \leq \hat{q}_{\alpha', K, n-r}) - \beta(r) \\ &= \mathbb{E}[\mathbb{P}_{r, \text{cal}, Z_{n+1}}^*(S_{n+1} \leq \hat{q}_{\alpha', K, n-r})] - \beta(r) \\ &\geq \mathbb{E}[\mathbb{P}_{r, \text{cal}^*, Z_{n+1}^*}^*(S_{n+1} \leq \hat{q}_{\alpha, K, n-r})] - 2 \frac{n-r}{K} \beta'_r(K) - \beta(r) \\ &= 1 - \alpha - \frac{1 + \alpha r}{n+1} - 2 \frac{n-r}{K} \beta'_r(K) - \beta(r) \end{aligned}$$

and similarly,

$$\mathbb{P}(S_{n+1} \leq \hat{q}_{\alpha, K, n}) \leq 1 - \alpha + \frac{K}{n-r} + \frac{1 + \alpha r}{n+1} + 2 \frac{n-r}{K} \beta'_r(K) + \beta(r)$$

□

A.6. Proof of Theorem 5.1

We focus our analysis on the "without restart" version as the "restart" version is simply an application of the former.

We suppose $K \gg 1$ (which will be verified further on) and we wish to minimize the total coverage gap as a function of K (i.e without the terms depending on r only)

$$\gamma_K = \frac{K}{n-r} + a \frac{n-r}{K} \rho^K$$

(replace ρ by $\sqrt{\rho}$ in the non stationary case)

For $u > 0$, let $f(u) = \frac{u}{n-r} + a \frac{n-r}{u} \rho^u$ then

$$\begin{aligned} f'(u) &= 0 \\ \iff \frac{1}{n-r} + a \frac{n-r}{u} e^{u \ln \rho} (\ln \rho - \frac{1}{u}) &= 0 \\ \iff \frac{u^2}{n-r} + a(n-r) e^{u \ln \rho} (u \ln \rho - 1) &= 0 \end{aligned}$$

Suppose furthermore that the optimal K (and therefore u) verifies $|K \ln \rho| \gg 1$ then the last equation is asymptotically equal to

$$\begin{aligned} \frac{u^2}{n-r} + a(n-r) u \ln(\rho) e^{u \ln \rho} &= 0 \\ \iff -u \ln(\rho) e^{-u \ln \rho} &= a(n-r)^2 (\ln \rho)^2 \\ \iff u = \frac{W_0(a(n-r)^2 (\ln \rho)^2)}{\ln \frac{1}{\rho}} &= \mathcal{O}_n\left(\frac{\ln(a(n-r)^2 (\ln \rho)^2)}{\ln \frac{1}{\rho}}\right) = \mathcal{O}_n\left(\frac{\ln(n-r)}{\ln \frac{1}{\rho}}\right) \end{aligned}$$

Let us again consider a few important points when choosing the optimal value of r :

- The coverage gap is an increasing function of r
- When $r = 0$, $K = \frac{W_0(a n^2 (\ln \rho)^2)}{\ln \frac{1}{\rho}}$ (as it indeed verifies $K \gg 1$ and $|K \ln \rho| \gg 1$) therefore

$$\frac{K}{n} = \mathcal{O}_n\left(\frac{\ln n}{n \ln \frac{1}{\rho}}\right)$$

and

$$\begin{aligned} a \frac{n}{K} \rho^K &= \mathcal{O}_n\left(a \frac{n \ln \frac{1}{\rho}}{\ln n} \exp\left[\frac{\ln(a n^2 (\ln \rho)^2)}{\ln \frac{1}{\rho}} \ln \rho\right]\right) \\ &= \mathcal{O}_n\left(a \frac{n \ln \frac{1}{\rho}}{\ln n} \frac{1}{a n^2 (\ln \rho)^2}\right) \\ &= \mathcal{O}_n\left(\frac{1}{n \ln(n) \ln(\frac{1}{\rho})}\right) \\ &= \mathcal{O}_n\left(\frac{\ln n}{n \ln \frac{1}{\rho}}\right) \end{aligned}$$

which finally gives $\gamma_K = \mathcal{O}_n\left(\frac{K}{n}\right) = \mathcal{O}_n\left(\frac{\ln n}{n \ln \frac{1}{\rho}}\right)$

From this calculus, we can already conclude in the restart case

$$\gamma_{\text{low}} = \gamma = \mathcal{O}_n\left(\frac{1}{n \ln(n) \ln(\frac{1}{\rho})}\right) \quad \gamma_{\text{up}} = \mathcal{O}_n\left(\frac{\ln n}{n \ln \frac{1}{\rho}}\right)$$

- Let $\gamma_{\text{as}} = \frac{\ln n}{n \ln \frac{1}{\rho}}$. Hence, we are looking for $r \in [n]$ such that asymptotically
 - $r = o(n)$ which will ensure $n - r = \mathcal{O}_n(n)$ and also $K \gg 1$ and $|K \ln \rho| \gg 1$
 - the remaining terms function of r only are also negligible i.e $\frac{1+\alpha r}{n+1} = \mathcal{O}_n(\gamma_{\text{as}})$ and $\beta(r) = \mathcal{O}_n(\gamma_{\text{as}})$

Taking $r = \frac{\ln n}{\ln \frac{1}{\rho}}$ convenes.

Finally, we also achieve an asymptotic optimal coverage gap in the non restart case

$$\gamma_{\text{low}} = \gamma = \mathcal{O}_n\left(\frac{1}{n \ln(n) \ln(\frac{1}{\rho})}\right) \quad \gamma_{\text{up}} = \mathcal{O}_n\left(\frac{\ln n}{n \ln \frac{1}{\rho}}\right)$$

A.7. Proof of Proposition 5.3

Proof. For a given \hat{K}_N , we wish to compute the asymptotic coverage gap $\gamma_{\text{low}}(\hat{K}_N)$ and $\gamma_{\text{up}}(\hat{K}_N)$. We start by bounding the estimation error on \hat{K}_N .

For $\delta \geq 0$ and a universal constant C , let

$$w = w(N, \delta) = C \sqrt{\frac{\ln \frac{|X|}{\delta} \ln N}{\pi^* (1 - \rho) N}}$$

We have with probability at least $1 - \delta$

$$\begin{aligned} \hat{K}_N &\leq -\frac{\ln n}{\ln(\rho + w)} = -\frac{\ln n}{\ln(\rho(1 + \frac{w}{\rho}))} = -\frac{\ln n}{\ln(\rho) + \frac{w}{\rho} + o_N(w)} \\ &= -\frac{\ln n}{\ln(\rho)(1 + \frac{w}{\rho \ln \rho} + o_N(w))} \\ &= -\frac{\ln n}{\ln(\rho)} \left(1 - \frac{w}{\rho \ln \rho} + o_N(w)\right) \\ &= \frac{\ln n}{\ln(\frac{1}{\rho})} + \frac{w \ln n}{\rho (\ln \rho)^2} + o_N(w) \\ &= K + \frac{w \ln n}{\rho (\ln \rho)^2} + o_N(w) \end{aligned}$$

Similarly, $\hat{K}_N \geq K - \frac{w \ln n}{\rho (\ln \rho)^2} + o_N(w)$.

Therefore there exists a universal constant which will again denoted by $C > 0$ such that

$$\mathbb{P}_{\text{tr}}\left(|\hat{K}_N - K| \leq C \frac{w(N, \delta) \ln n}{\rho (\ln \rho)^2}\right) \geq 1 - \delta$$

and for simplicity, we will ignore it.

For the lower bound of the coverage gap, we have with probability at least $1 - \delta$

$$\begin{aligned}
 \gamma_{\text{low}}(\hat{K}_N) &\geq \frac{n}{K + \frac{w \ln n}{\rho(\ln \rho)^2}} \rho^{\hat{K}_N} = \frac{n}{K} \left(1 - \frac{w \ln n}{\rho(\ln \rho)^2 K} + o_N(w)\right) \rho^K \\
 &= \gamma_{\text{low}} - \frac{w \ln n}{\rho(\ln \rho)^2 K} \gamma_{\text{low}} + o_N(w) \\
 &= \gamma_{\text{low}} - \frac{w}{\rho(\ln \rho)^2 n \ln n} + o_N(w)
 \end{aligned}$$

And similarly $\hat{\gamma}_l \leq \gamma_{\text{low}} + \frac{w}{\rho(\ln \rho)^2 n \ln n} + o_N(w)$.

Doing the same calculation for the upper bound of the coverage gap $\gamma_{\text{up}}(\hat{K}_N) = \frac{\hat{K}_N}{n}$, we finally obtain for a universal constant $C > 0$

$$\mathbb{P}_{\text{tr}} \left(\left| \gamma_{\text{low}}(\hat{K}_N) - \gamma_{\text{low}} \right| \leq C \frac{w(N, \delta)}{\rho(\ln \rho)^2 n \ln n}, \left| \gamma_{\text{up}}(\hat{K}_N) - \gamma_{\text{up}} \right| \leq C \frac{w(N, \delta)}{n} \right) \geq 1 - \delta$$

Let $I = \{|\gamma_{\text{low}}(\hat{K}_N) - \gamma_{\text{low}}| \leq C \frac{w(N, \delta)}{\rho(\ln \rho)^2 n \ln n}, |\gamma_{\text{up}}(\hat{K}_N) - \gamma_{\text{up}}| \leq C \frac{w(N, \delta)}{n}\}$, we have $\mathbb{P}_{\text{tr}}(I) \geq 1 - \delta$

By combining both probabilities,

$$\begin{aligned}
 \mathbb{P}(Y_{n+1} \in C_{\hat{K}_N}(X_{n+1})) &\geq \mathbb{P}(Y_{n+1} \in C_{\hat{K}_N}(X_{n+1}), I) \\
 &= \mathbb{E}_{\text{tr}}[\mathbb{P}(Y_{n+1} \in C_{\hat{K}_N}(X_{n+1}) | D_{\text{tr}})] \\
 &\geq \mathbb{E}_{\text{tr}}[1 - \alpha - \gamma_{\text{low}}(\hat{K}_N)] \\
 &\geq \mathbb{E}_{\text{tr}}[1 - \alpha - \gamma_{\text{low}} - C \frac{w(N, \delta)}{\rho(\ln \rho)^2 n \ln n}] \\
 &\geq 1 - \alpha - \gamma_{\text{low}} - C \frac{w(N, \delta)}{\rho(\ln \rho)^2 n \ln n} - \delta
 \end{aligned}$$

and similarly for the upper bound, we have

$$\mathbb{P}(Y_{n+1} \in C_{\hat{K}_N}(X_{n+1})) \leq 1 - \alpha + \gamma_{\text{up}} + C \frac{w(N, \delta)}{n} + \delta$$

Taking $\delta = \frac{|\mathcal{X}| \ln N}{N}$, this gives $w(N, \delta) = \sqrt{\frac{\ln \frac{N}{|\mathcal{X}|} \ln N}{\pi^*(1-\rho)N}} = \mathcal{O}_N\left(\frac{\ln N}{\sqrt{N}}\right)$ □

B. Experiments

B.1. Lazy random walk

For a given $w > 0$, the associated transition matrix is $P : [1, w]^2 \rightarrow [0, 1]$

$$P(x, y) = \begin{cases} \frac{1}{4} & \text{if } y = x + 1 \text{ or } y = x - 1 \pmod{w}, \\ \frac{1}{2} & \text{if } y = x, \\ 0 & \text{otherwise.} \end{cases}$$

We consider a simple linear regression for illustration purposes and take $\mu(x) = ax$ for a certain $a \in [0, 1]$. The independent noise follows $\varepsilon_t \sim N(0, 1)$.

B.2. Gaussian AR(1)

It can be easily shown that the corresponding transition kernel P and all its iterates are also Gaussian

$$P(a, dy) = \frac{1}{\sqrt{2\pi\omega}} \exp\left(-\frac{(y - \theta a)^2}{2\omega^2}\right) dy$$

and

$$\forall n \quad P^n(x, dy) \sim N\left(\theta^n x, \omega^2 \frac{1 - \theta^{2n}}{1 - \theta^2}\right)$$

In the stationary case, we know that β can be written as

$$\beta(n) = \frac{1}{2} \int_{\mathcal{R}} \pi(dx) \|P^n(x, \cdot) - \pi\|_{TV}$$

Furthermore, we also have the following Pinsker inequality on the total variation distance

$$\|P^n(x, \cdot) - \pi\|_{TV} \leq \sqrt{\frac{1}{2} \text{KL}(P^n(x, \cdot) \| \pi)}$$

where KL is the usual KL divergence between two distributions. Or for Gaussian distributions $N(m_1, \sigma_1^2), N(m_2, \sigma_2^2)$, its expression is well known

$$\frac{1}{2} \log \frac{\sigma_2^2}{\sigma_1^2} + \frac{\sigma_1^2 + (m_1 - m_2)^2}{2\sigma_2^2} - \frac{1}{2}$$

In our case, for a given starting state $x \in \mathbb{R}$,

$$\begin{aligned} \text{KL}(P^n(x, \cdot) \| \pi) &= -\frac{1}{2} \log(1 - \theta^{2n}) + \frac{1 - \theta^{2n}}{2} + \frac{x^2 \theta^{2n} (1 - \theta^2)}{\omega^2} - \frac{1}{2} \\ &= -\frac{1}{2} \log(1 - \theta^{2n}) - \frac{\theta^{2n}}{2} + \frac{x^2 \theta^{2n} (1 - \theta^2)}{\omega^2} \\ &\approx \frac{x^2 \theta^{2n} (1 - \theta^2)}{\omega^2} \quad \text{for } n \gg 1 \end{aligned}$$

Therefore

$$\begin{aligned} \beta(n) &= \frac{1}{2} \int_{\mathcal{R}} \pi(dx) \|P^n(x, \cdot) - \pi\|_{TV} \\ &\leq \frac{1}{2} \int_{\mathbb{R}} \pi(dx) \sqrt{\frac{x^2 \theta^{2n} (1 - \theta^2)}{2\omega^2}} \\ &= \frac{\theta^n}{2} \sqrt{\frac{1 - \theta^2}{2\omega^2}} \int_{\mathbb{R}} |x| \pi(dx) \\ &= \theta^n \sqrt{\frac{1 - \theta^2}{2\omega^2}} \end{aligned}$$

Note also that if the initial distribution ν_0 is Gaussian then $\nu_0 P^n$ will also be Gaussian. Therefore given the geometric ergodicity function $Q(x) = |x|$, the condition of Theorem 5.1, $\int_{\mathbb{R}} Q(x) \nu_0 P^N(x) dx < \infty$ will be verified for any N .

B.3. Real-world dataset




Figure 6. Exchange rate EUR/SEK




Figure 7. French electricity price